

Explaining complex technology problems in a simplistic way is a skill I've been working to hone. Marketing approached me to explain to all of their organization "containers" in a simplistic and fun way, as part of a series they called "The Stark Hour", after Tony Stark.

I enjoyed making this deck so much, and the feedback was so positive, it's been used in countless other customer conversations as well as new hire, college graduate programs and more.

This presentation is part of a larger experiment in embracing public, visual personal brand building, available at:

<https://mattschneider-visualcv.github.io/>

VisualCV started as a Pathfinder's project at Dell Technologies while I was mentoring engineers through our career ladder into roles requiring panel & packet review, Principals & Distinguished.



# CONTAINERS

# CONTAINERS





# CONTAINERS



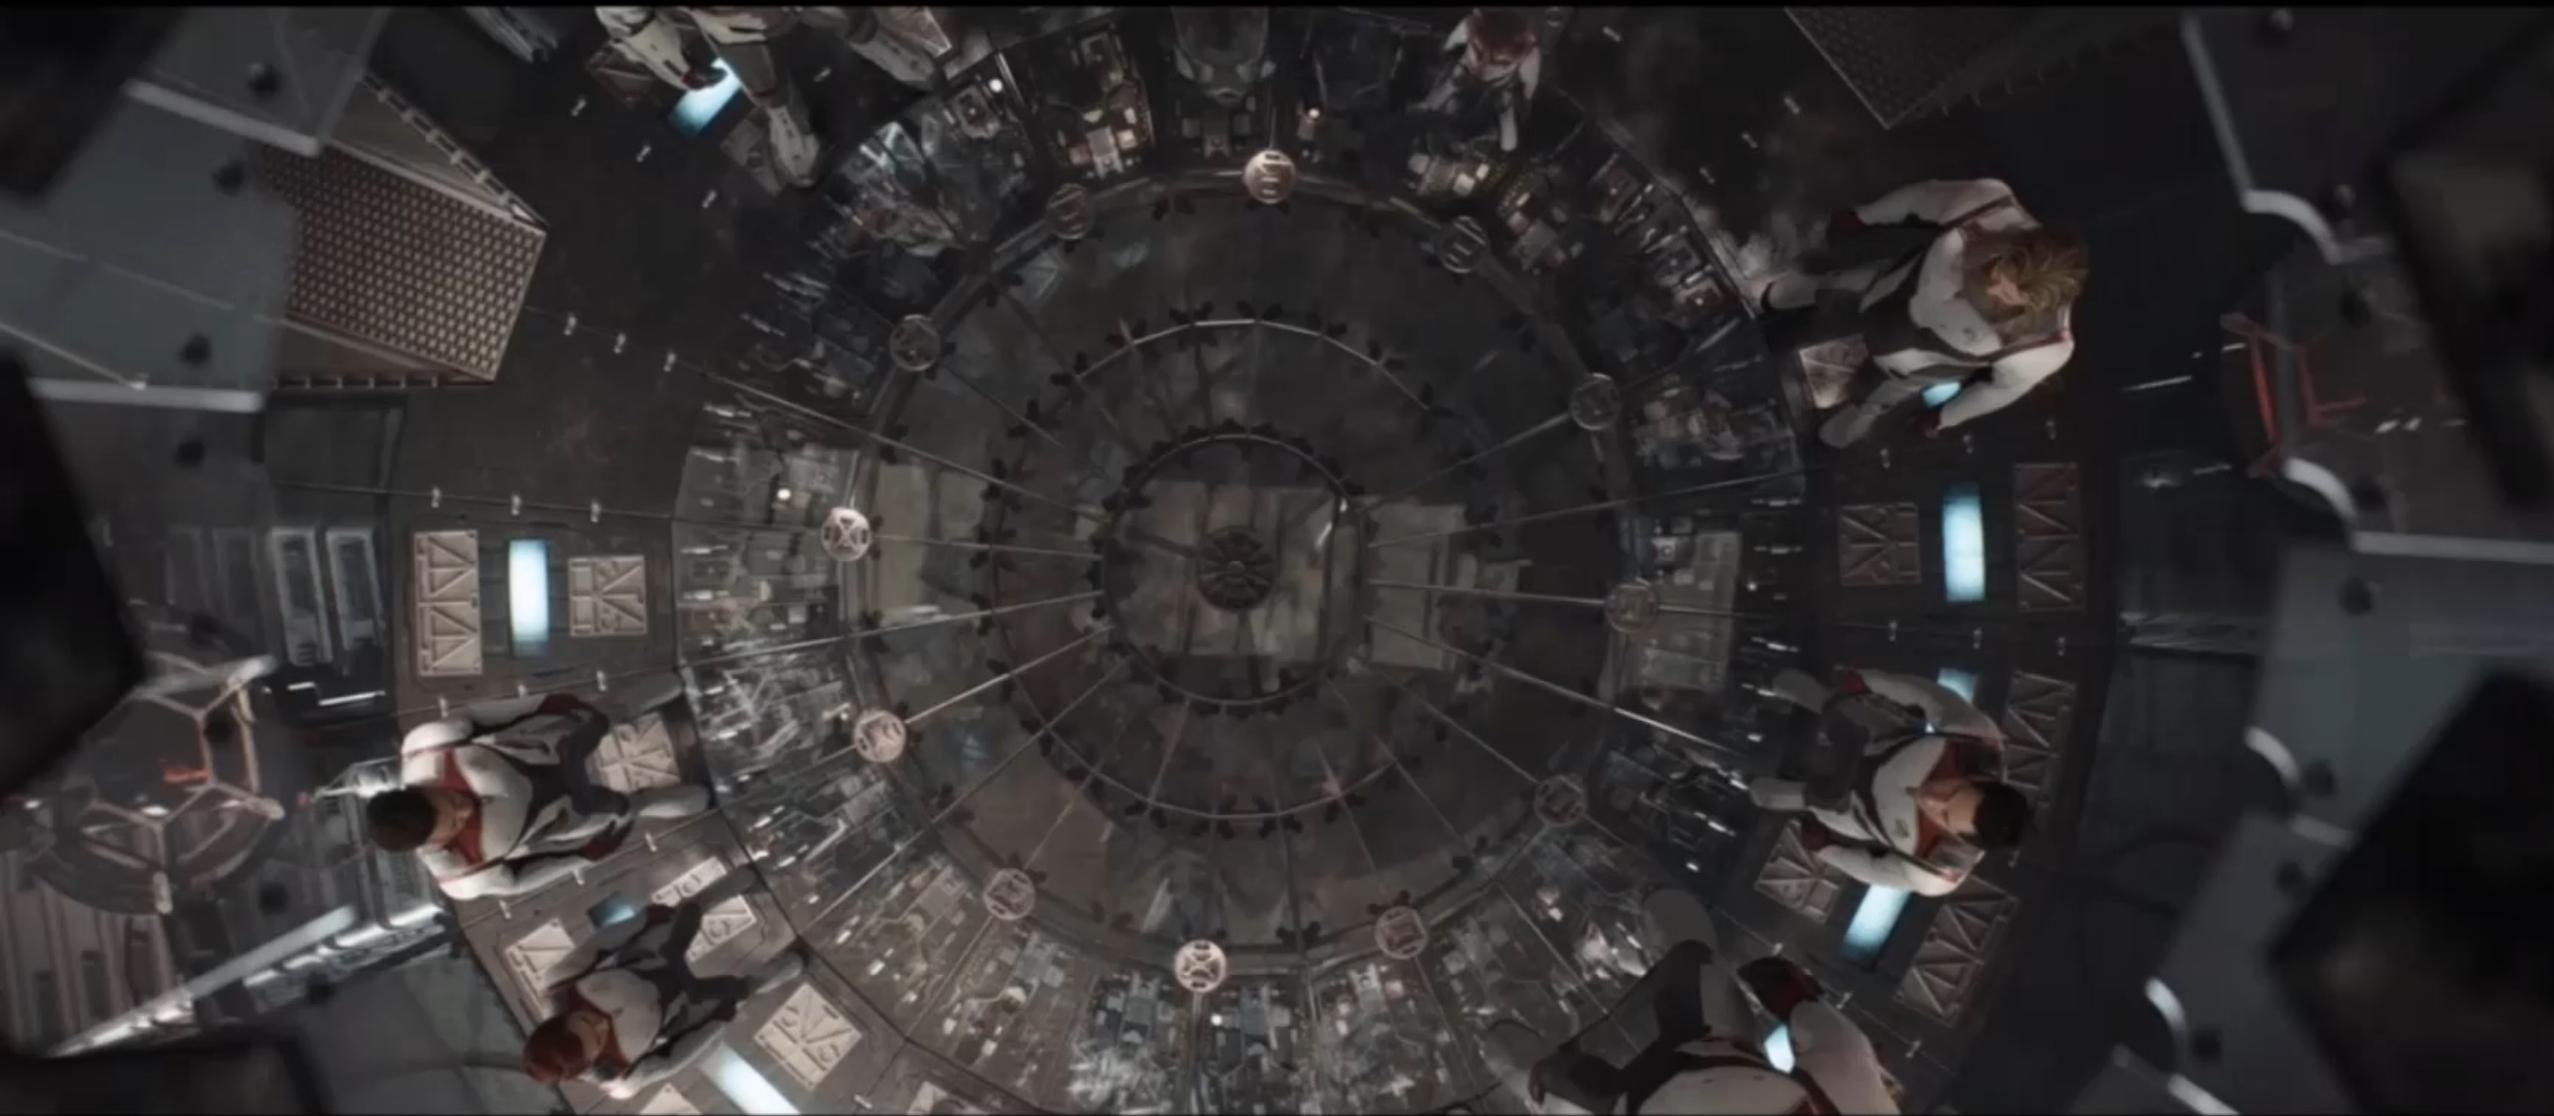


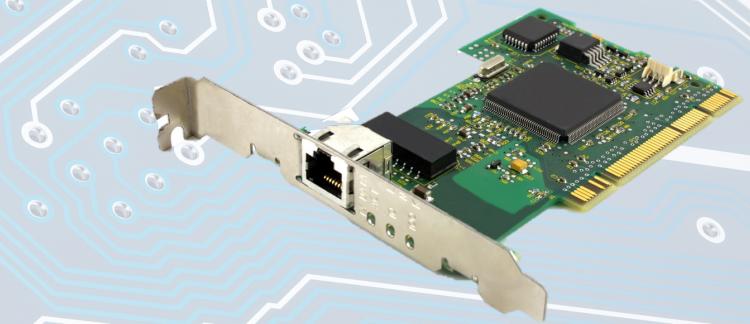
Matt Schneider  
Field CTO & Principal Engineer  
Enterprise Americas Presales

 @md\_schneider  
 /in/mdschneider

★★  
CTO Ambassadors  
• 2020 •







## Hardware



### Processor

CPU  
GPU  
FPGA  
ASIC



### Storage

PMEM  
RAM  
NVME/SSD  
Hard Drive



### Network

Ethernet  
Fibre Channel  
Infiniband  
RDMA

## Operating System



## Hardware



Processor



Storage



Network

CPU  
GPU  
FPGA  
ASIC

PMEM  
RAM  
NVME/SSD  
Hard Drive

Ethernet  
Fibre Channel  
Infiniband  
RDMA

## Platforms



## Operating System



## Hardware



Processor



Storage

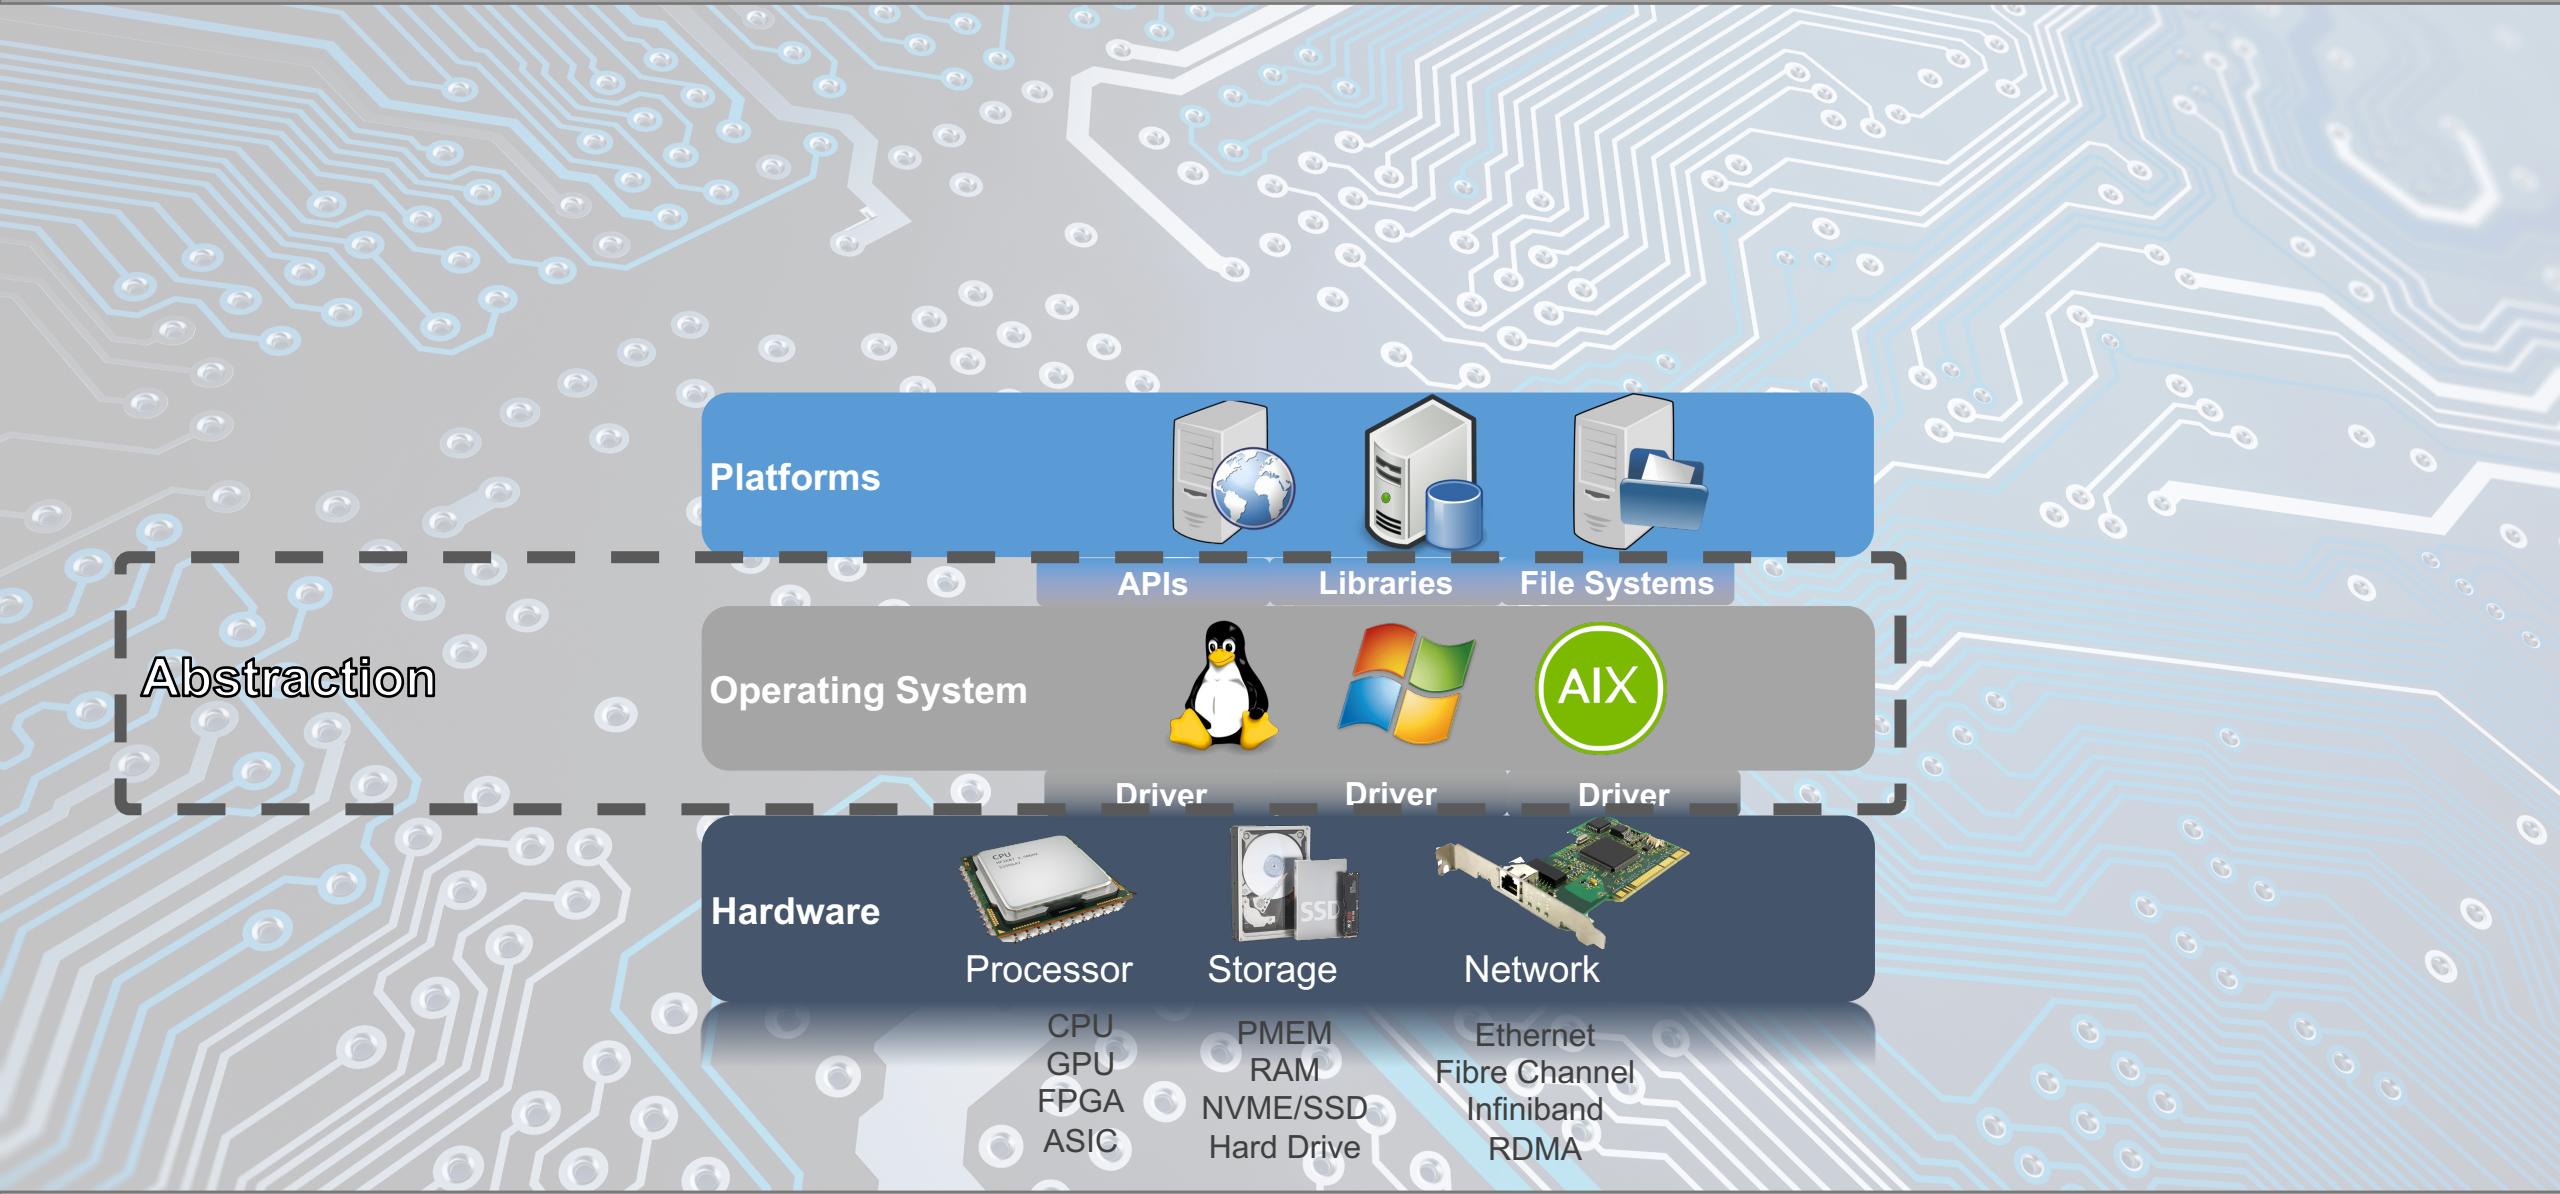


Network

CPU  
GPU  
FPGA  
ASIC

PMEM  
RAM  
NVME/SSD  
Hard Drive

Ethernet  
Fibre Channel  
Infiniband  
RDMA





## Platforms

APIs

Libraries

File Systems

## Operating System

Driver

Driver

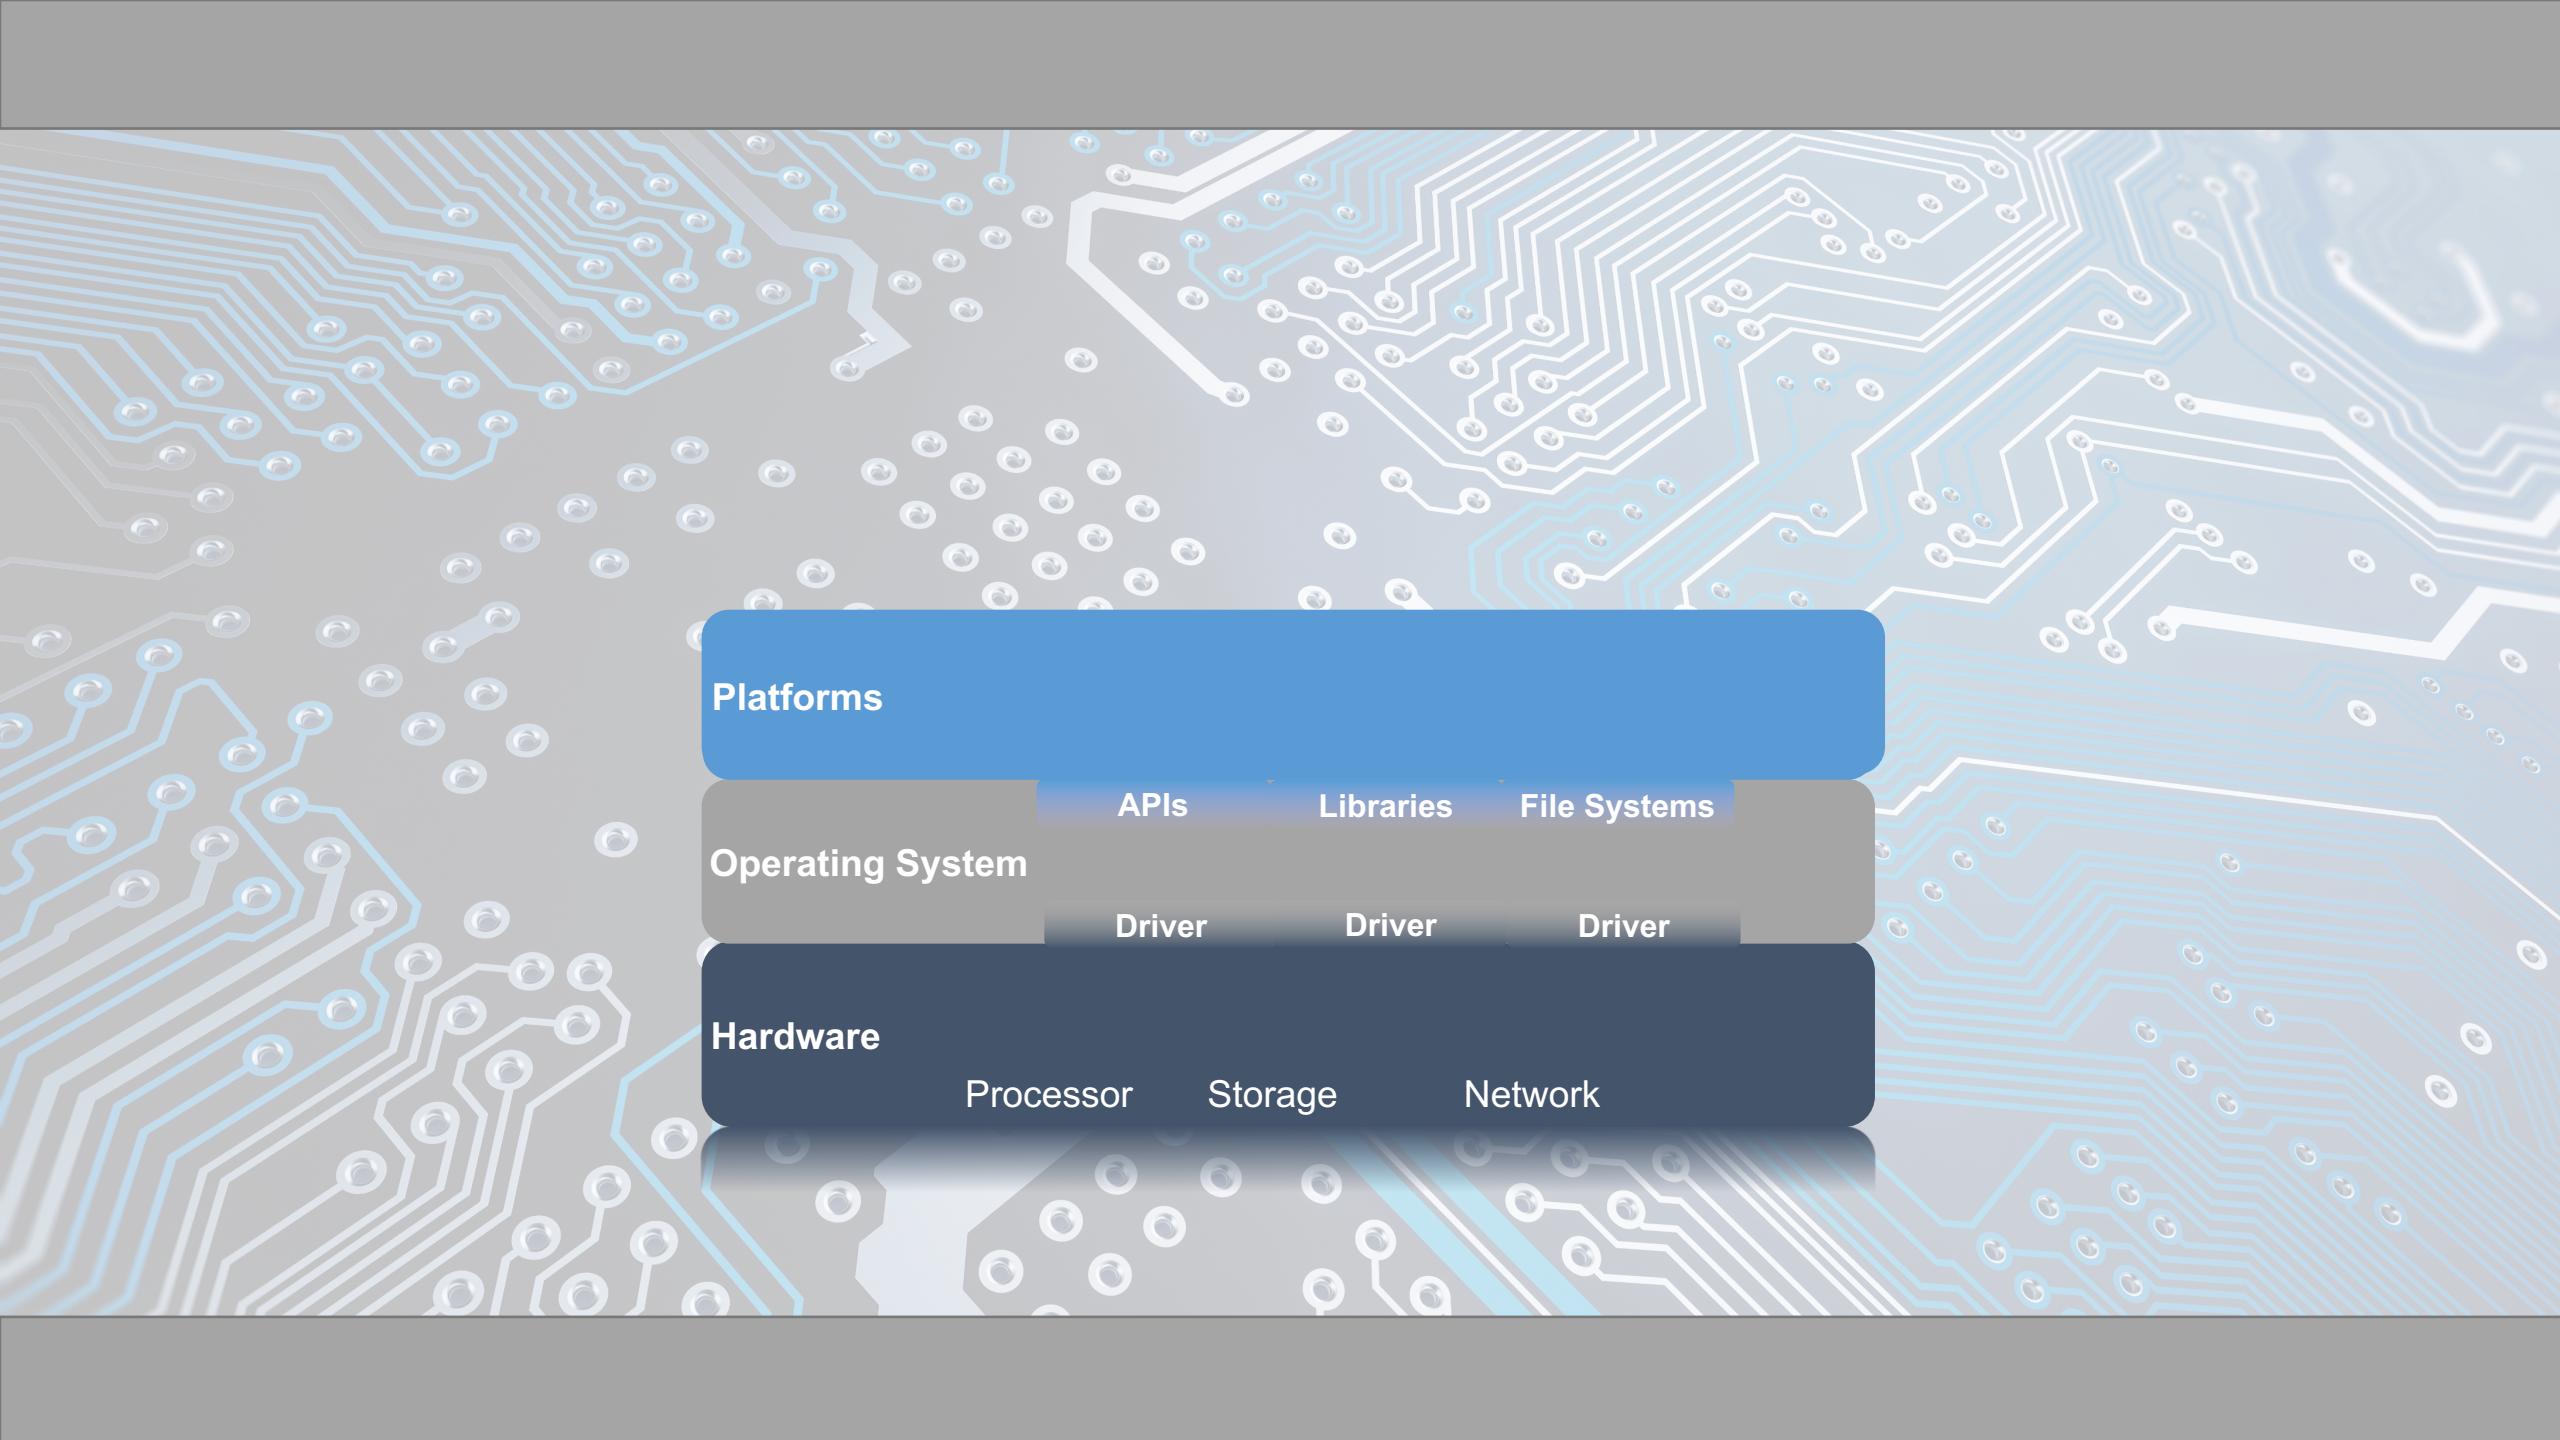
Driver

## Hardware

Processor

Storage

Network



## Line of Business Application

Web Server

Database

APIs

Libraries

File Systems

Operating System

Driver

Driver

Driver

Hardware

Processor

Storage

Network

## Line of Business Application

Web Server

Database

APIs

Libraries

File Systems

Operating System

Driver

Driver

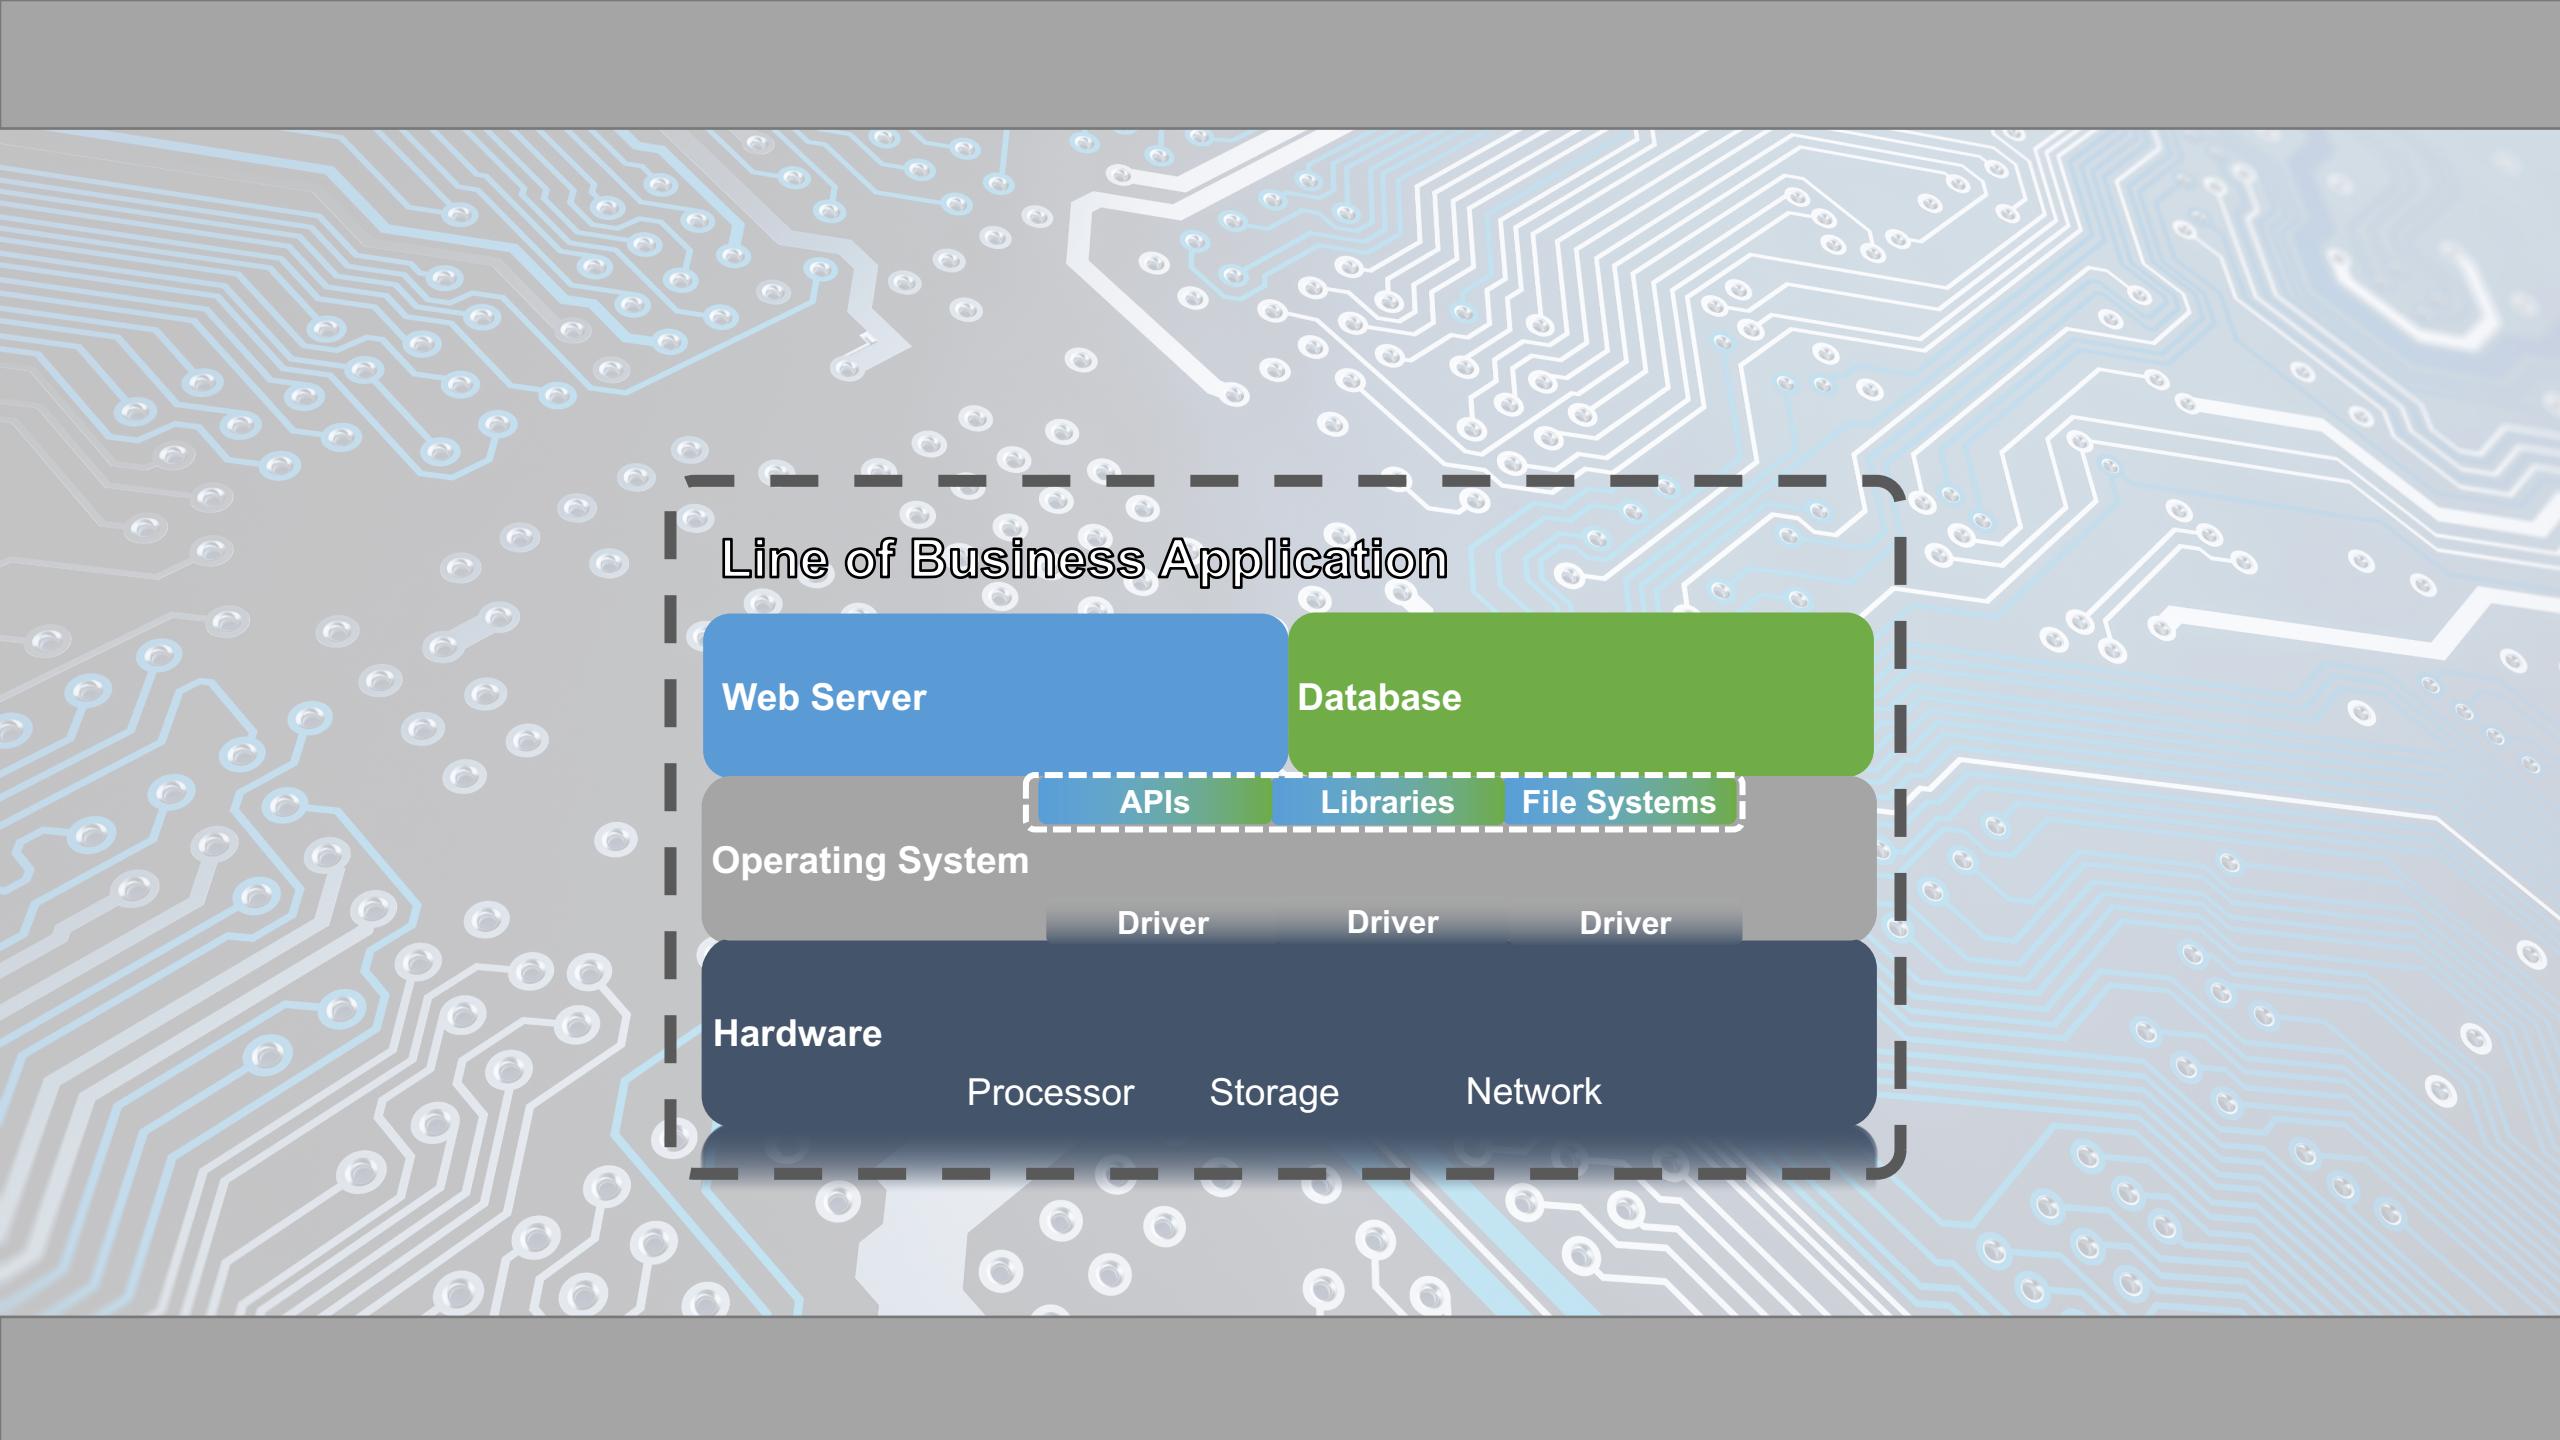
Driver

Hardware

Processor

Storage

Network



**Line of Business Application**

Web Server

Database

Operating System

APIs

Libraries

File Systems

Driver

Driver

Driver

Hardware

Processor

Storage

Network

## Line of Business Application

Web Server

APIs

Libraries

File Systems

Operating System

Driver

Driver

Driver

Hardware

Processor

Storage

Network

Database

APIs

Libraries

File Systems

Operating System

Driver

Driver

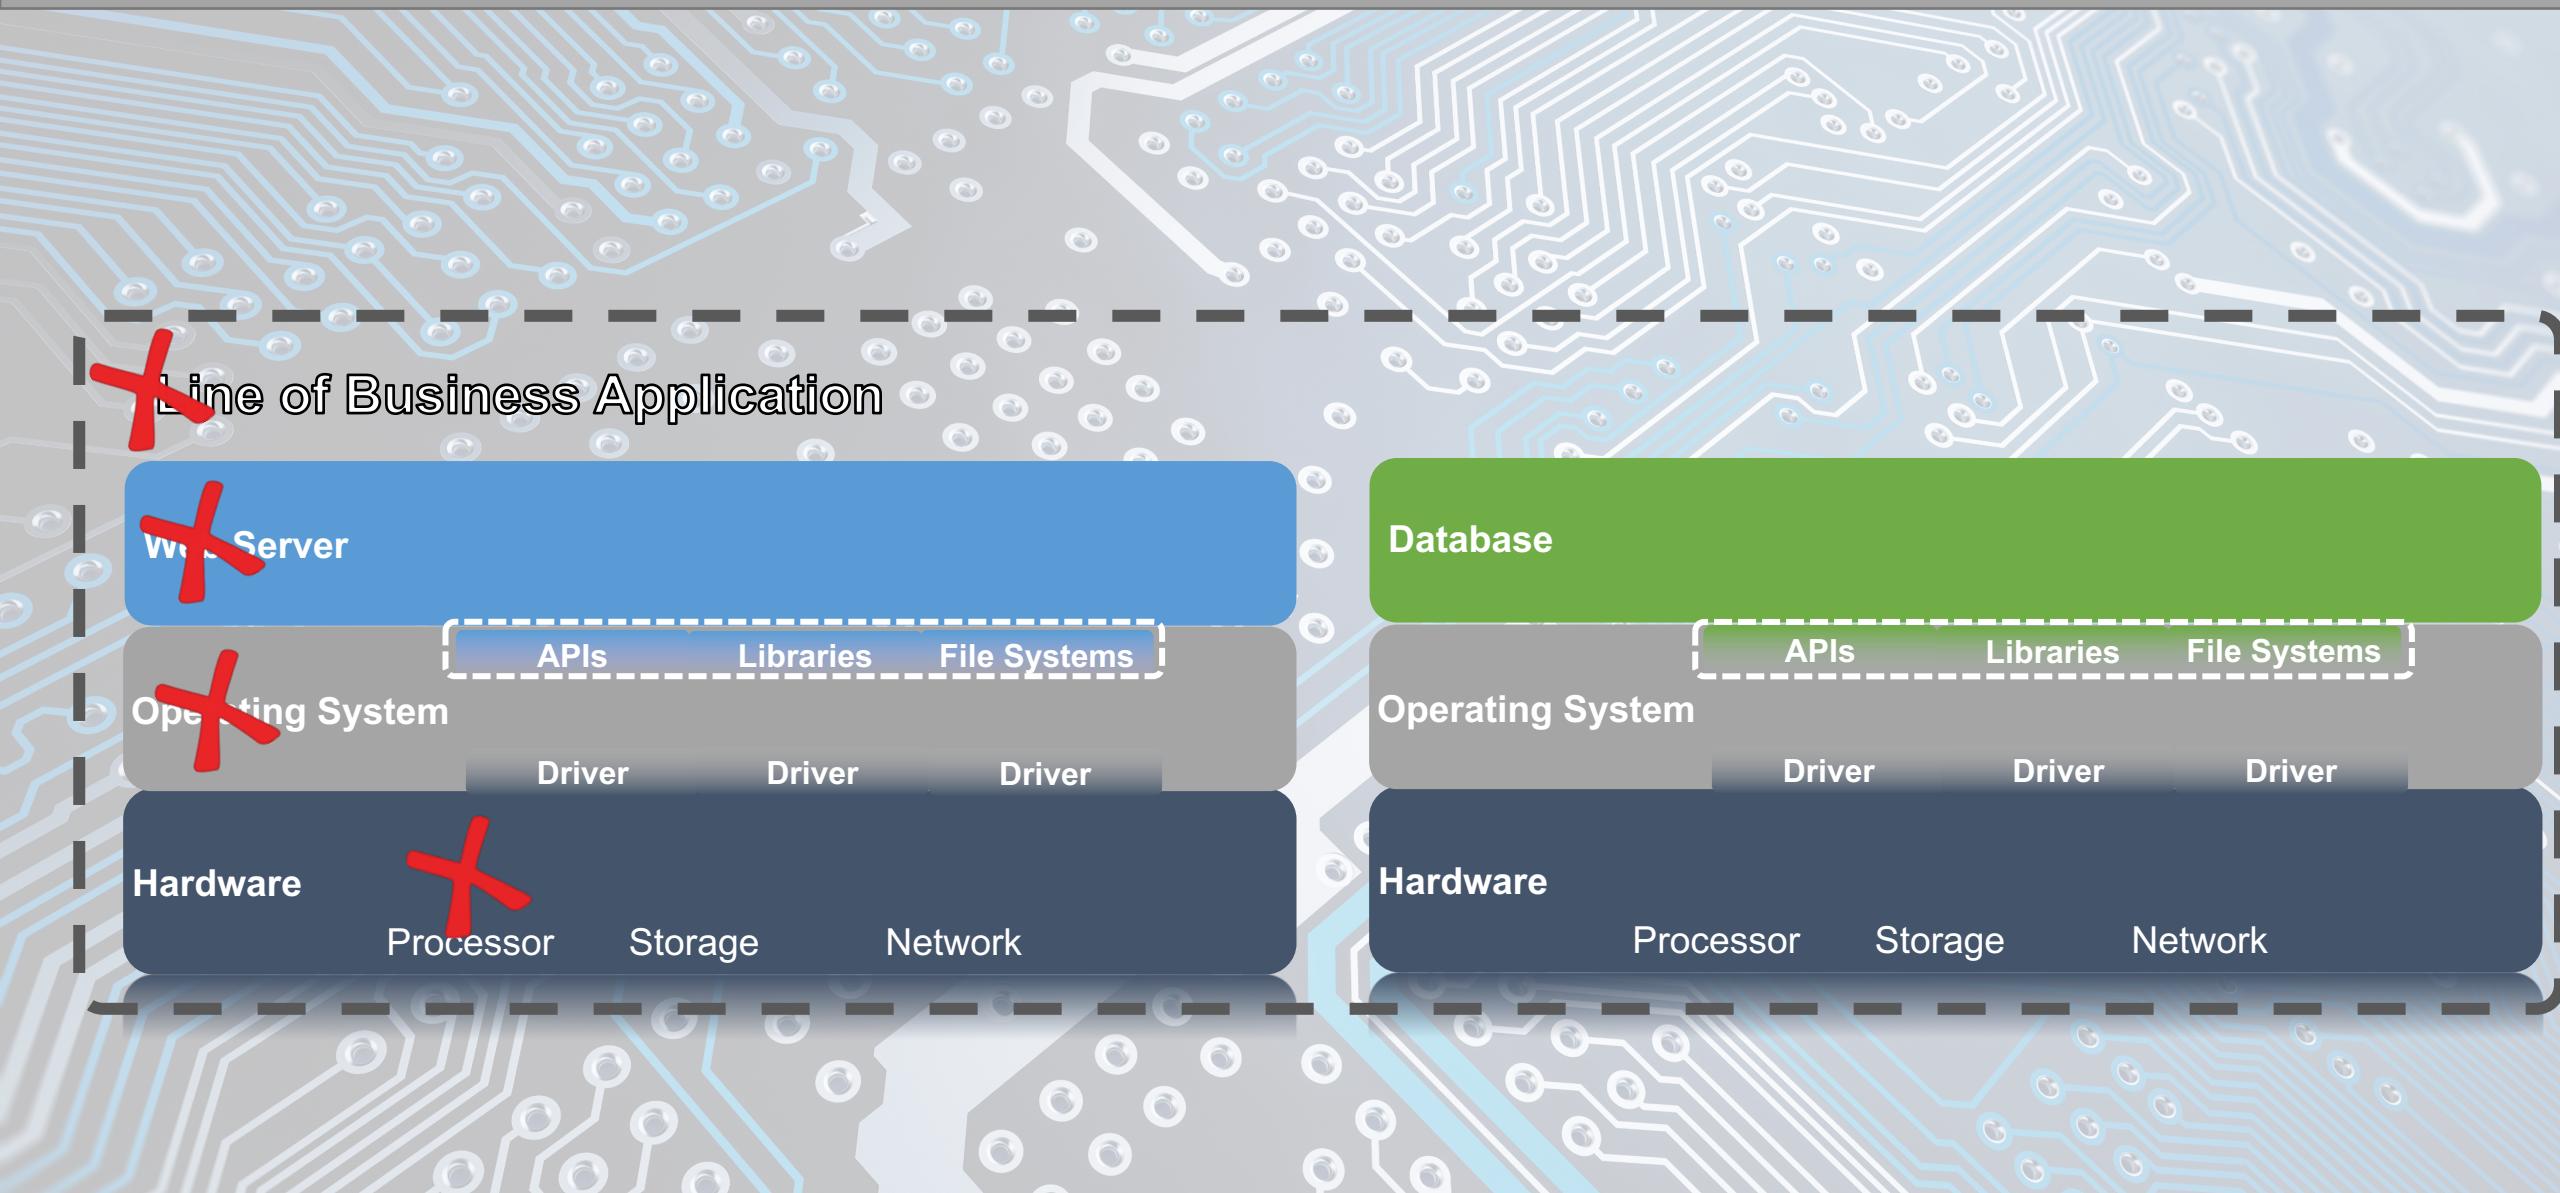
Driver

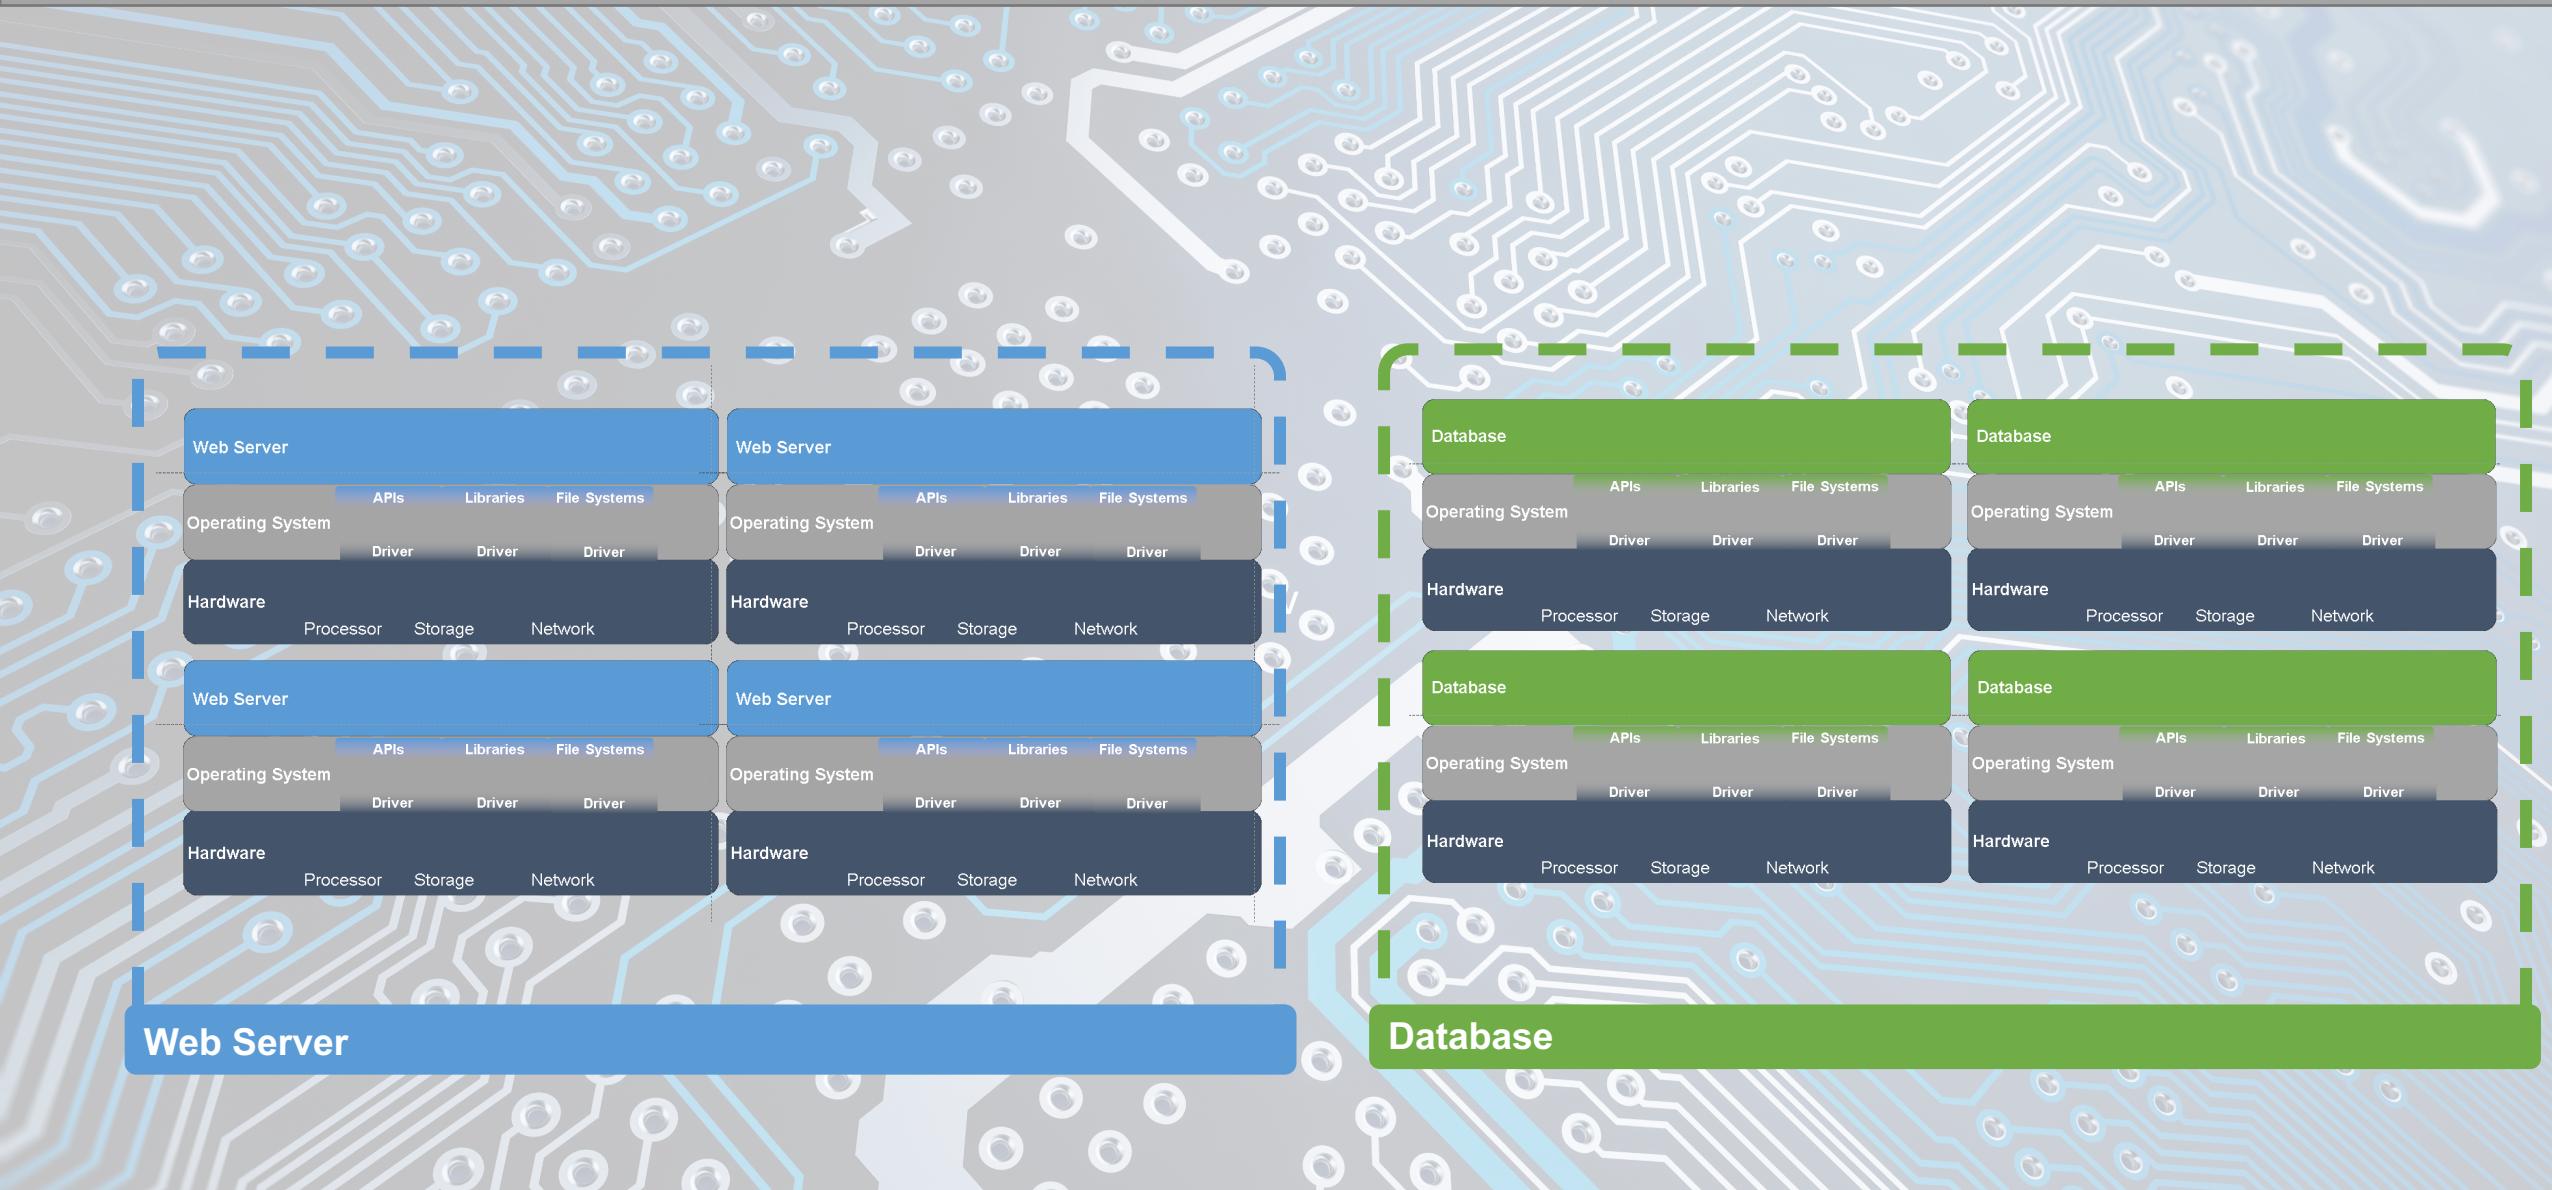
Hardware

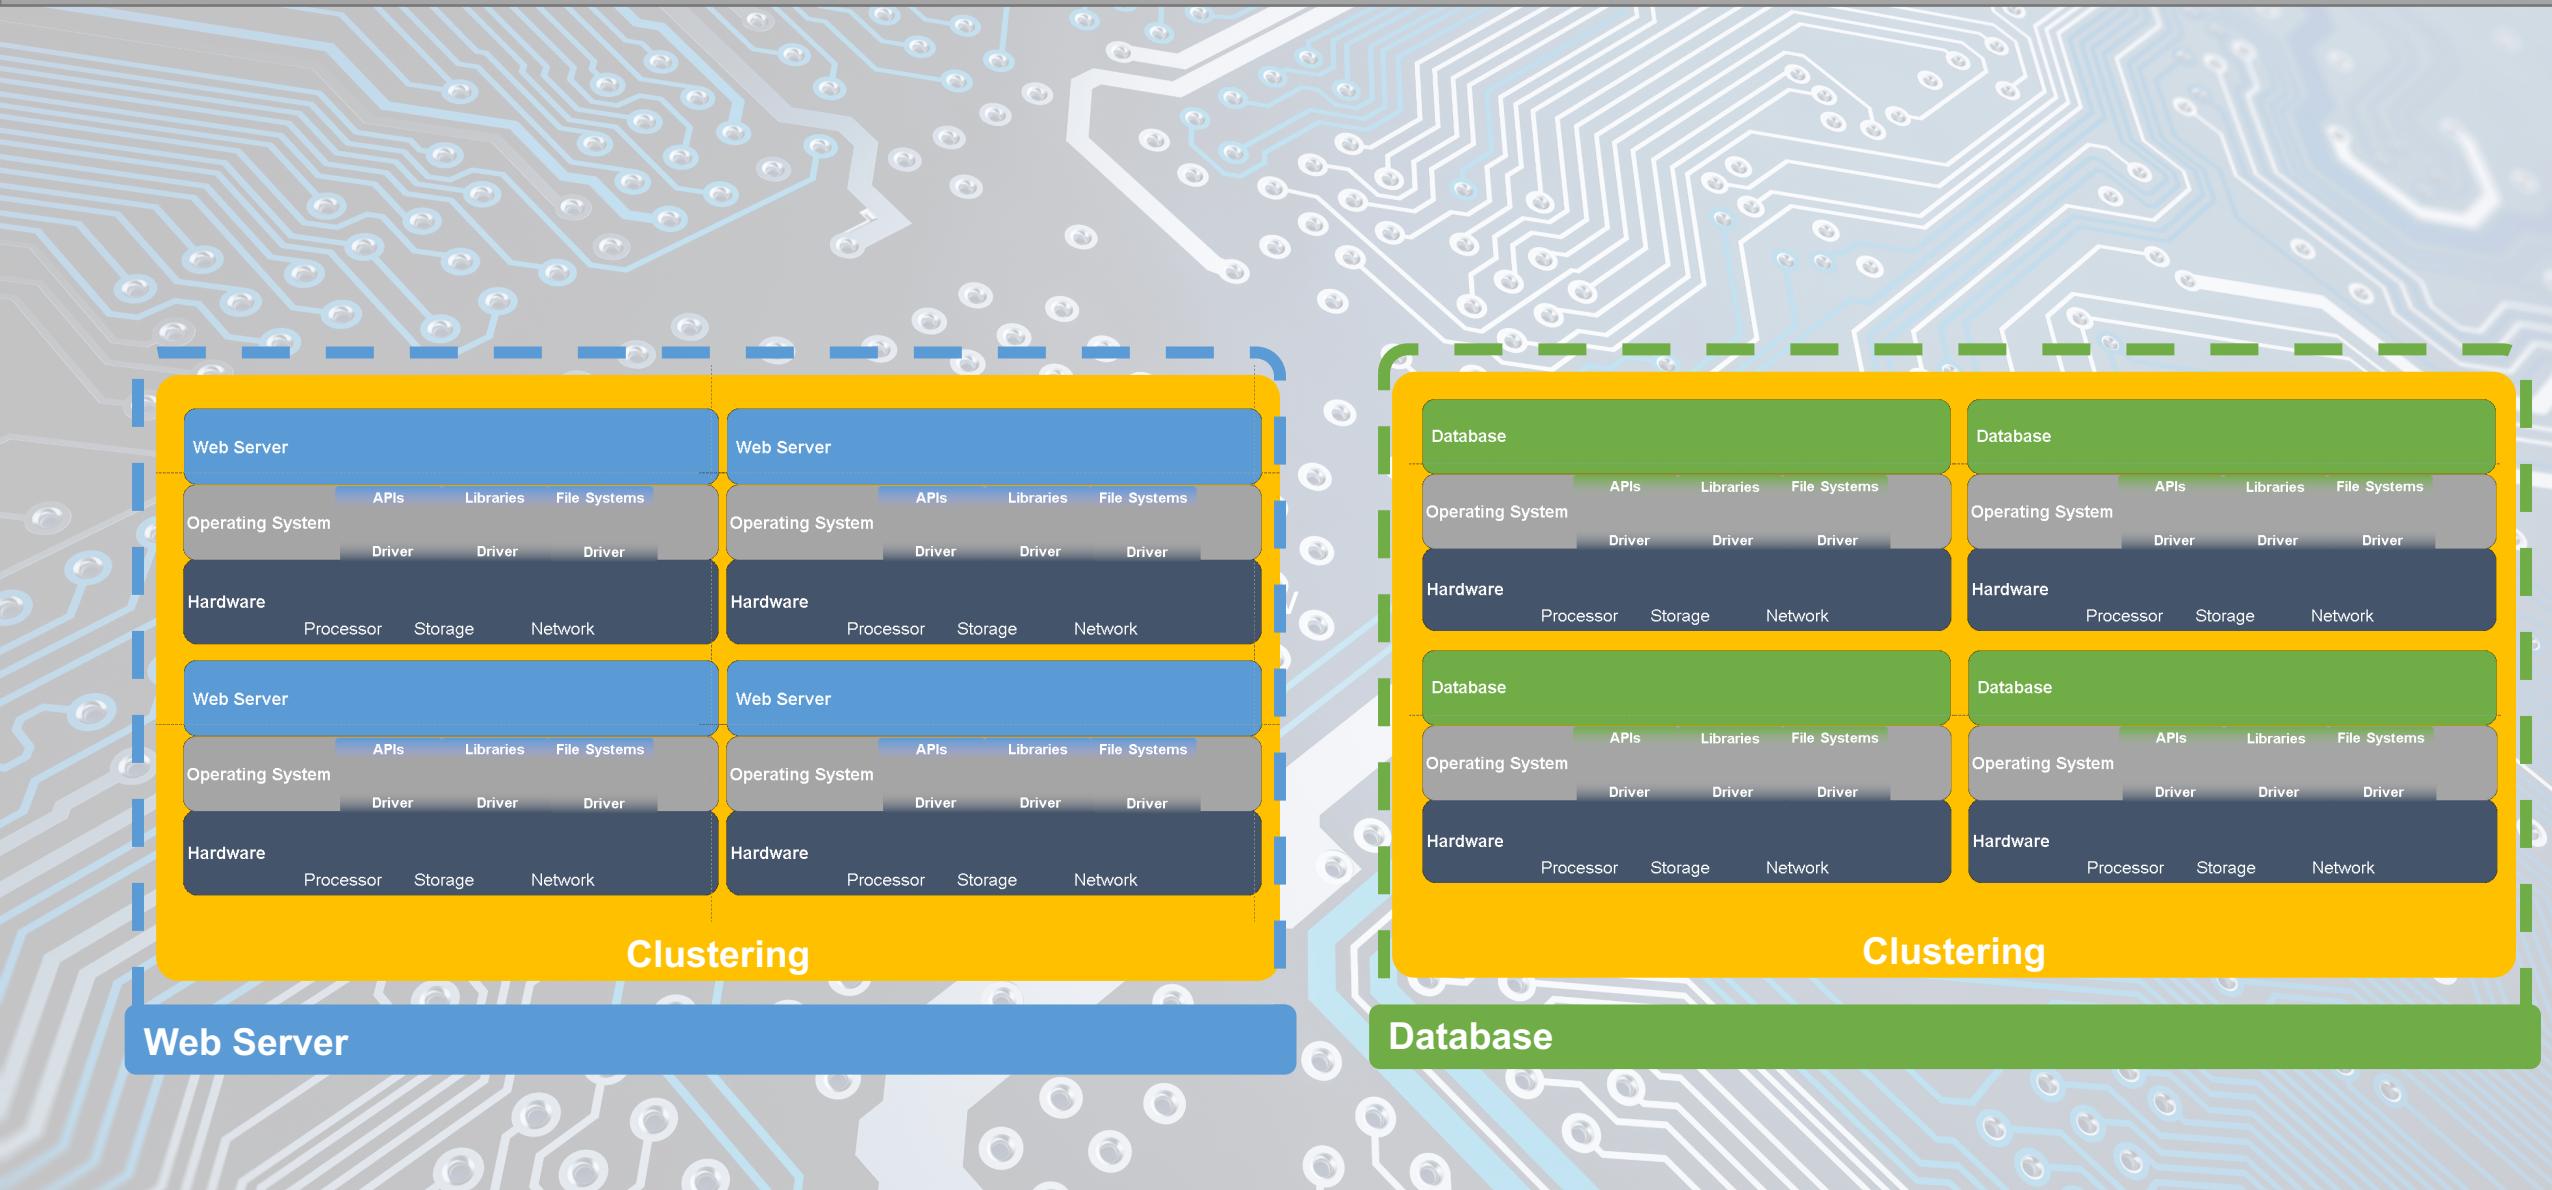
Processor

Storage

Network





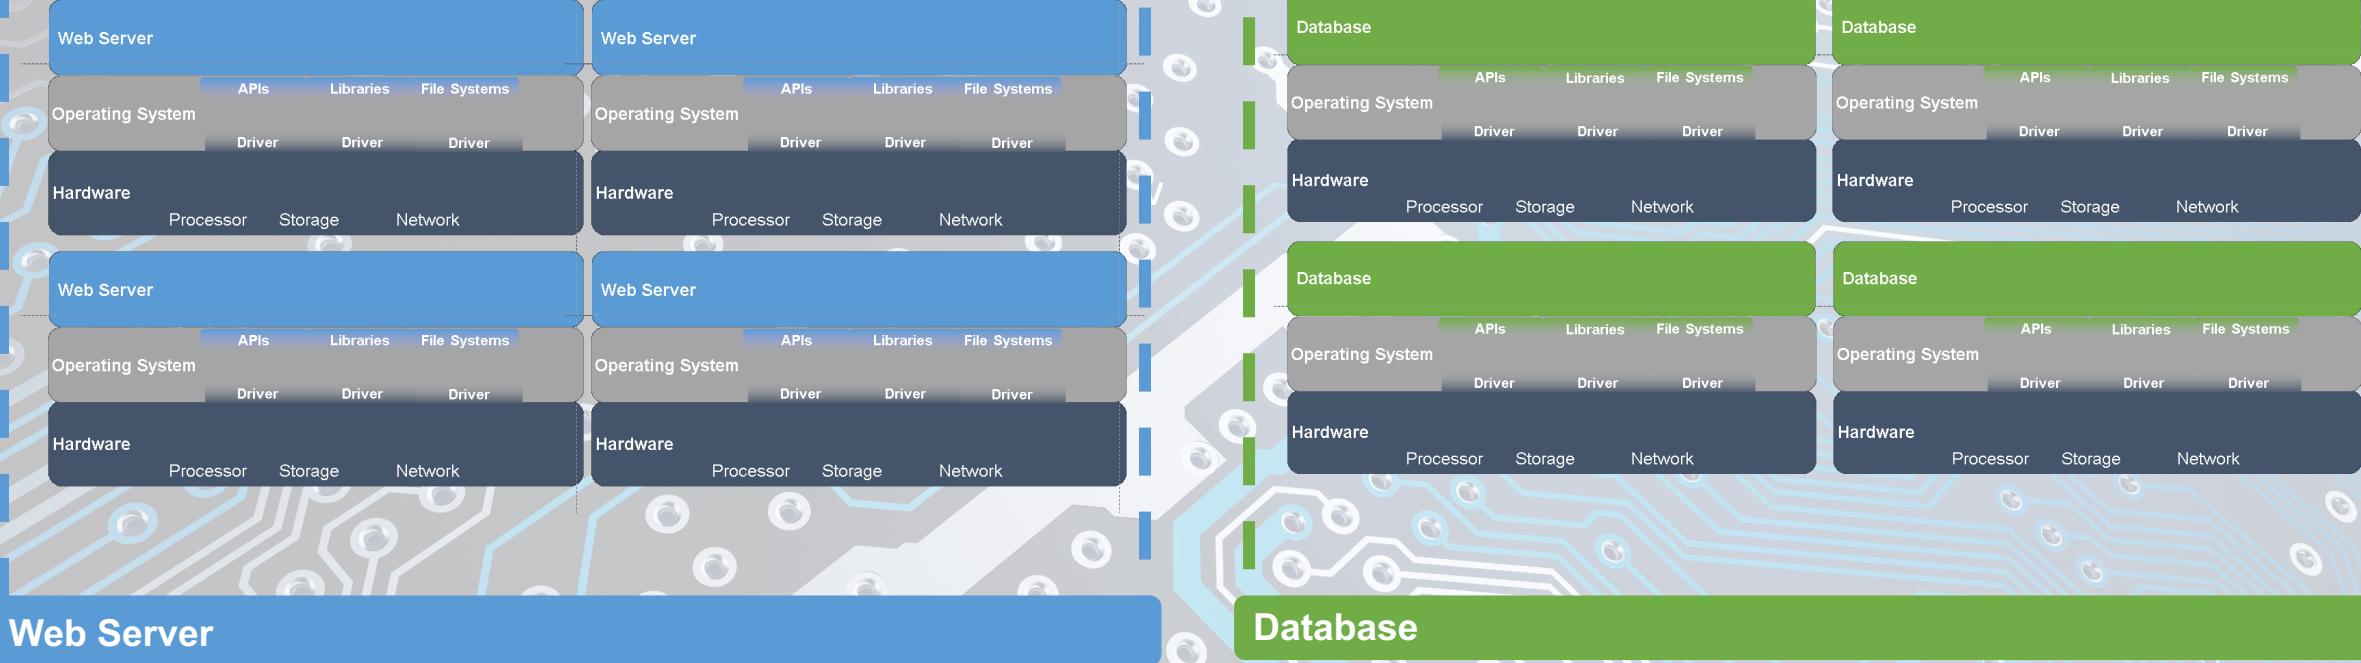




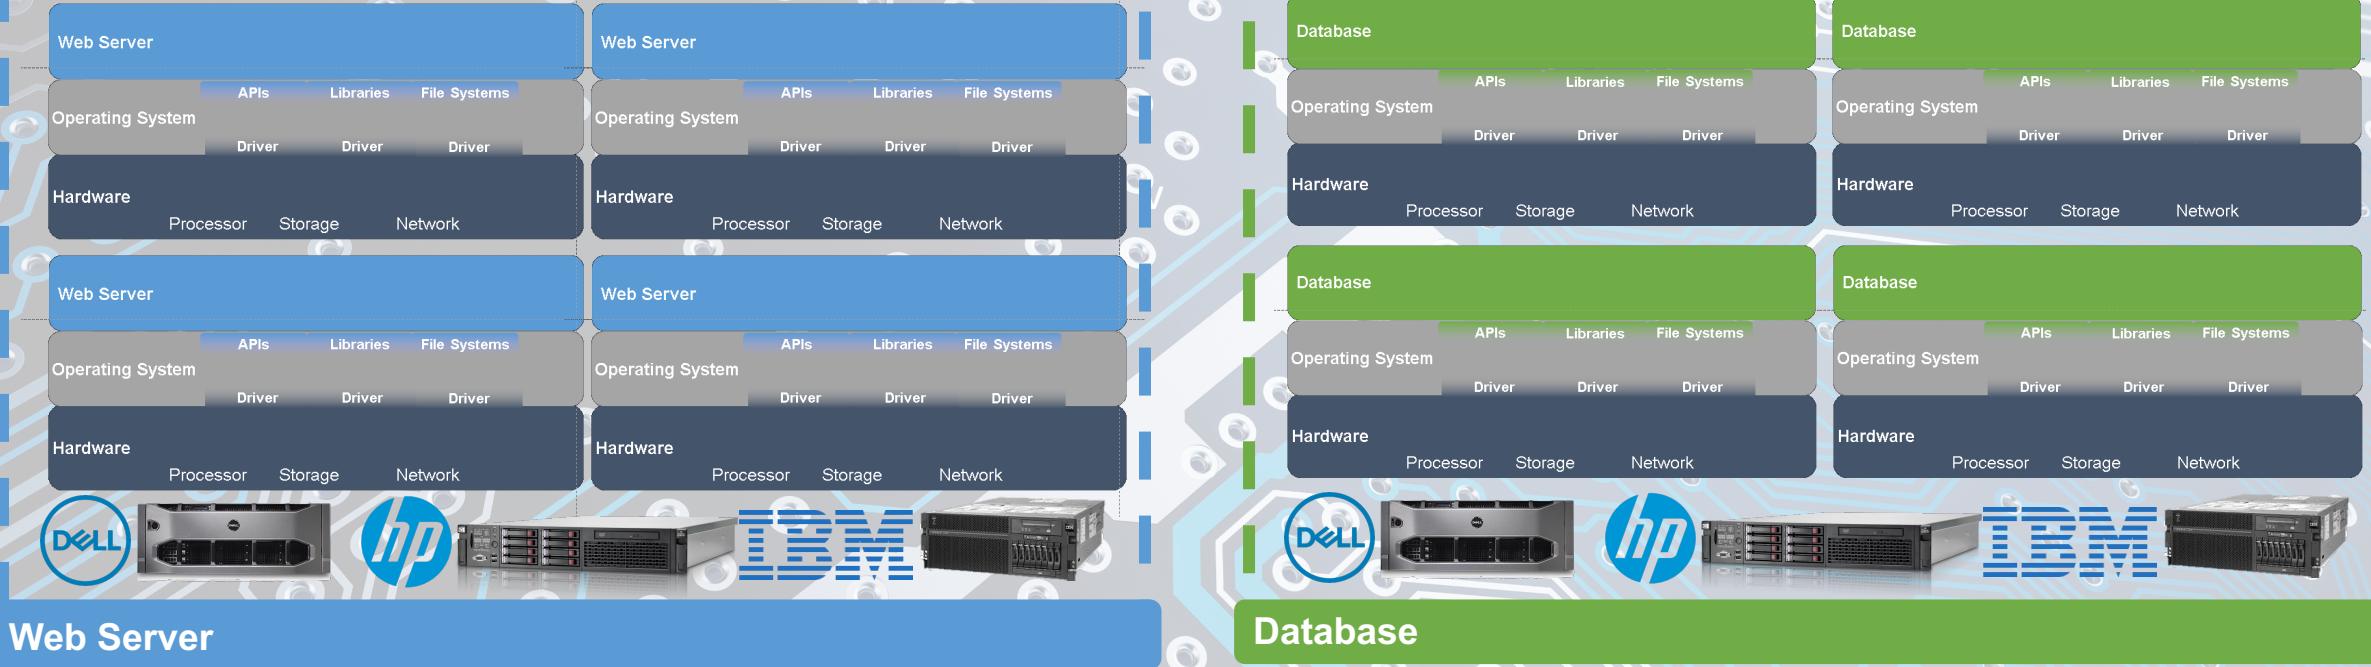
✓ LOB APP



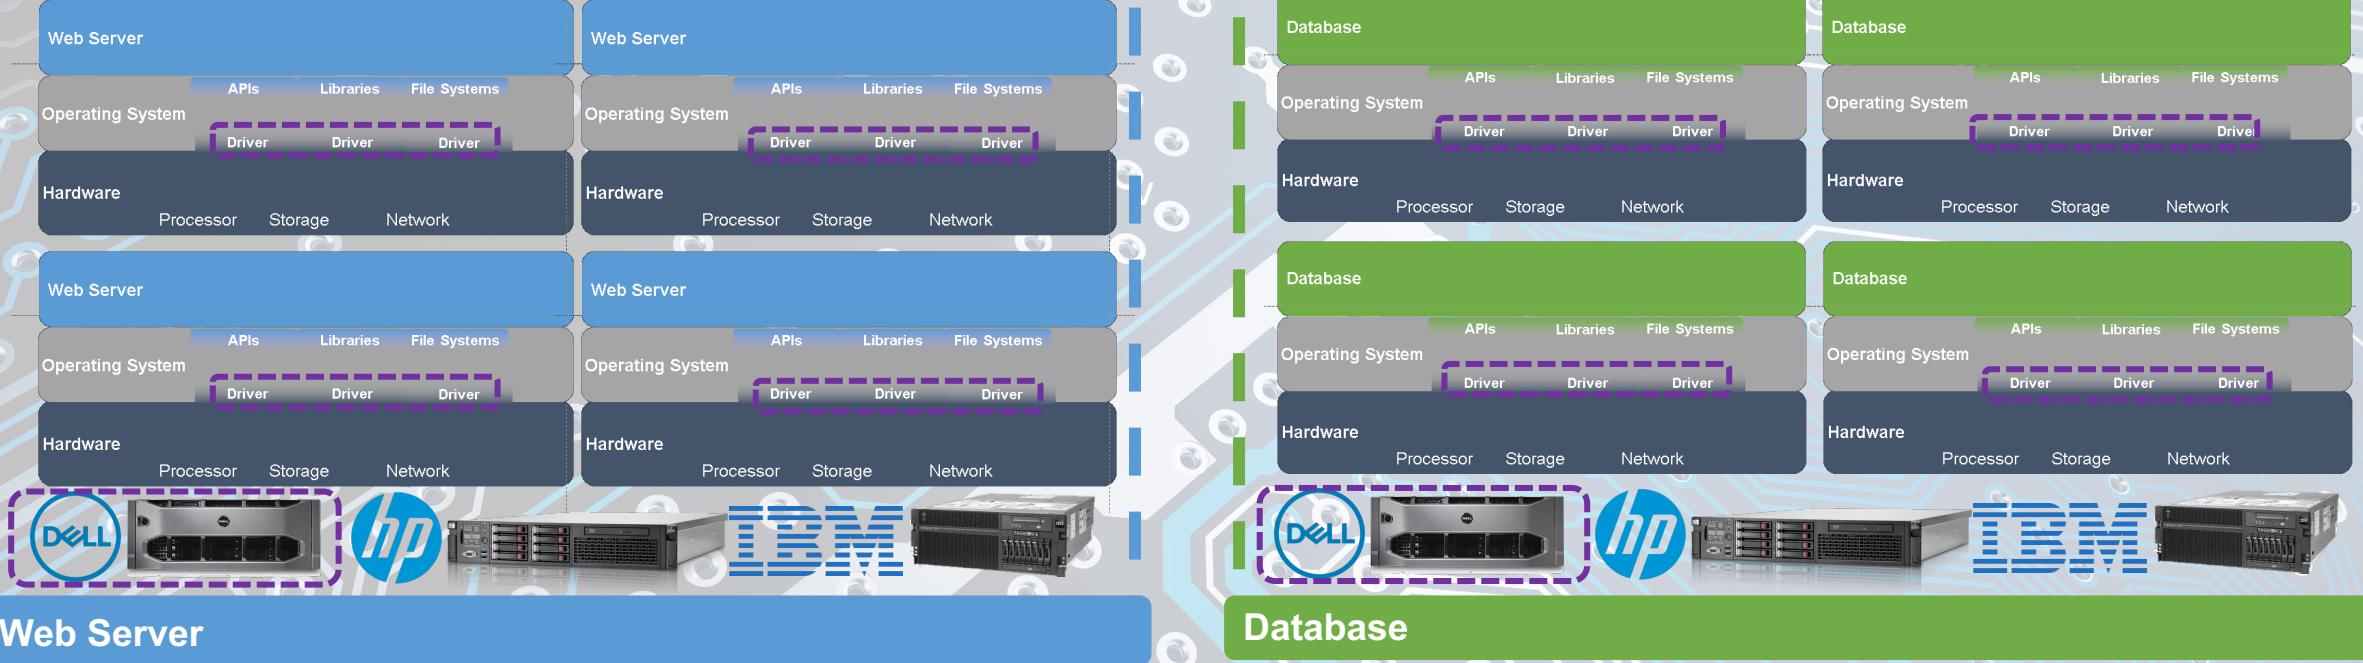
# LOB APP



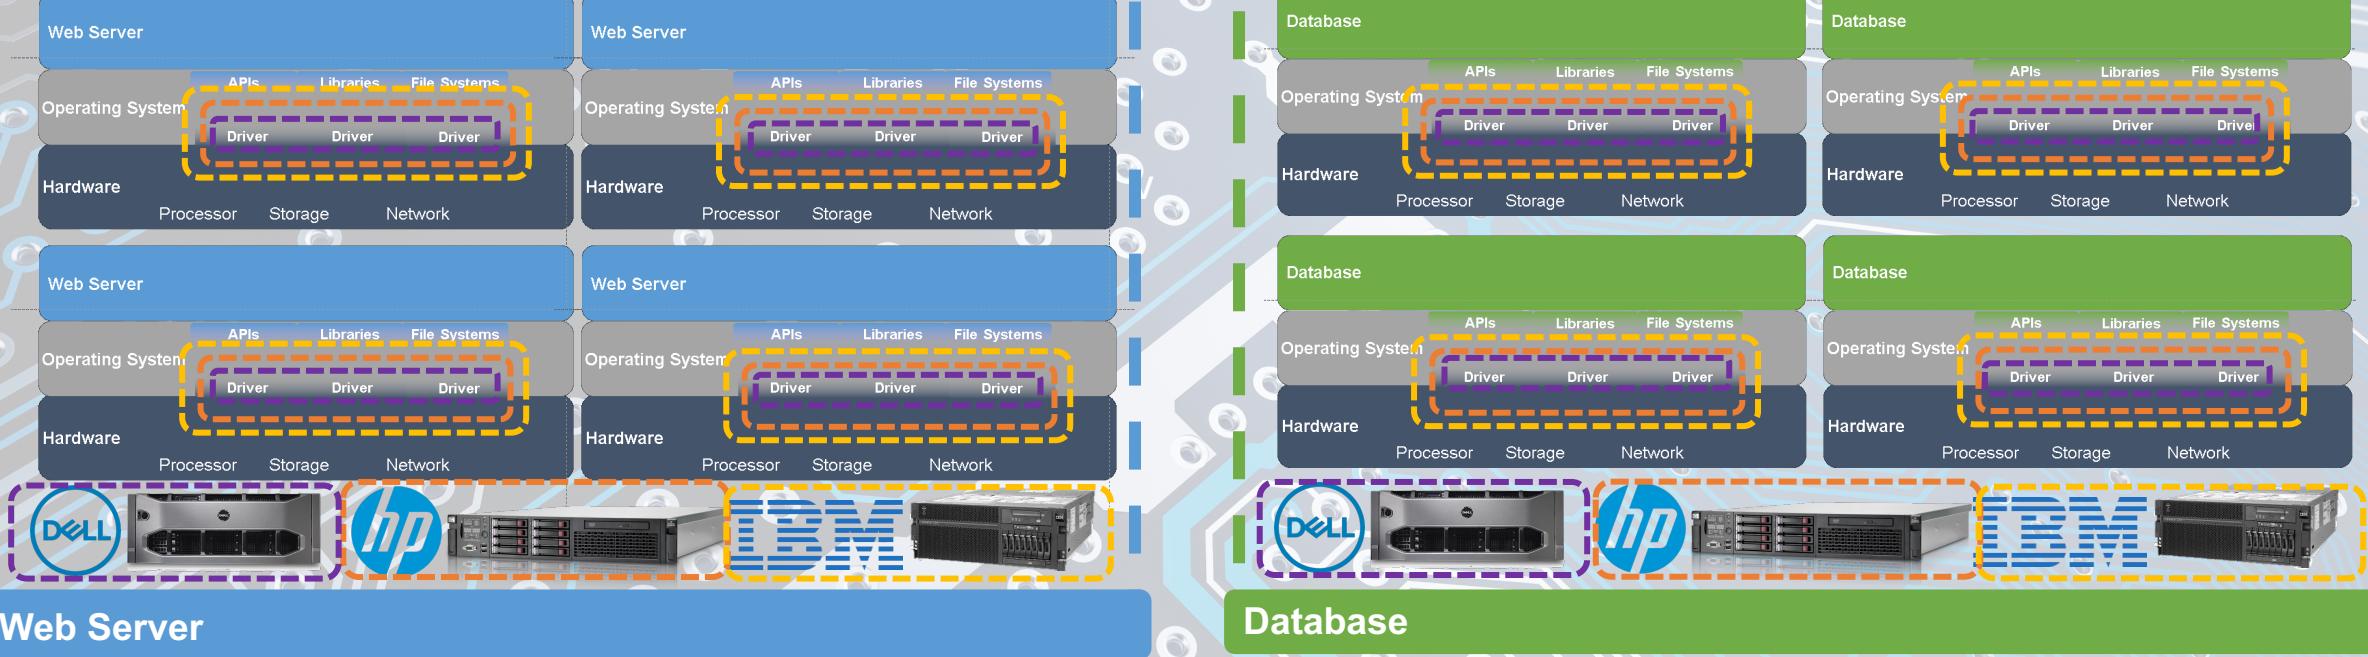
# LOB APP

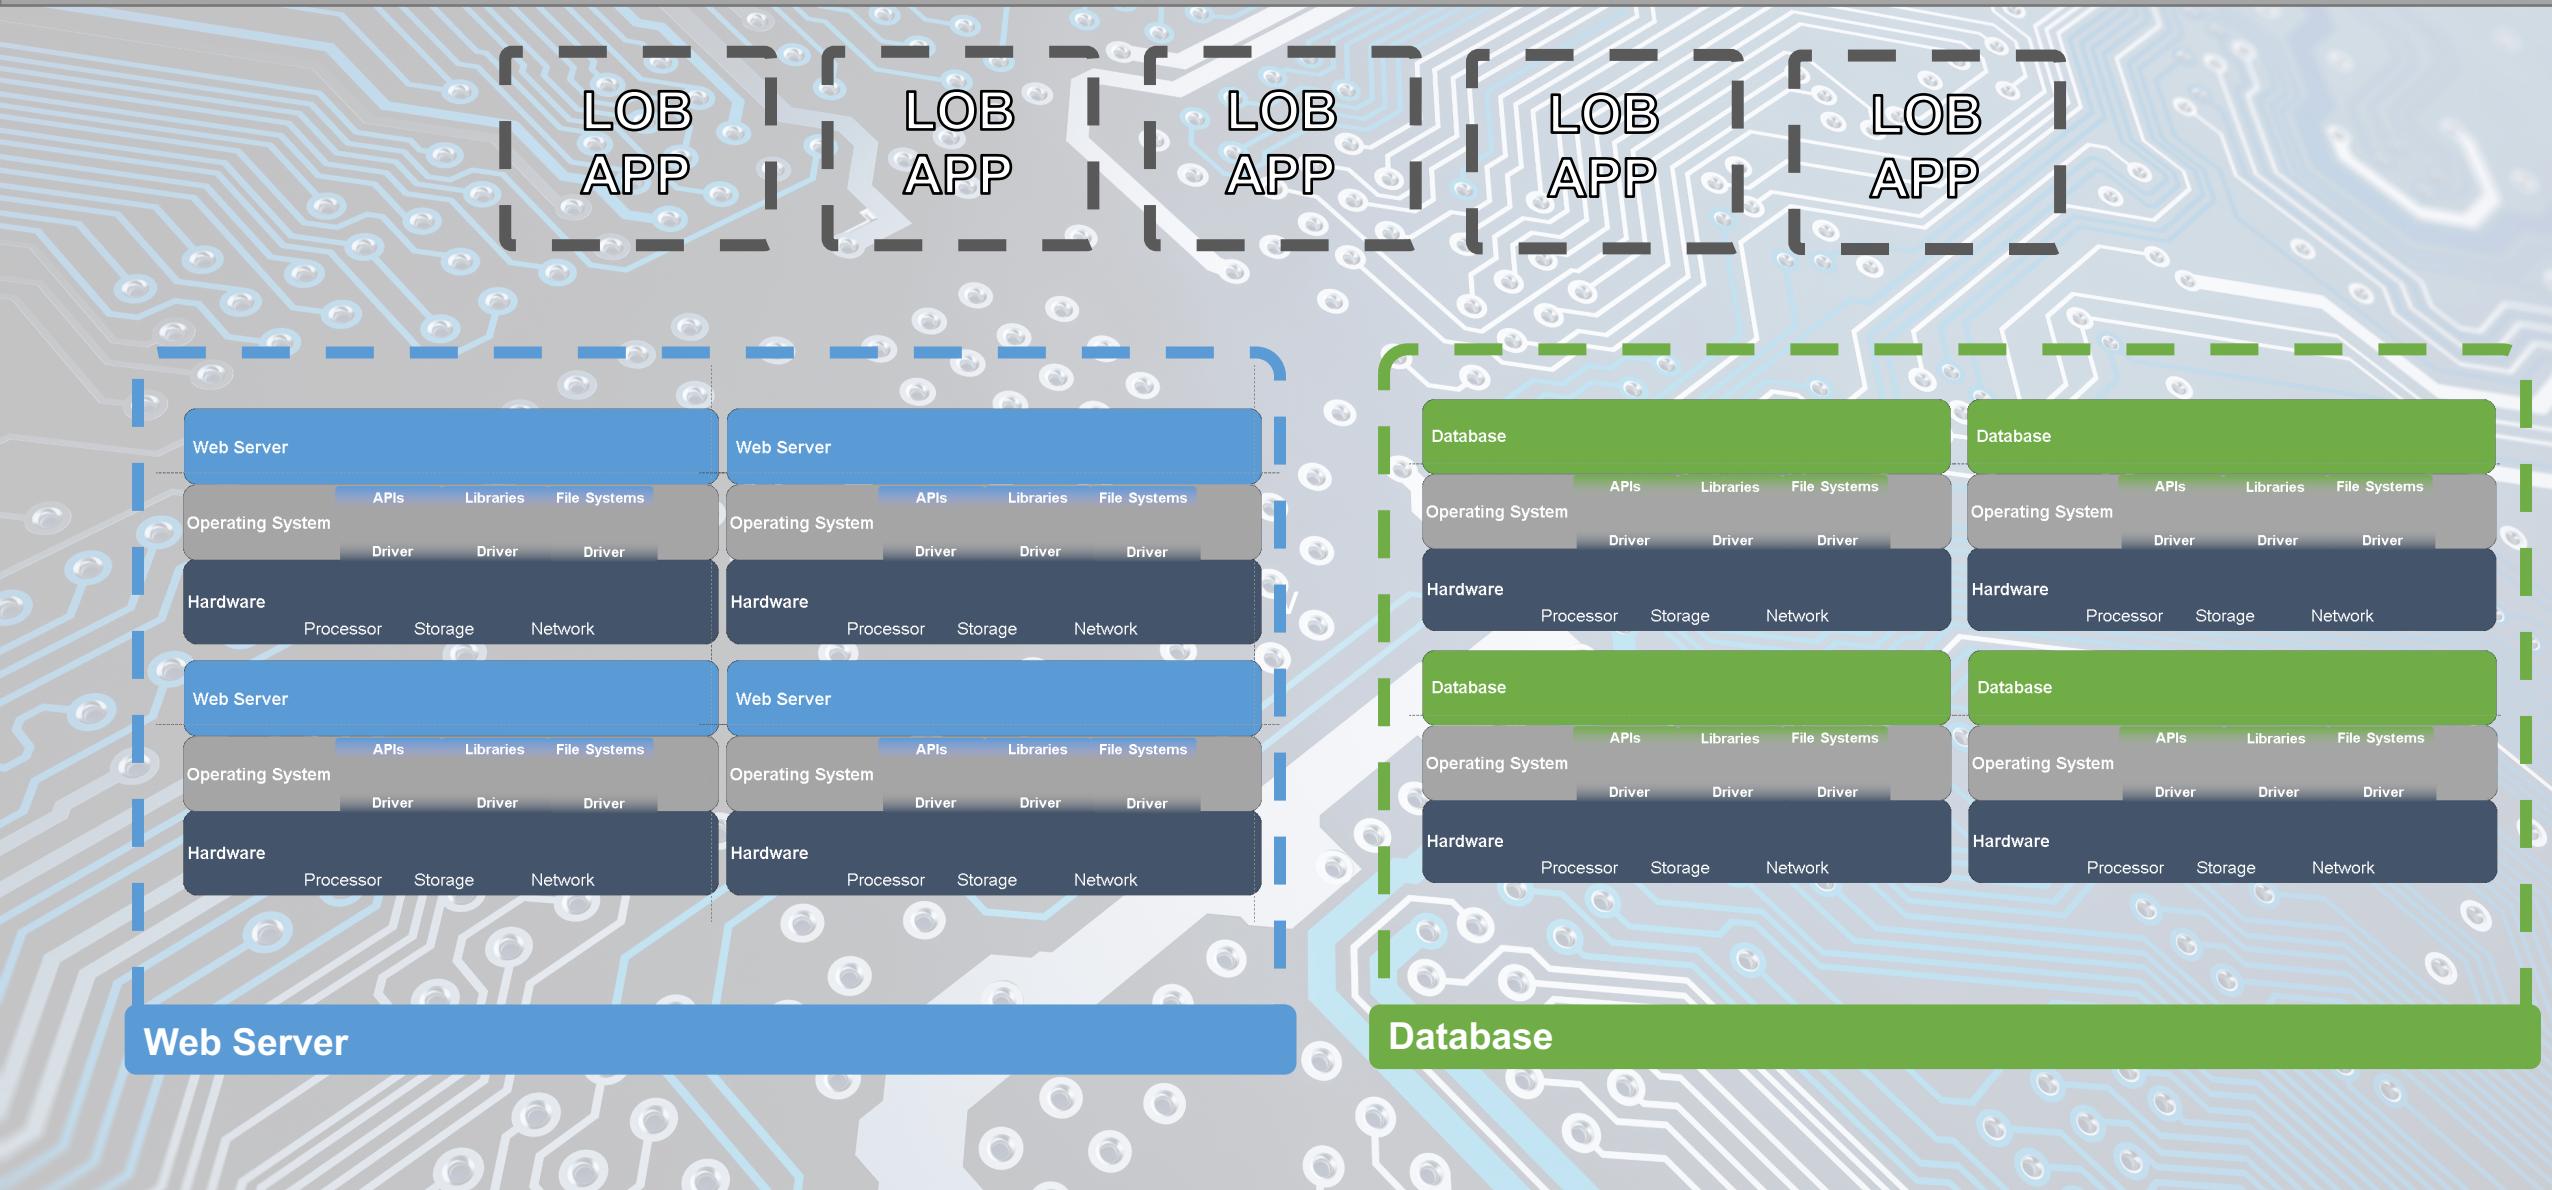


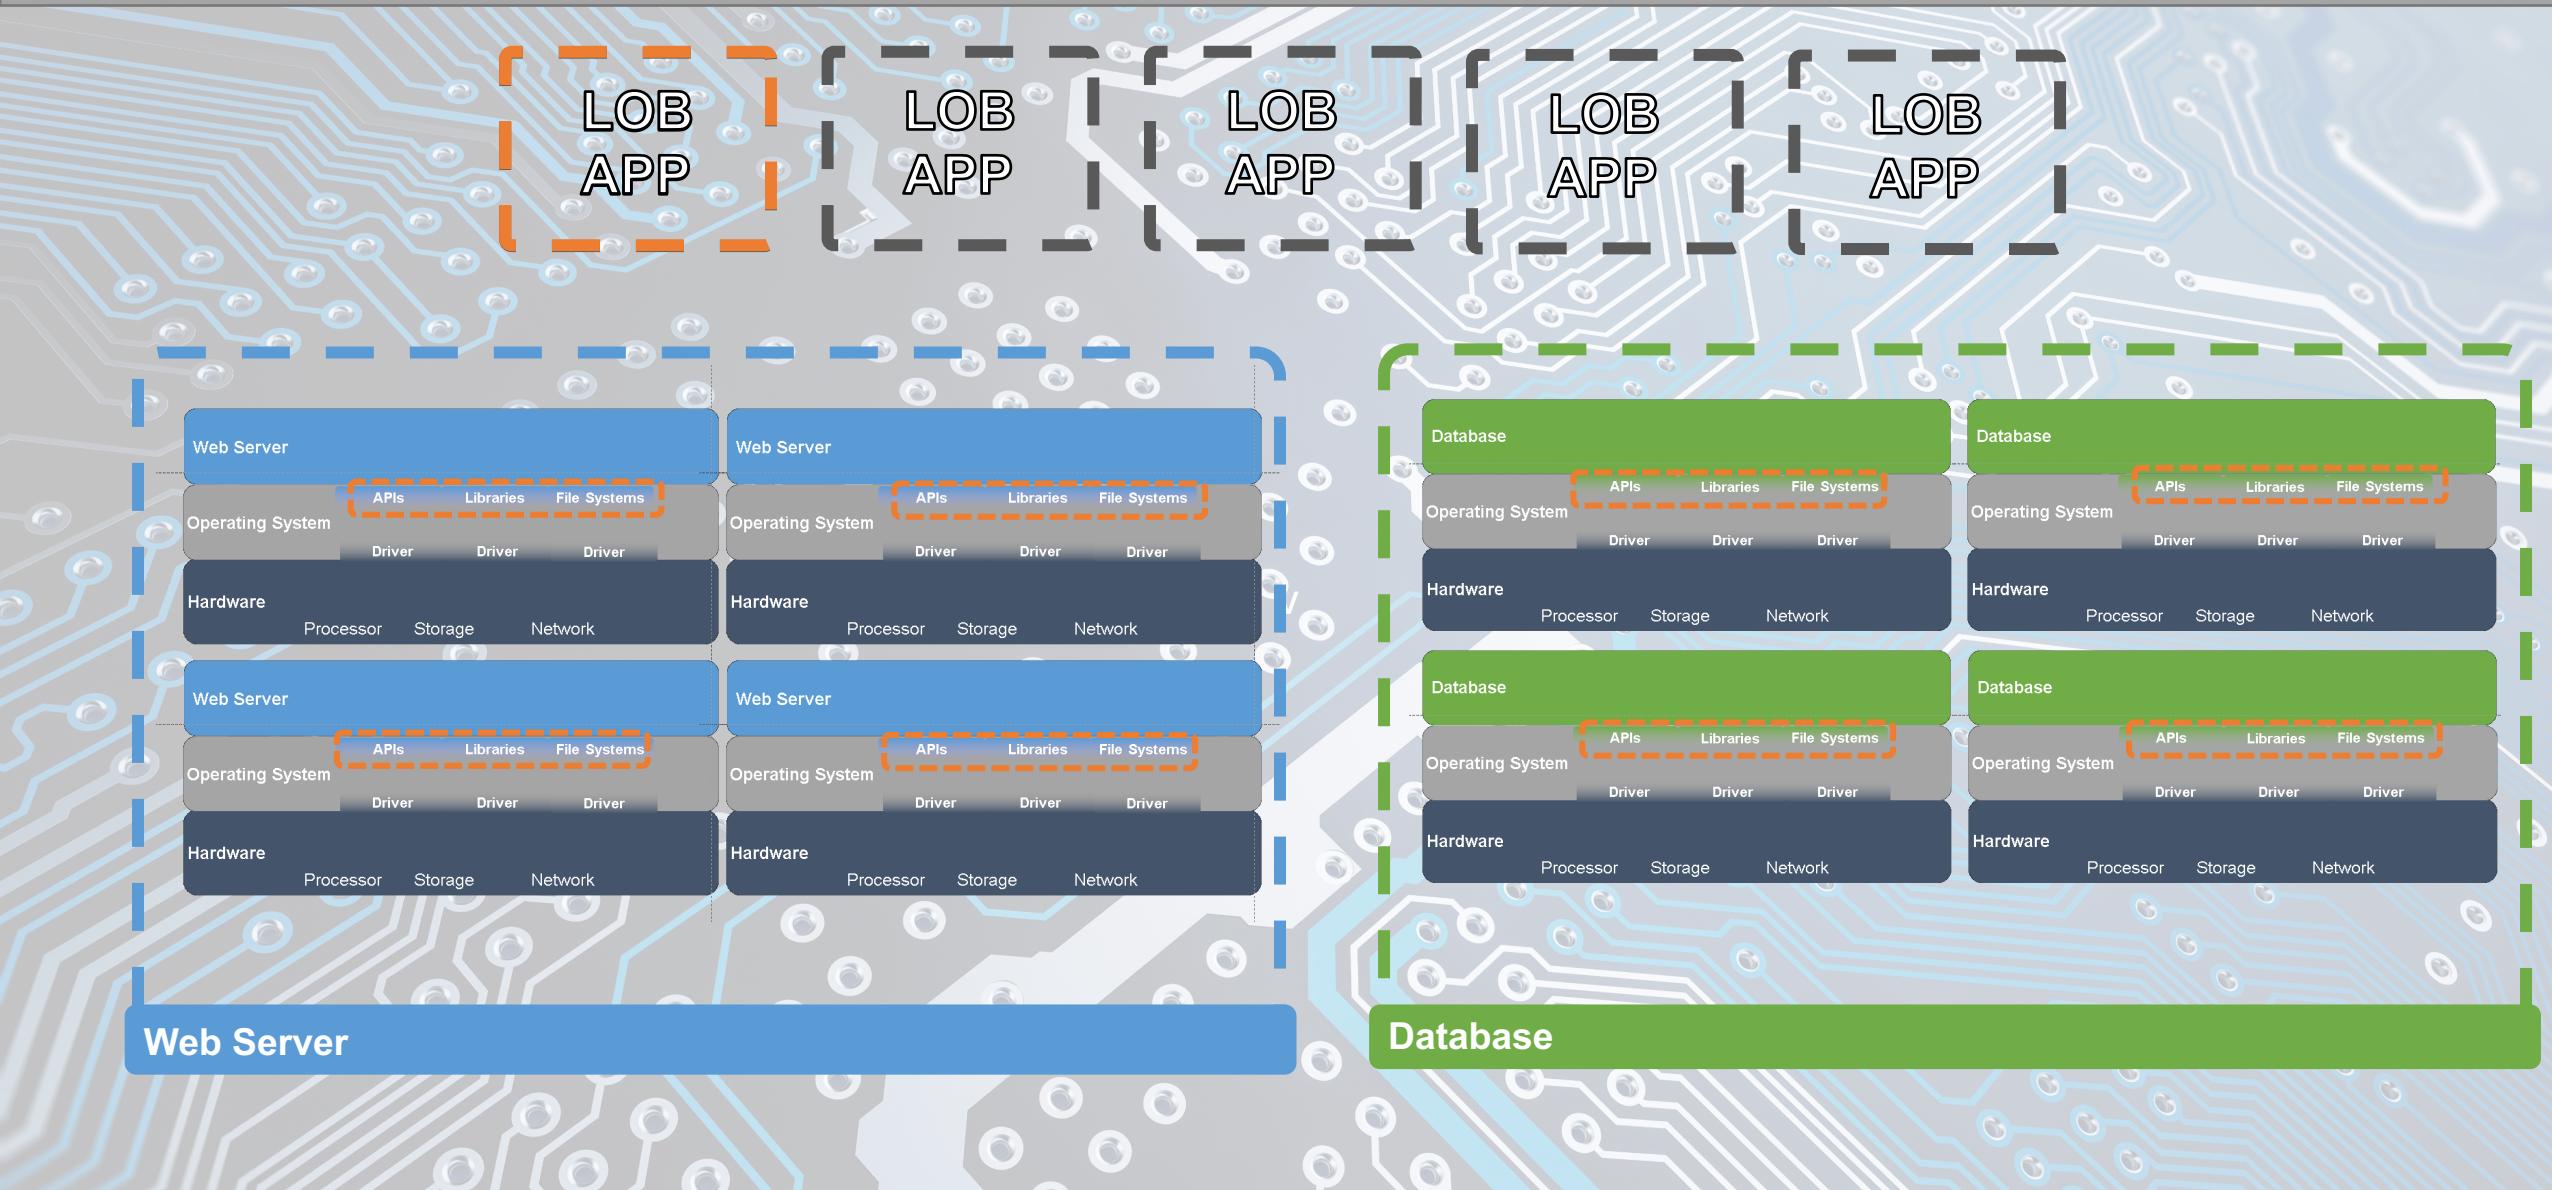
# LOB APP

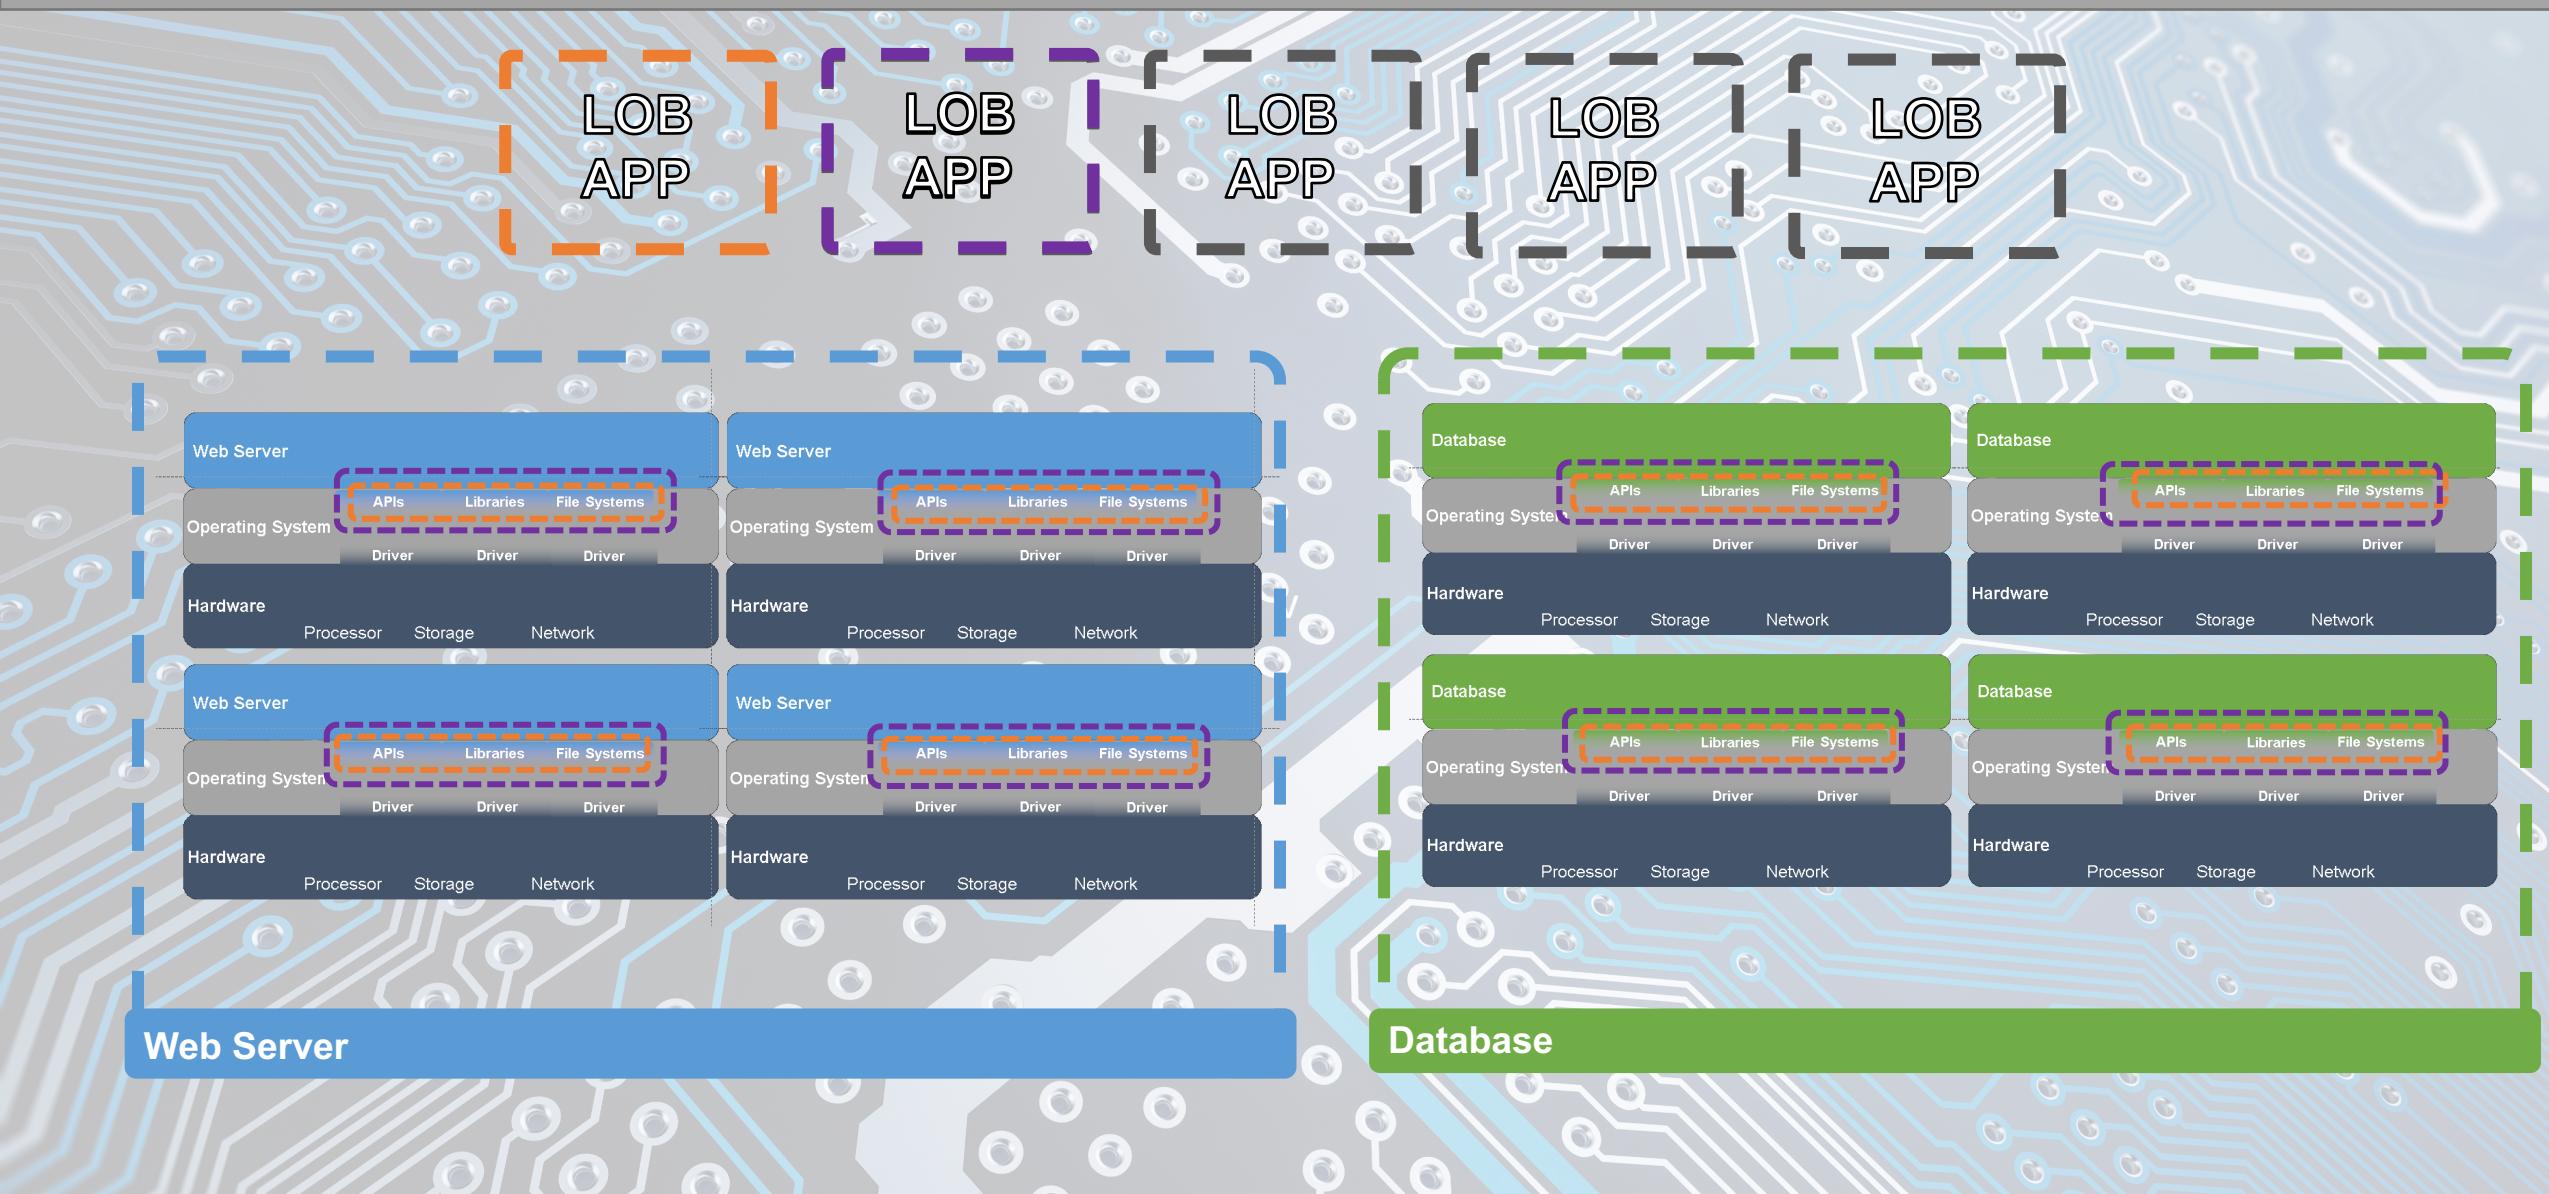


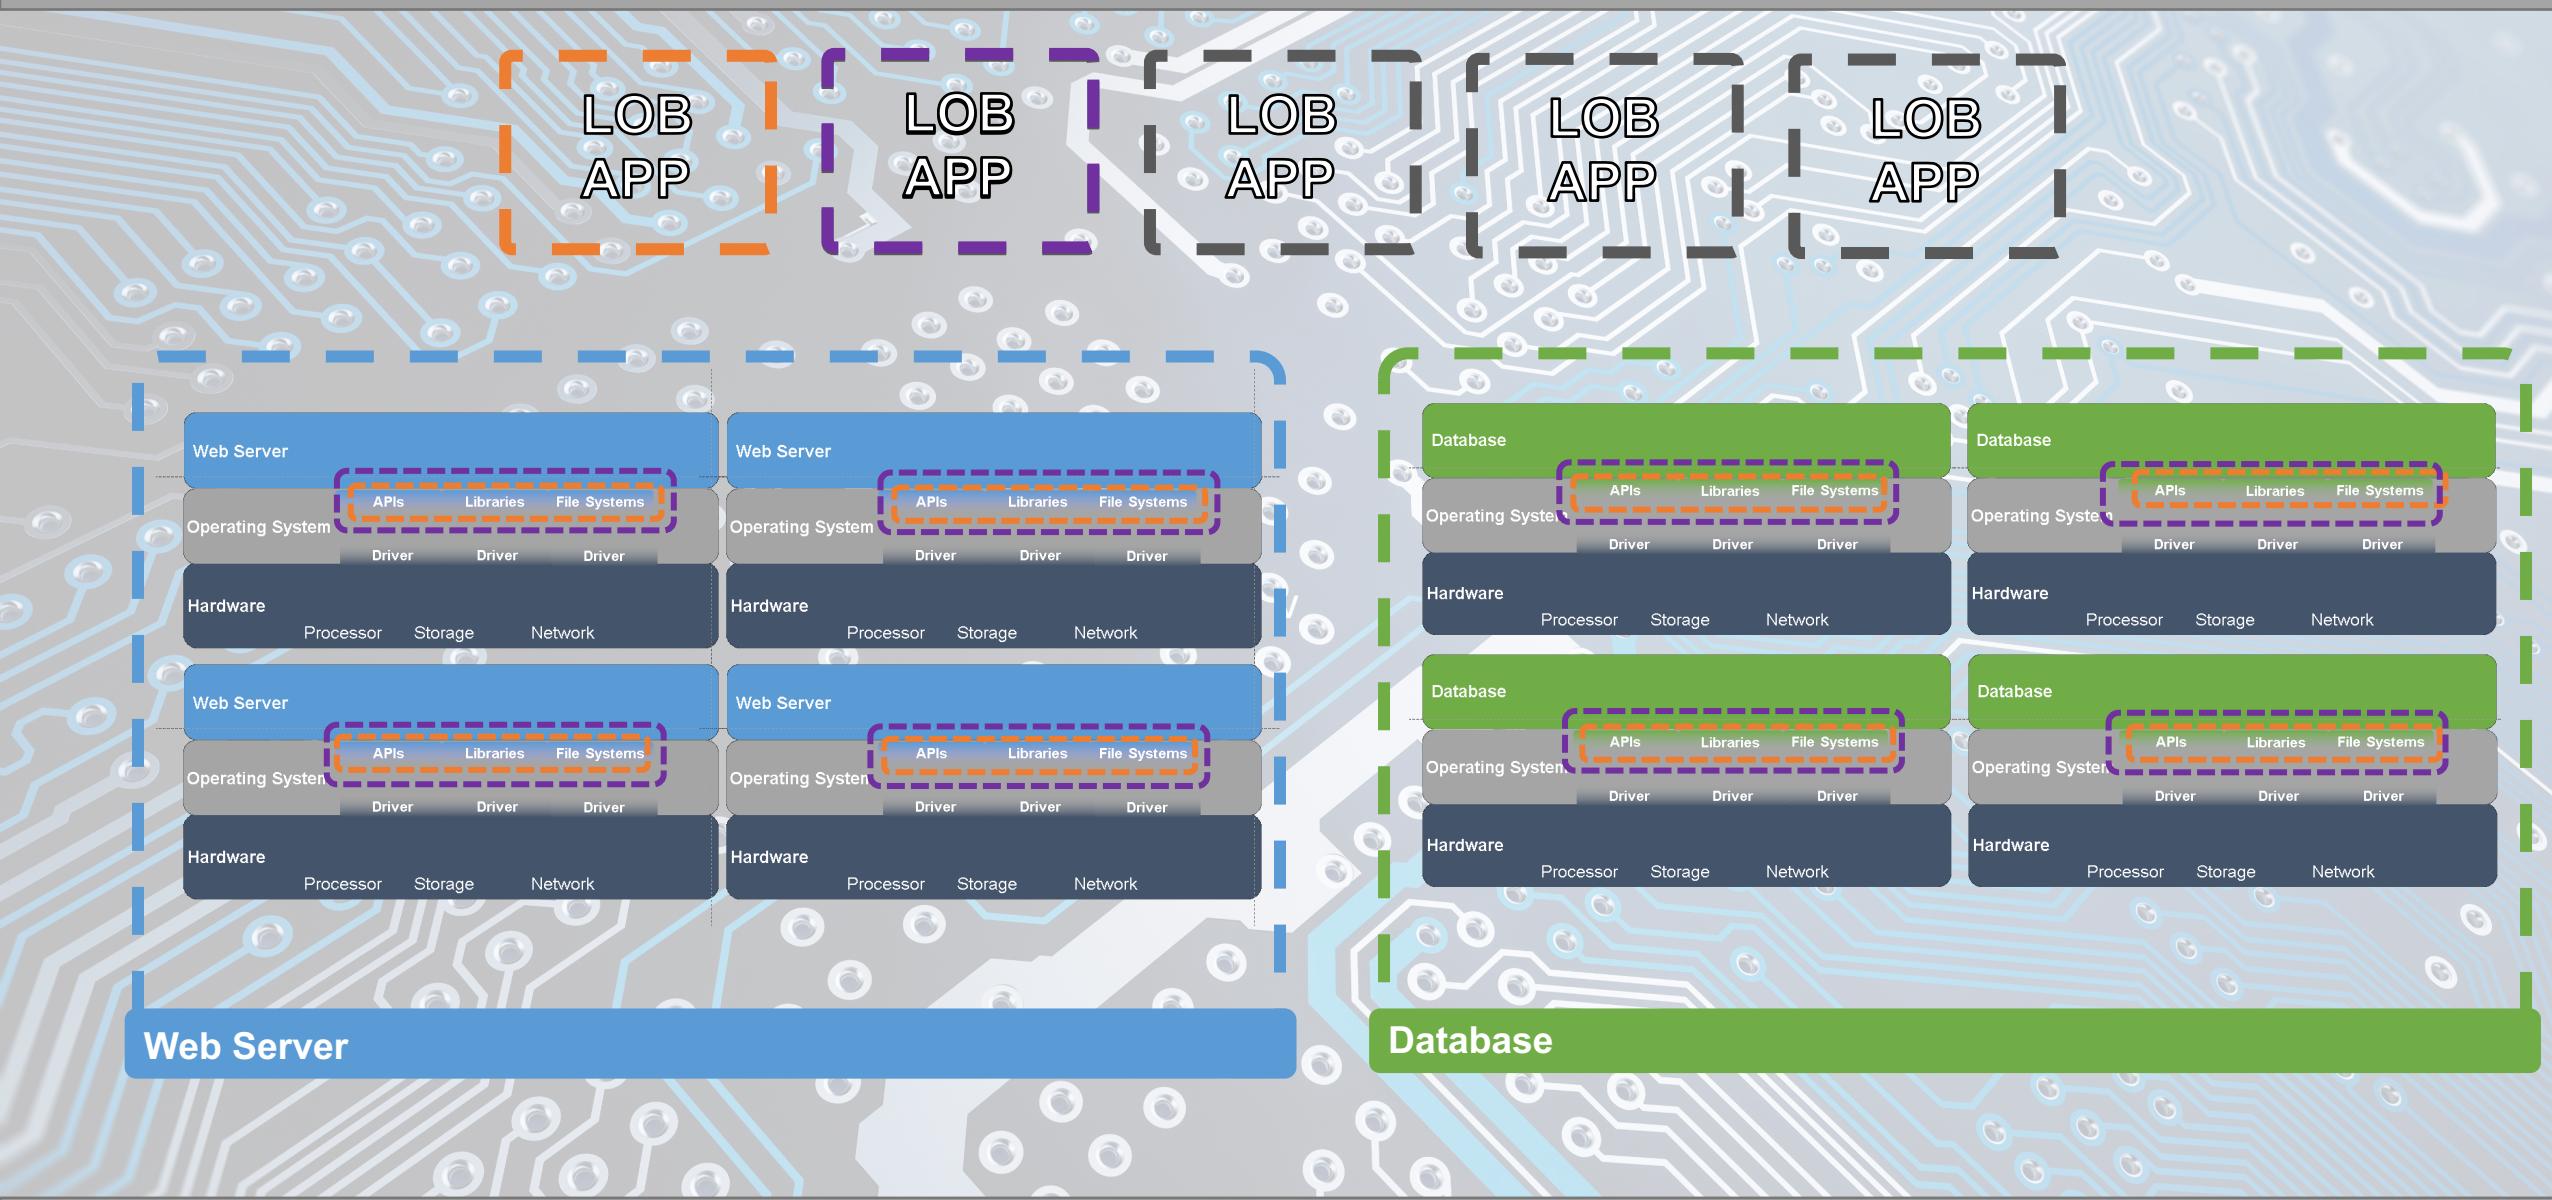
# LOB APP

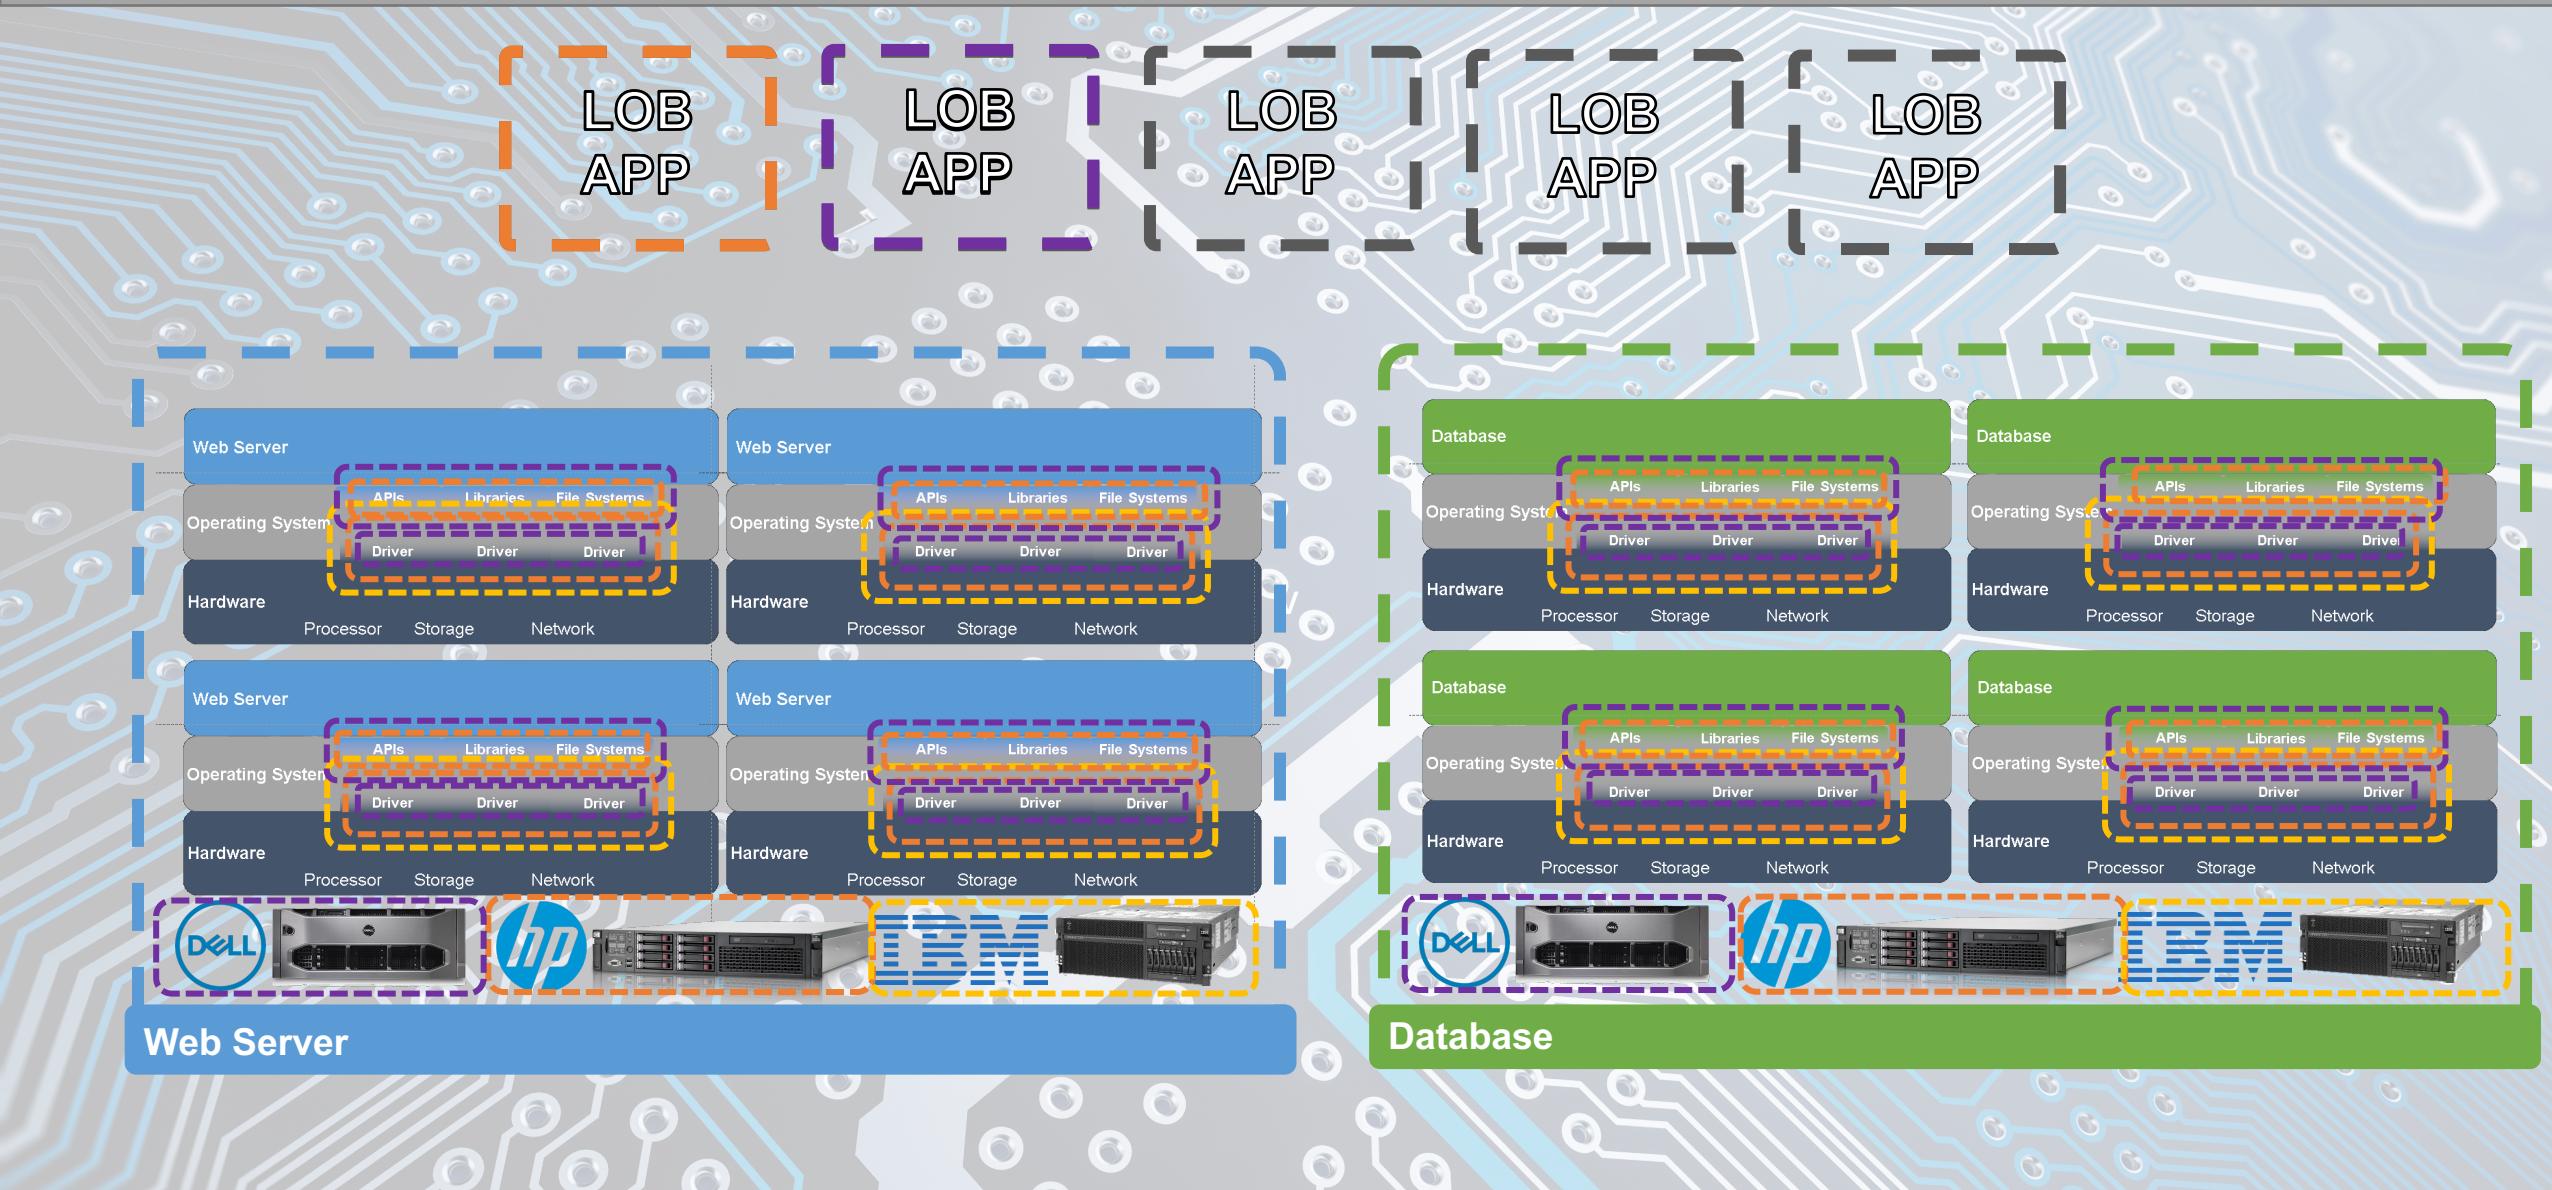














# Web Server



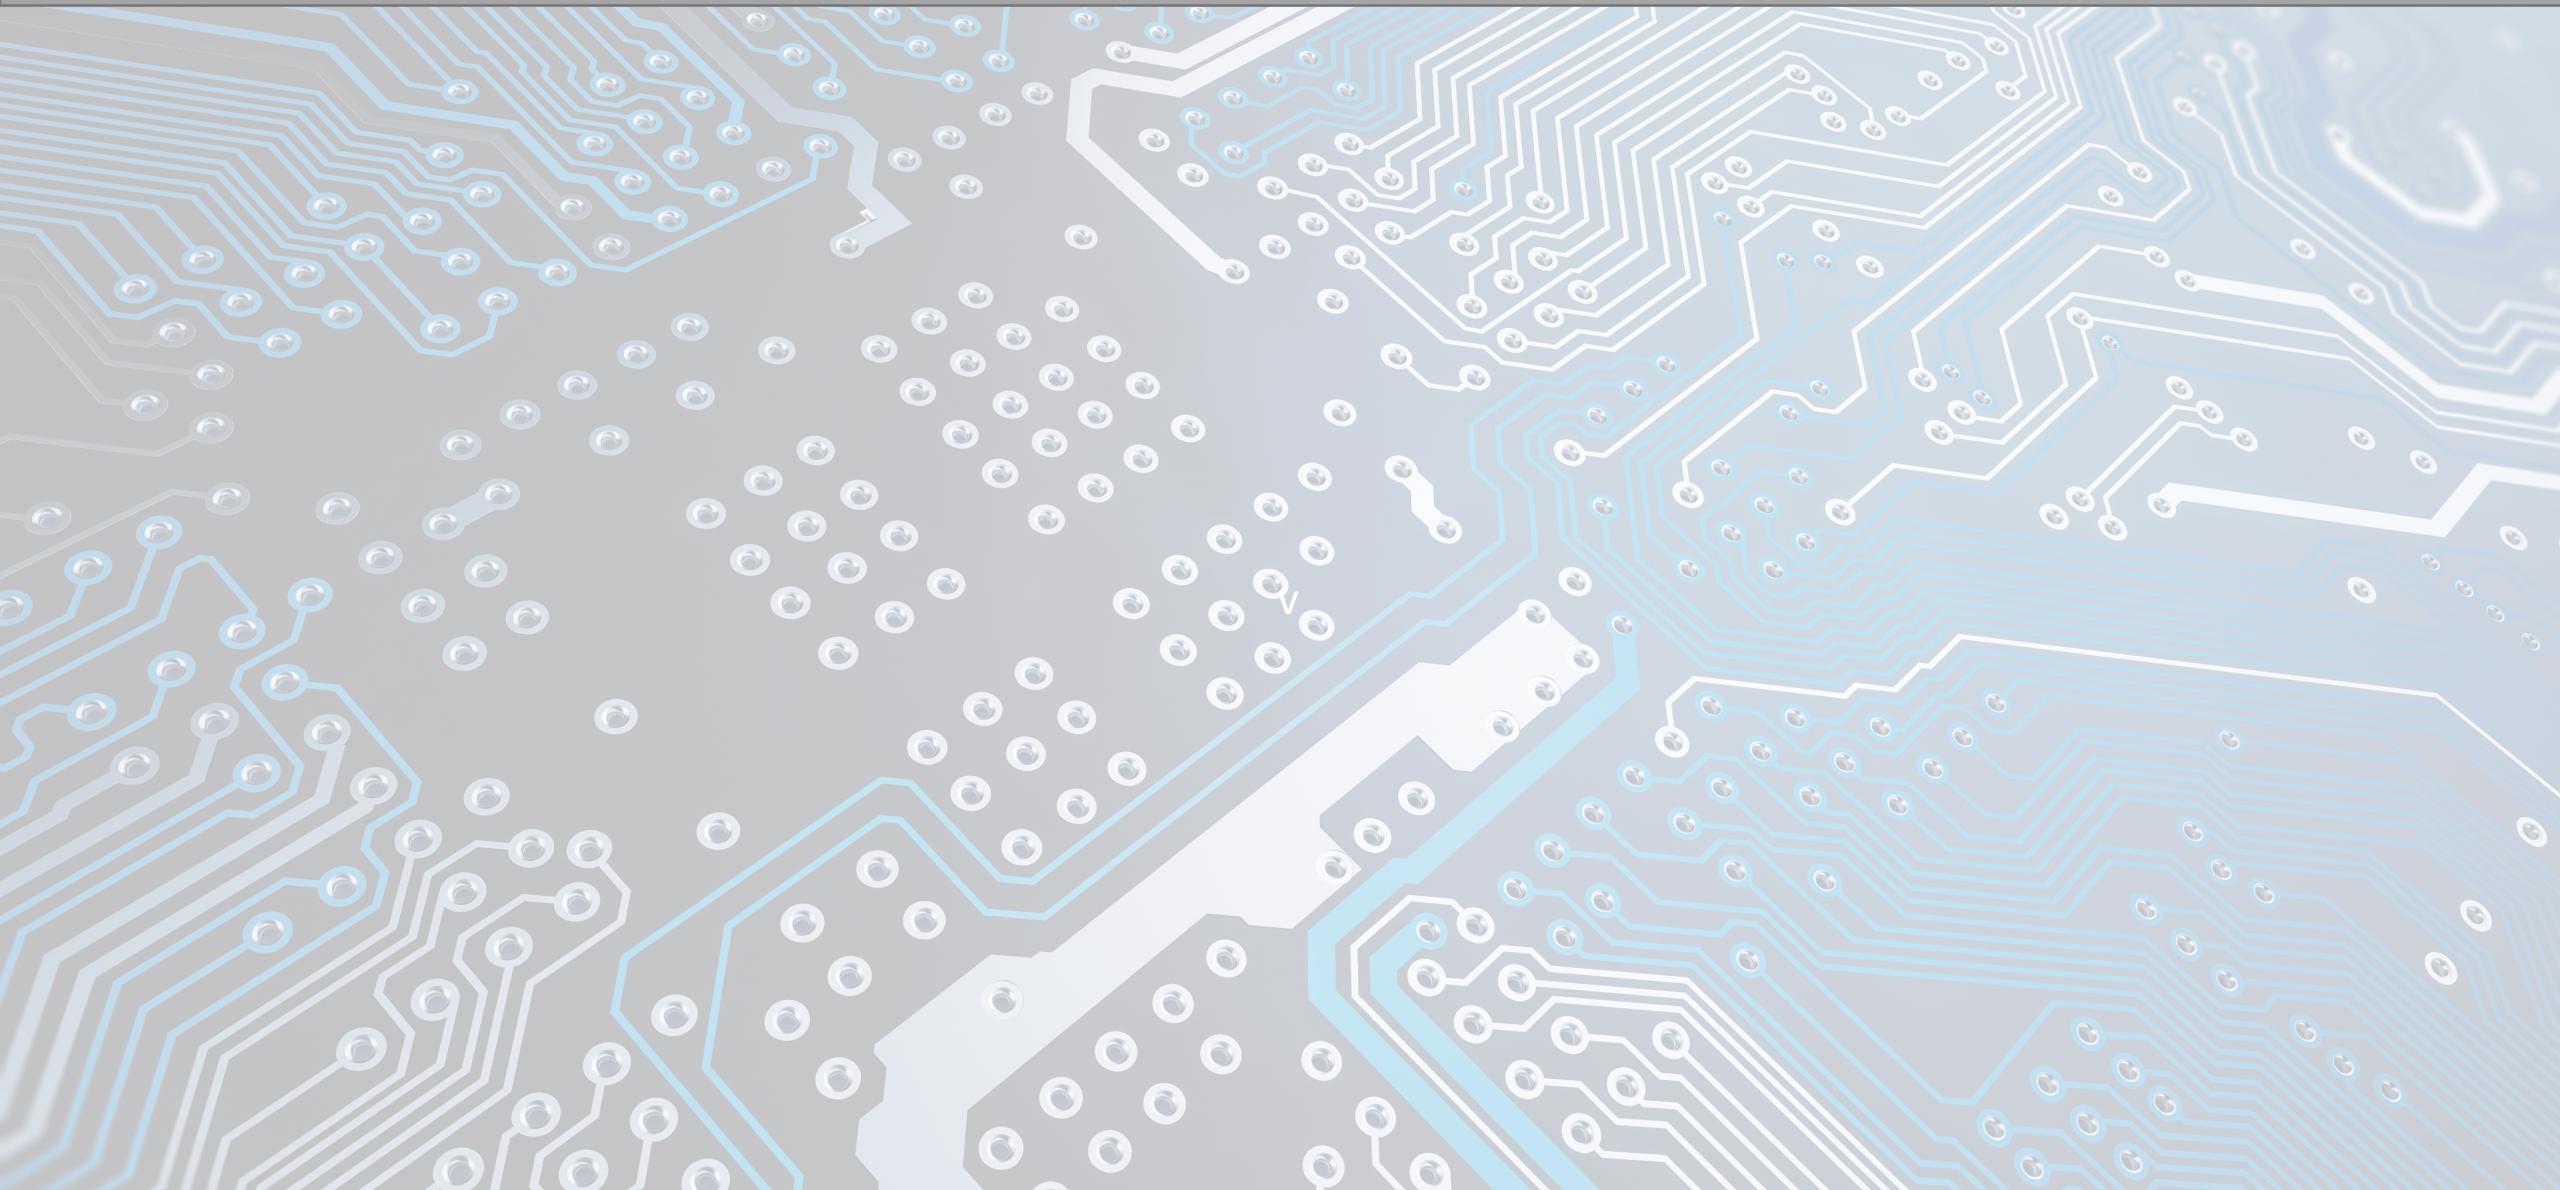
Hardware



rage Network



ork



## Operating System



## Hardware



Processor



Storage



Network



**Operating System**

Driver

Driver

Driver

**Virtualization Layer**

Driver

Driver

Driver

**Hardware**



**Operating System**

Driver

Driver

Driver

**Virtualization Layer**

Driver

Driver

Driver

**Hardware**



**Operating System**

Driver

Driver

Driver

**Virtualization Layer**

Driver

Driver

Driver

**Hardware**

IBM





Operating System

Driver

Driver

Driver

Hypervisor

Driver

Driver

Driver

Hardware

**Web Server**

**Database**

**Operating System**

**Hypervisor**

**Hardware**



**Web Server**

**Database**

**Operating System**

**Operating System**

**Hypervisor**

**Hardware**



## Line of Business Application

Web Server

Database

Operating System



Operating System



Hypervisor

Hardware

## Line of Business Application

Web Server

Database

APIs   Libraries   File Systems

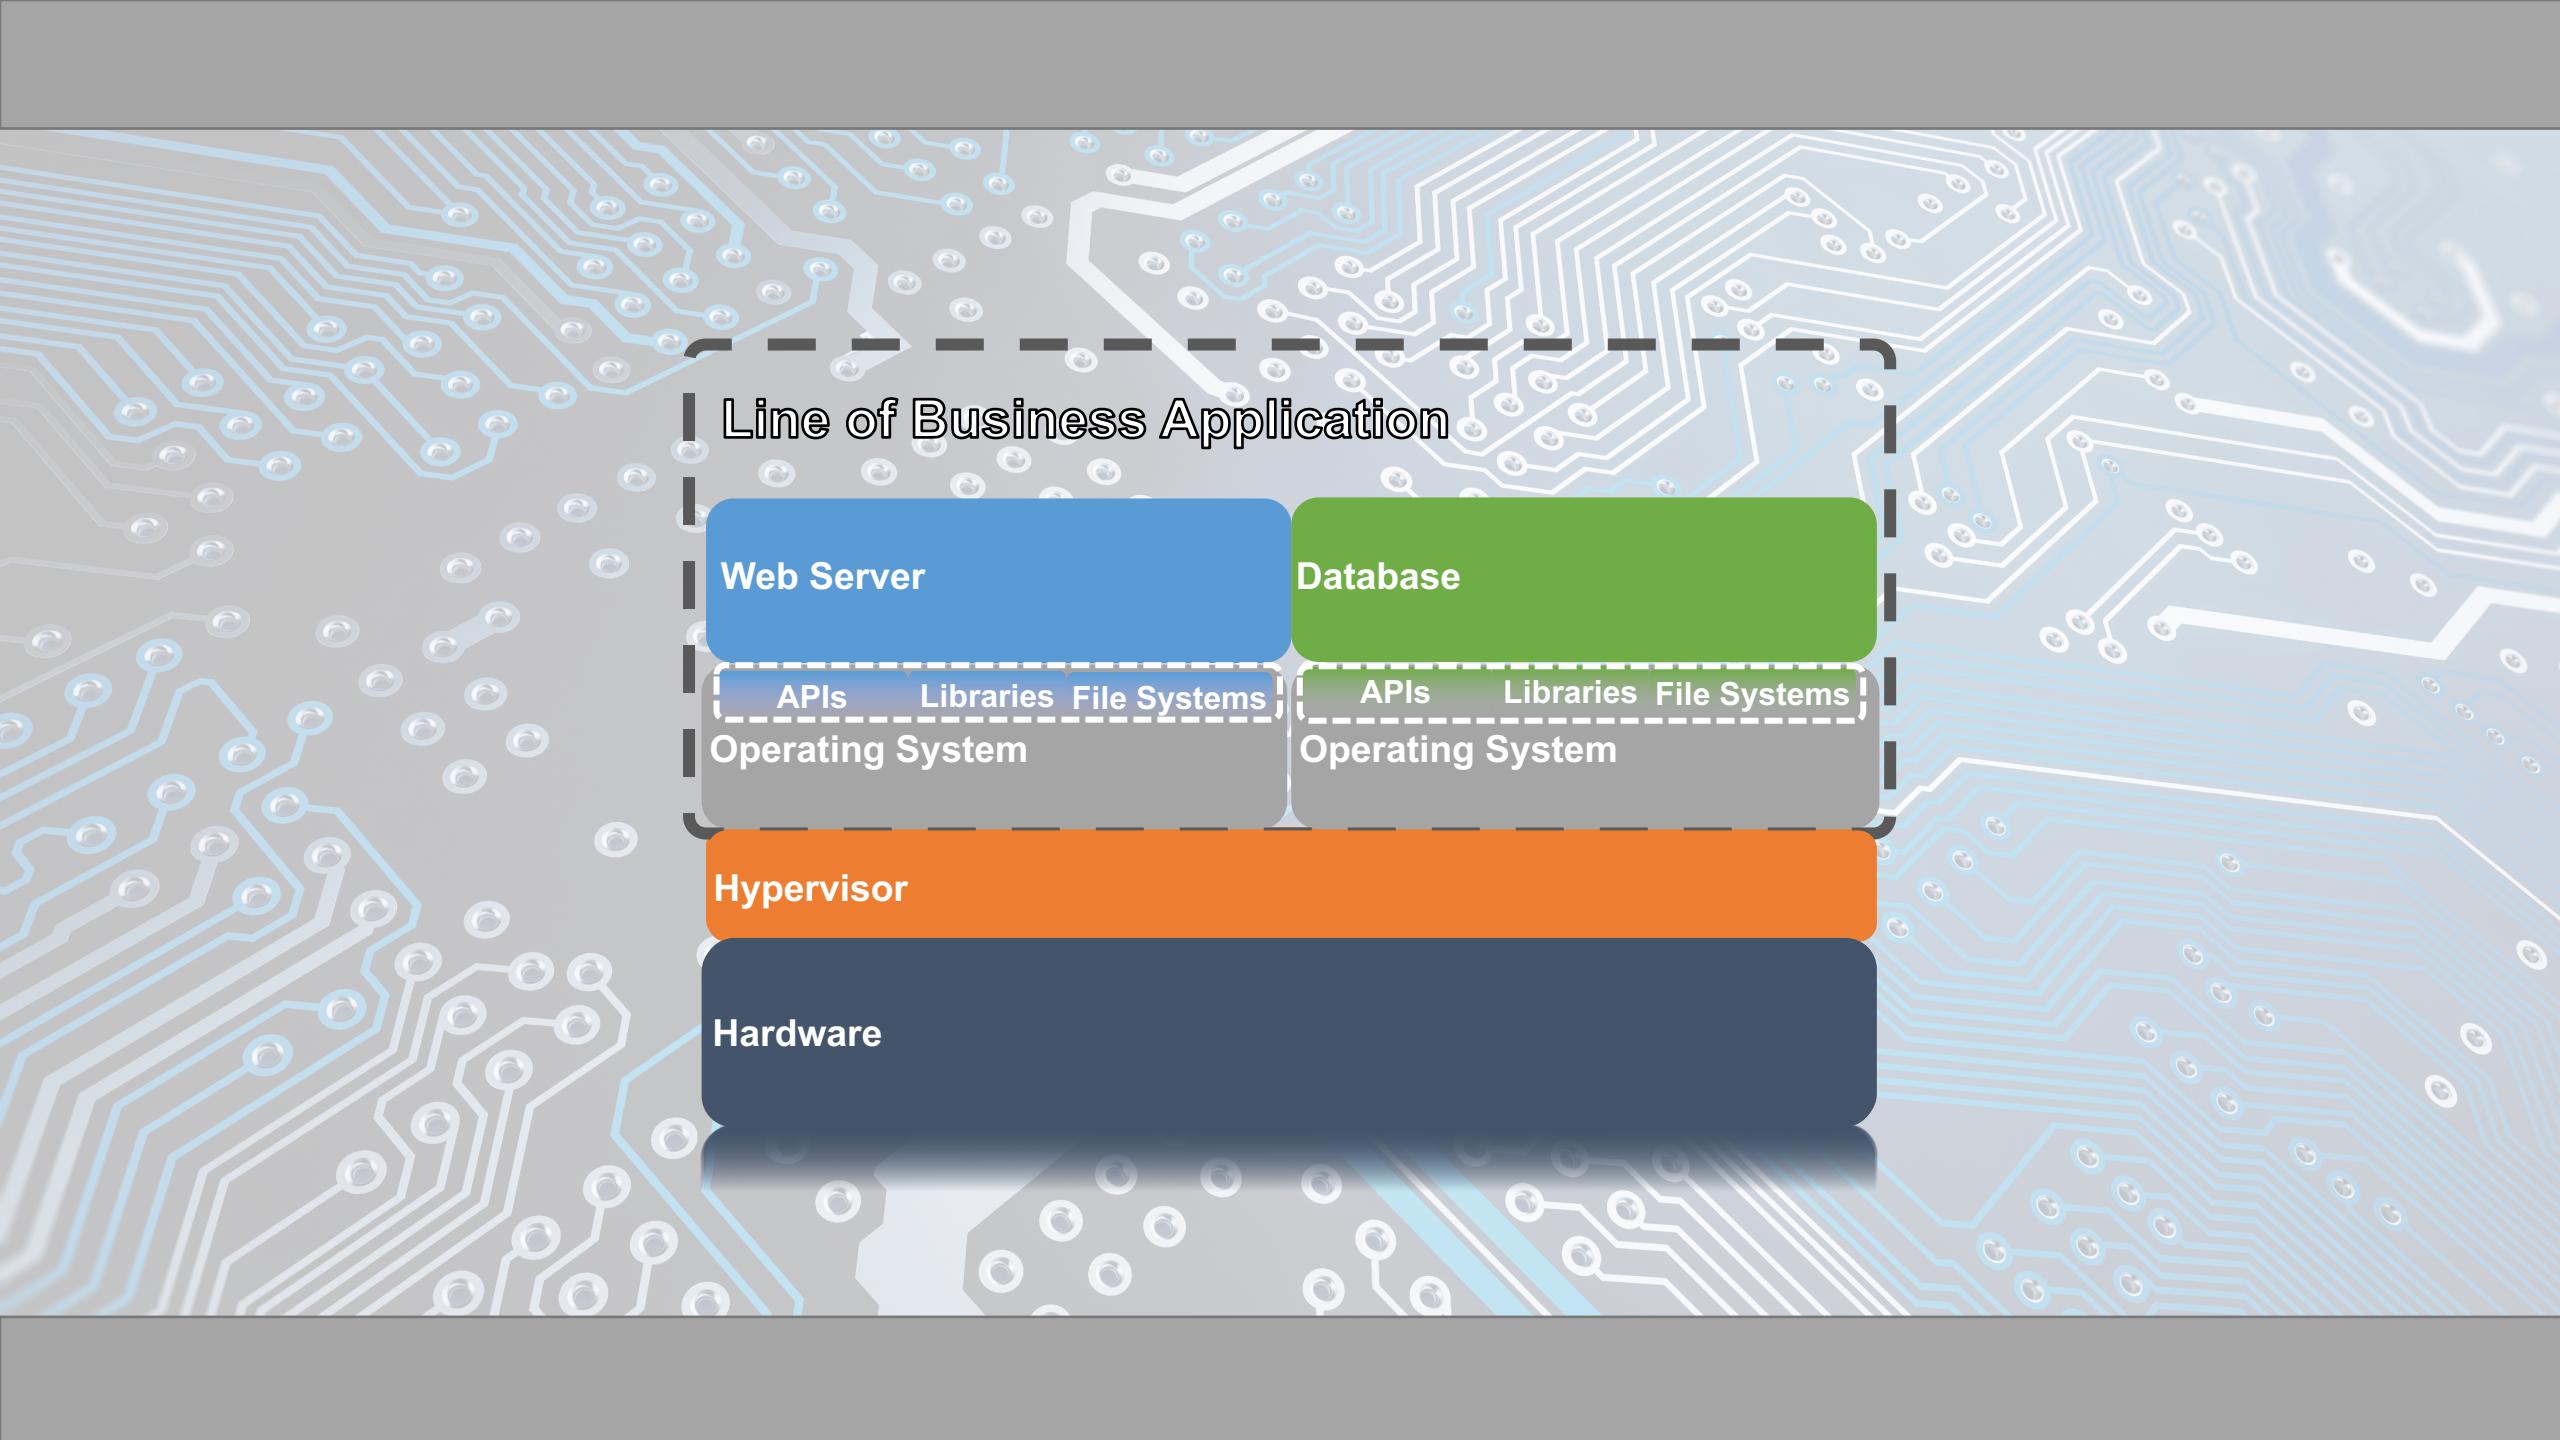
Operating System

APIs   Libraries   File Systems

Operating System

Hypervisor

Hardware



## Line of Business Application

Web Server

Database

Operating System

Operating System

Hypervisor

Hypervisor

Hardware

Hardware

## Application Development

Web Server

Database

Operating System

Operating System

Hypervisor

Hypervisor

Hardware

Hardware

Infrastructure Operations

## Application Development

Web Server

Database

Operating System

Operating System

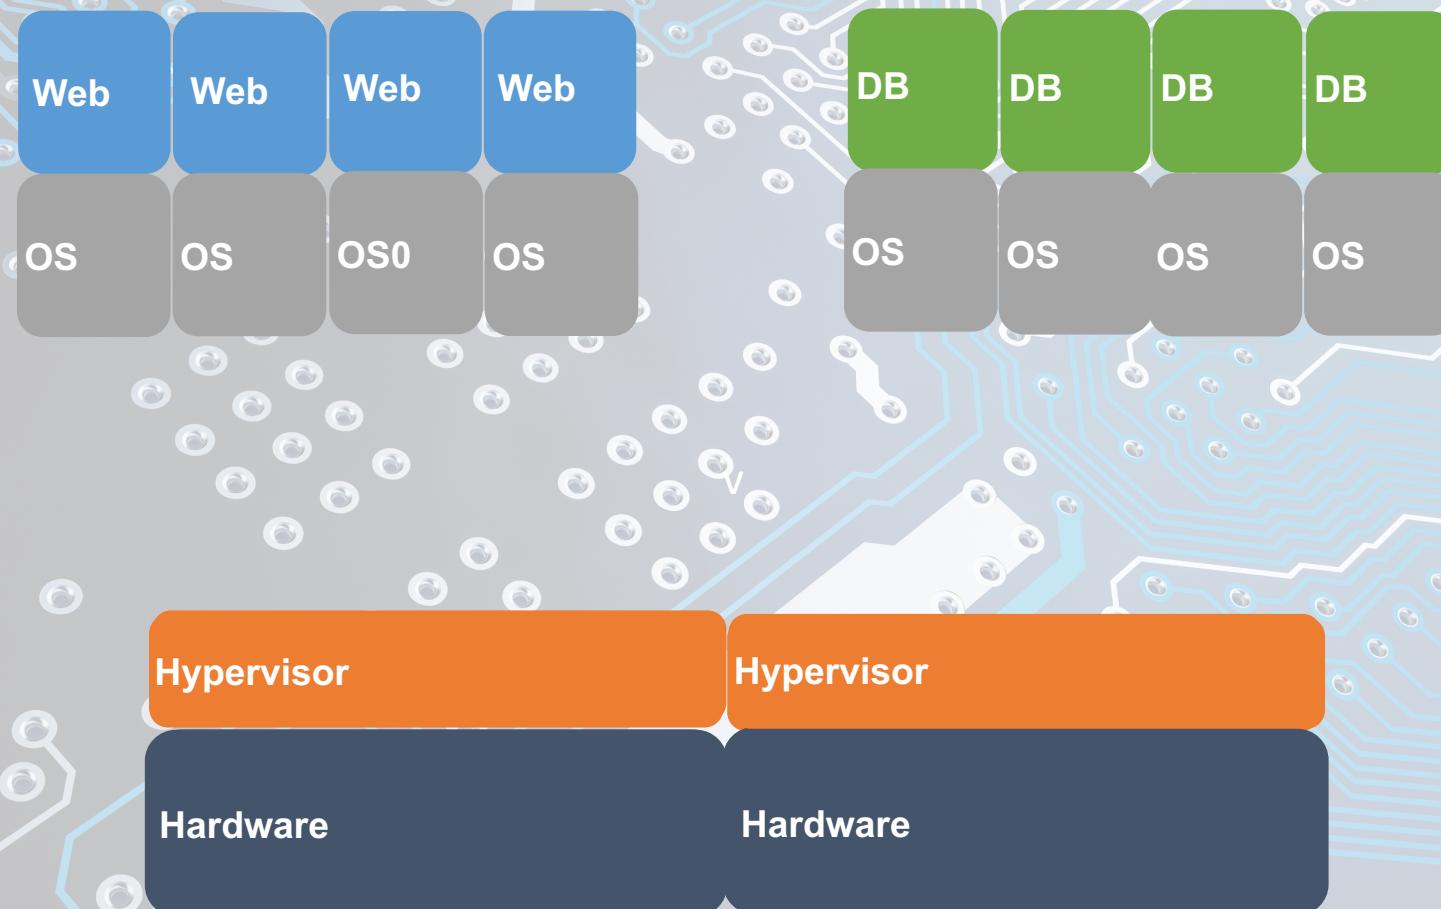
Hypervisor

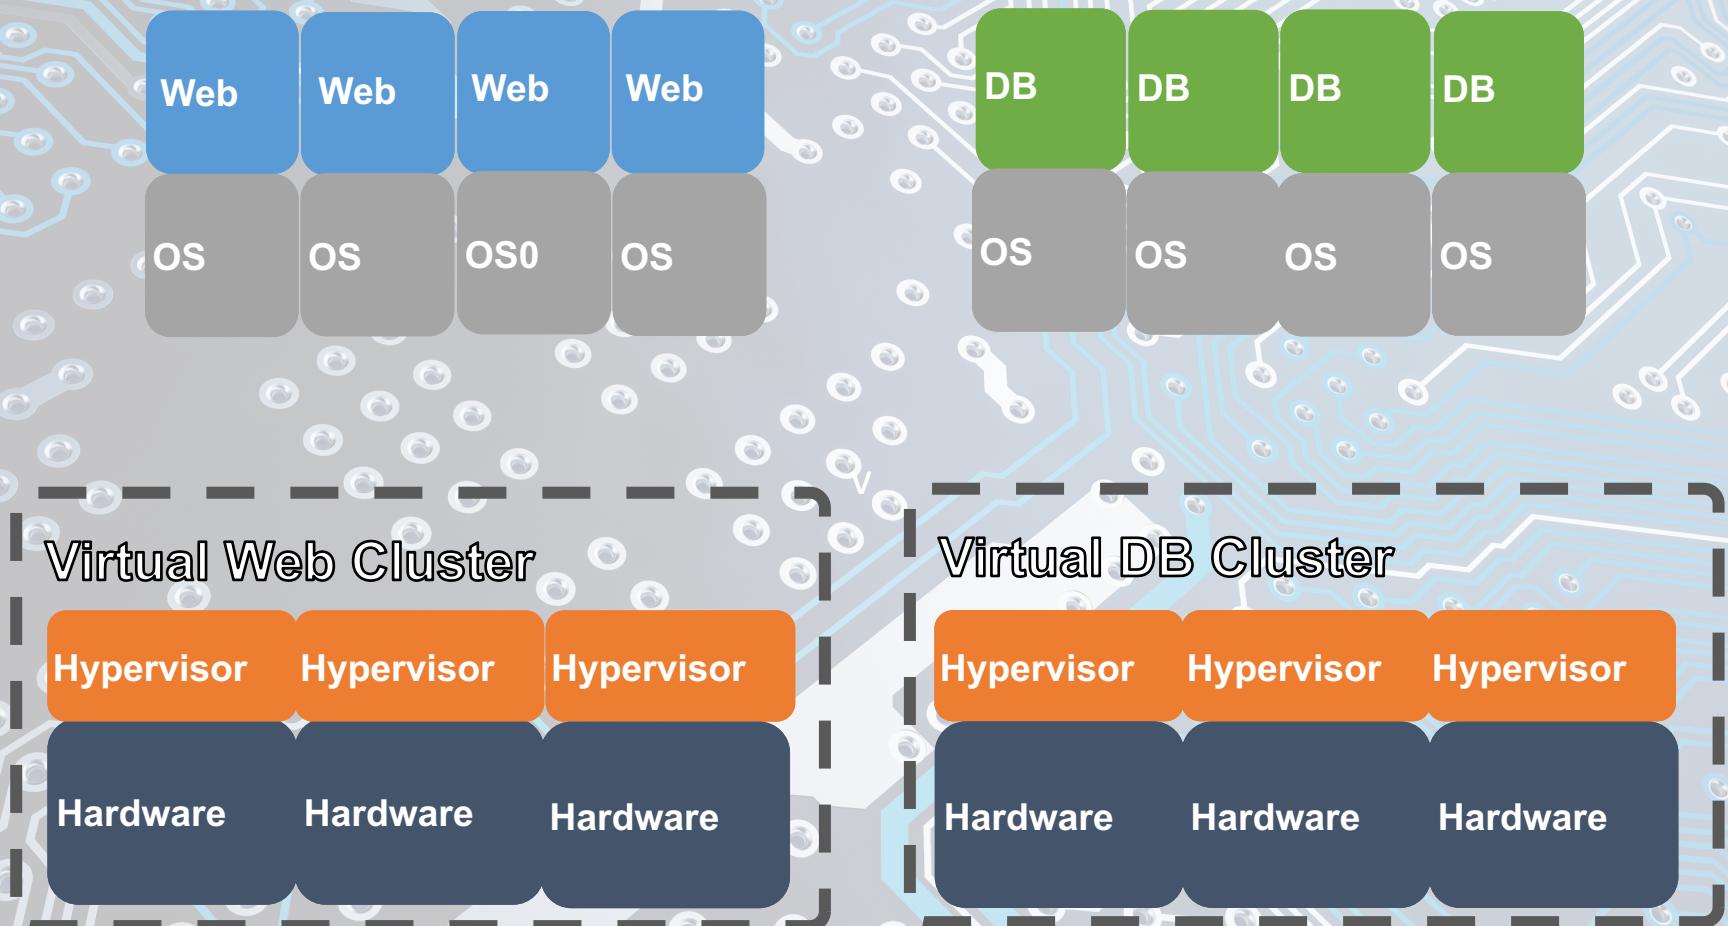
Hypervisor

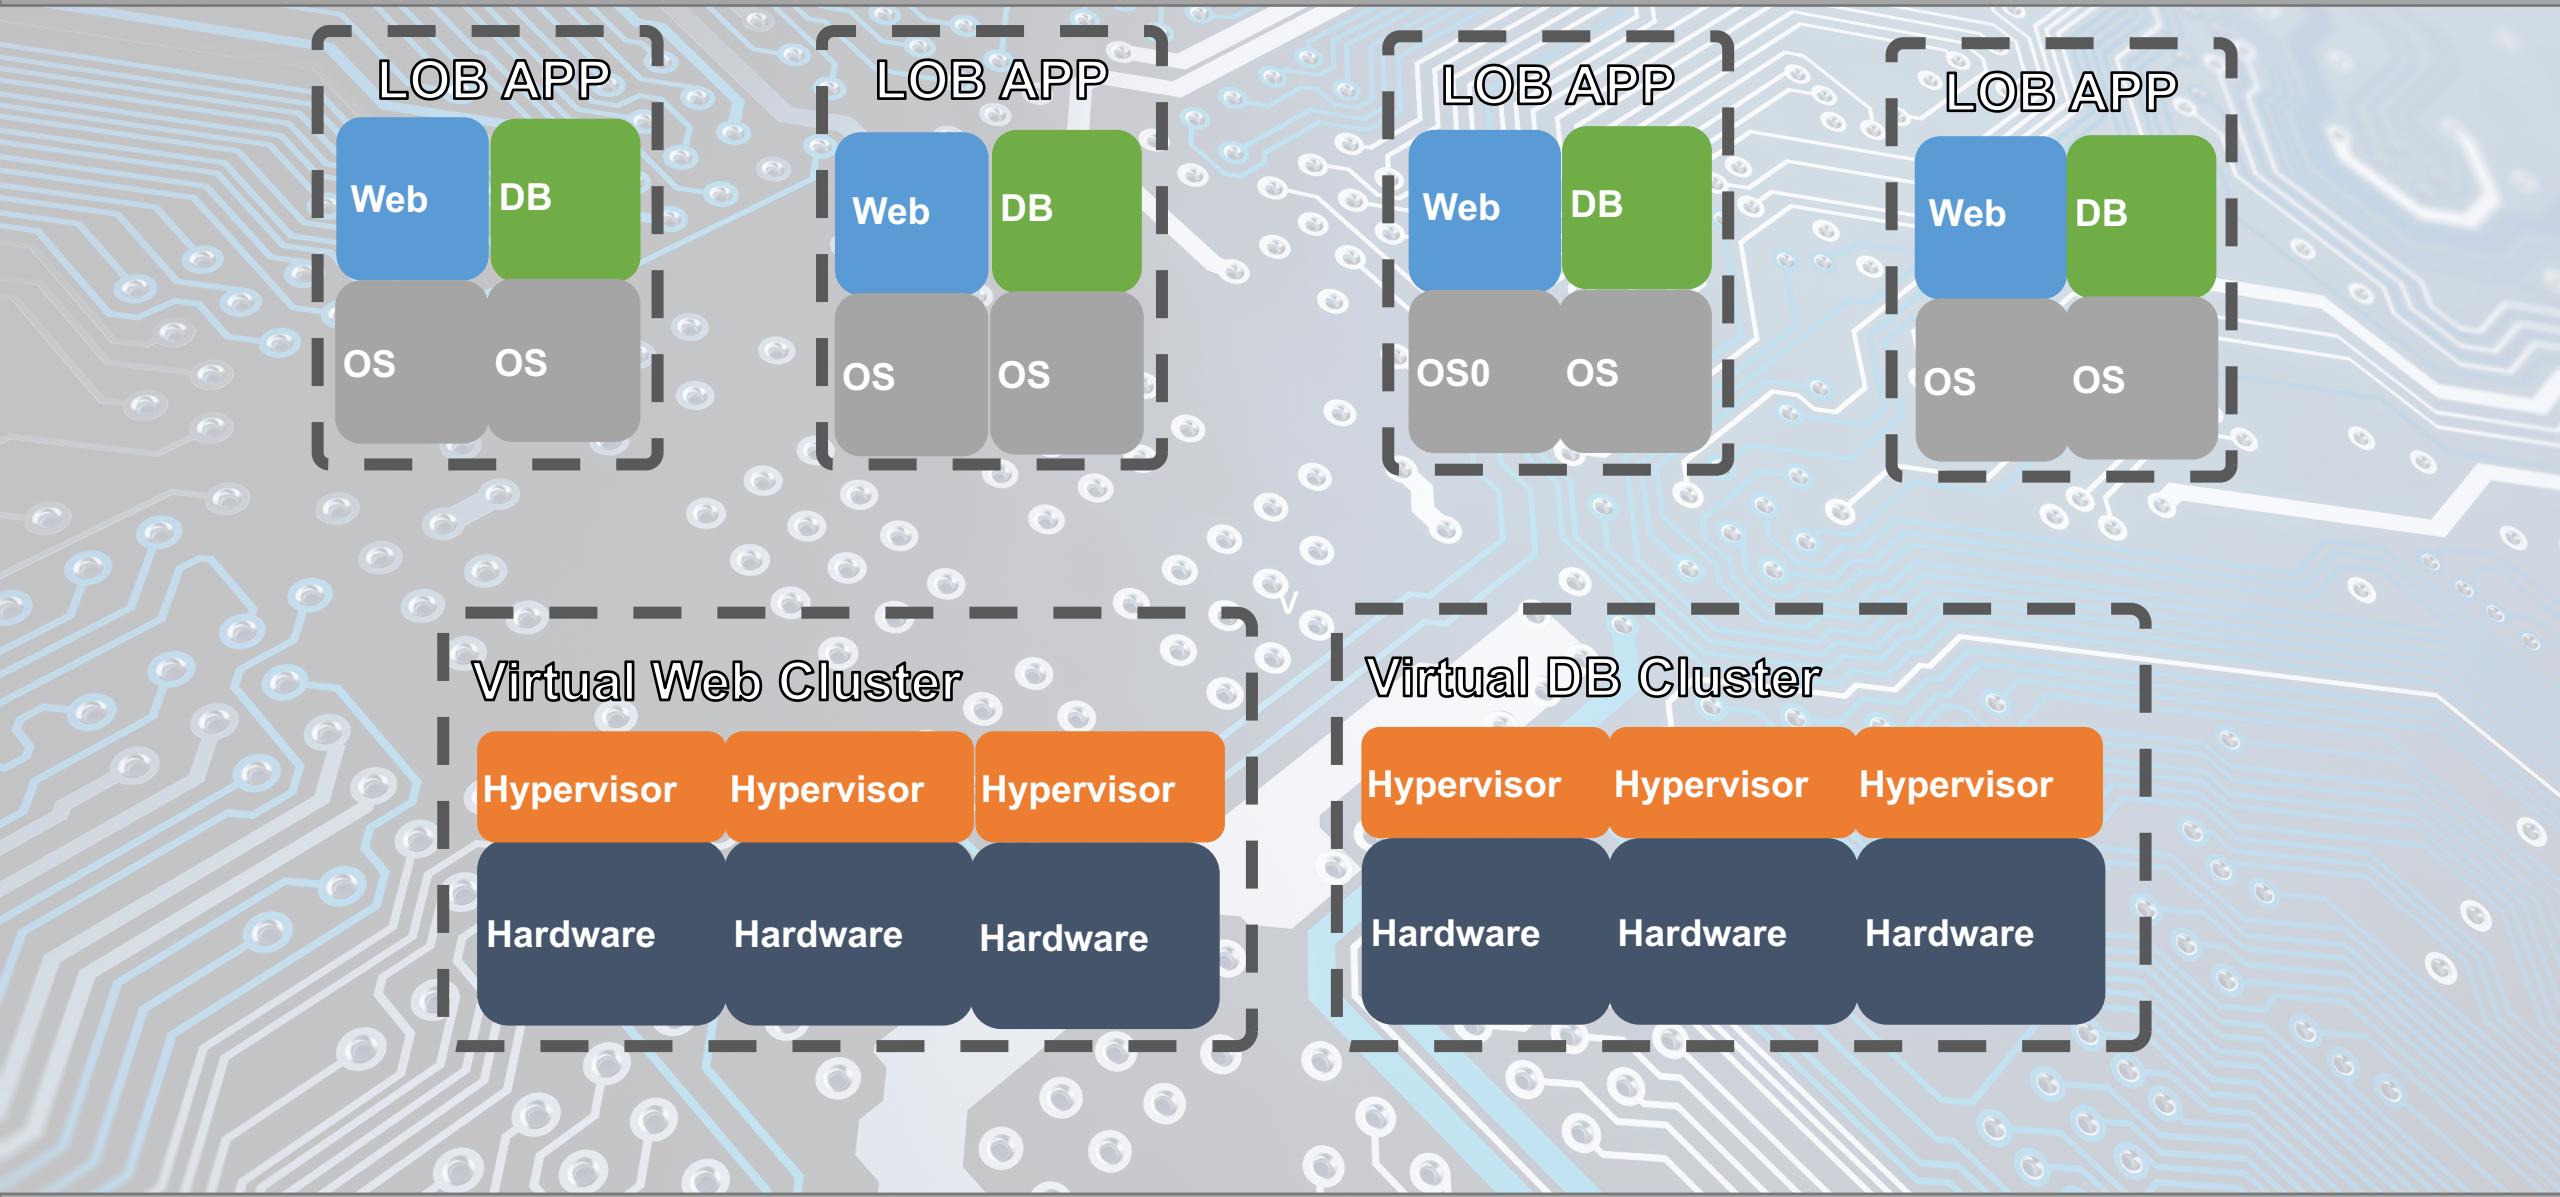
Hardware

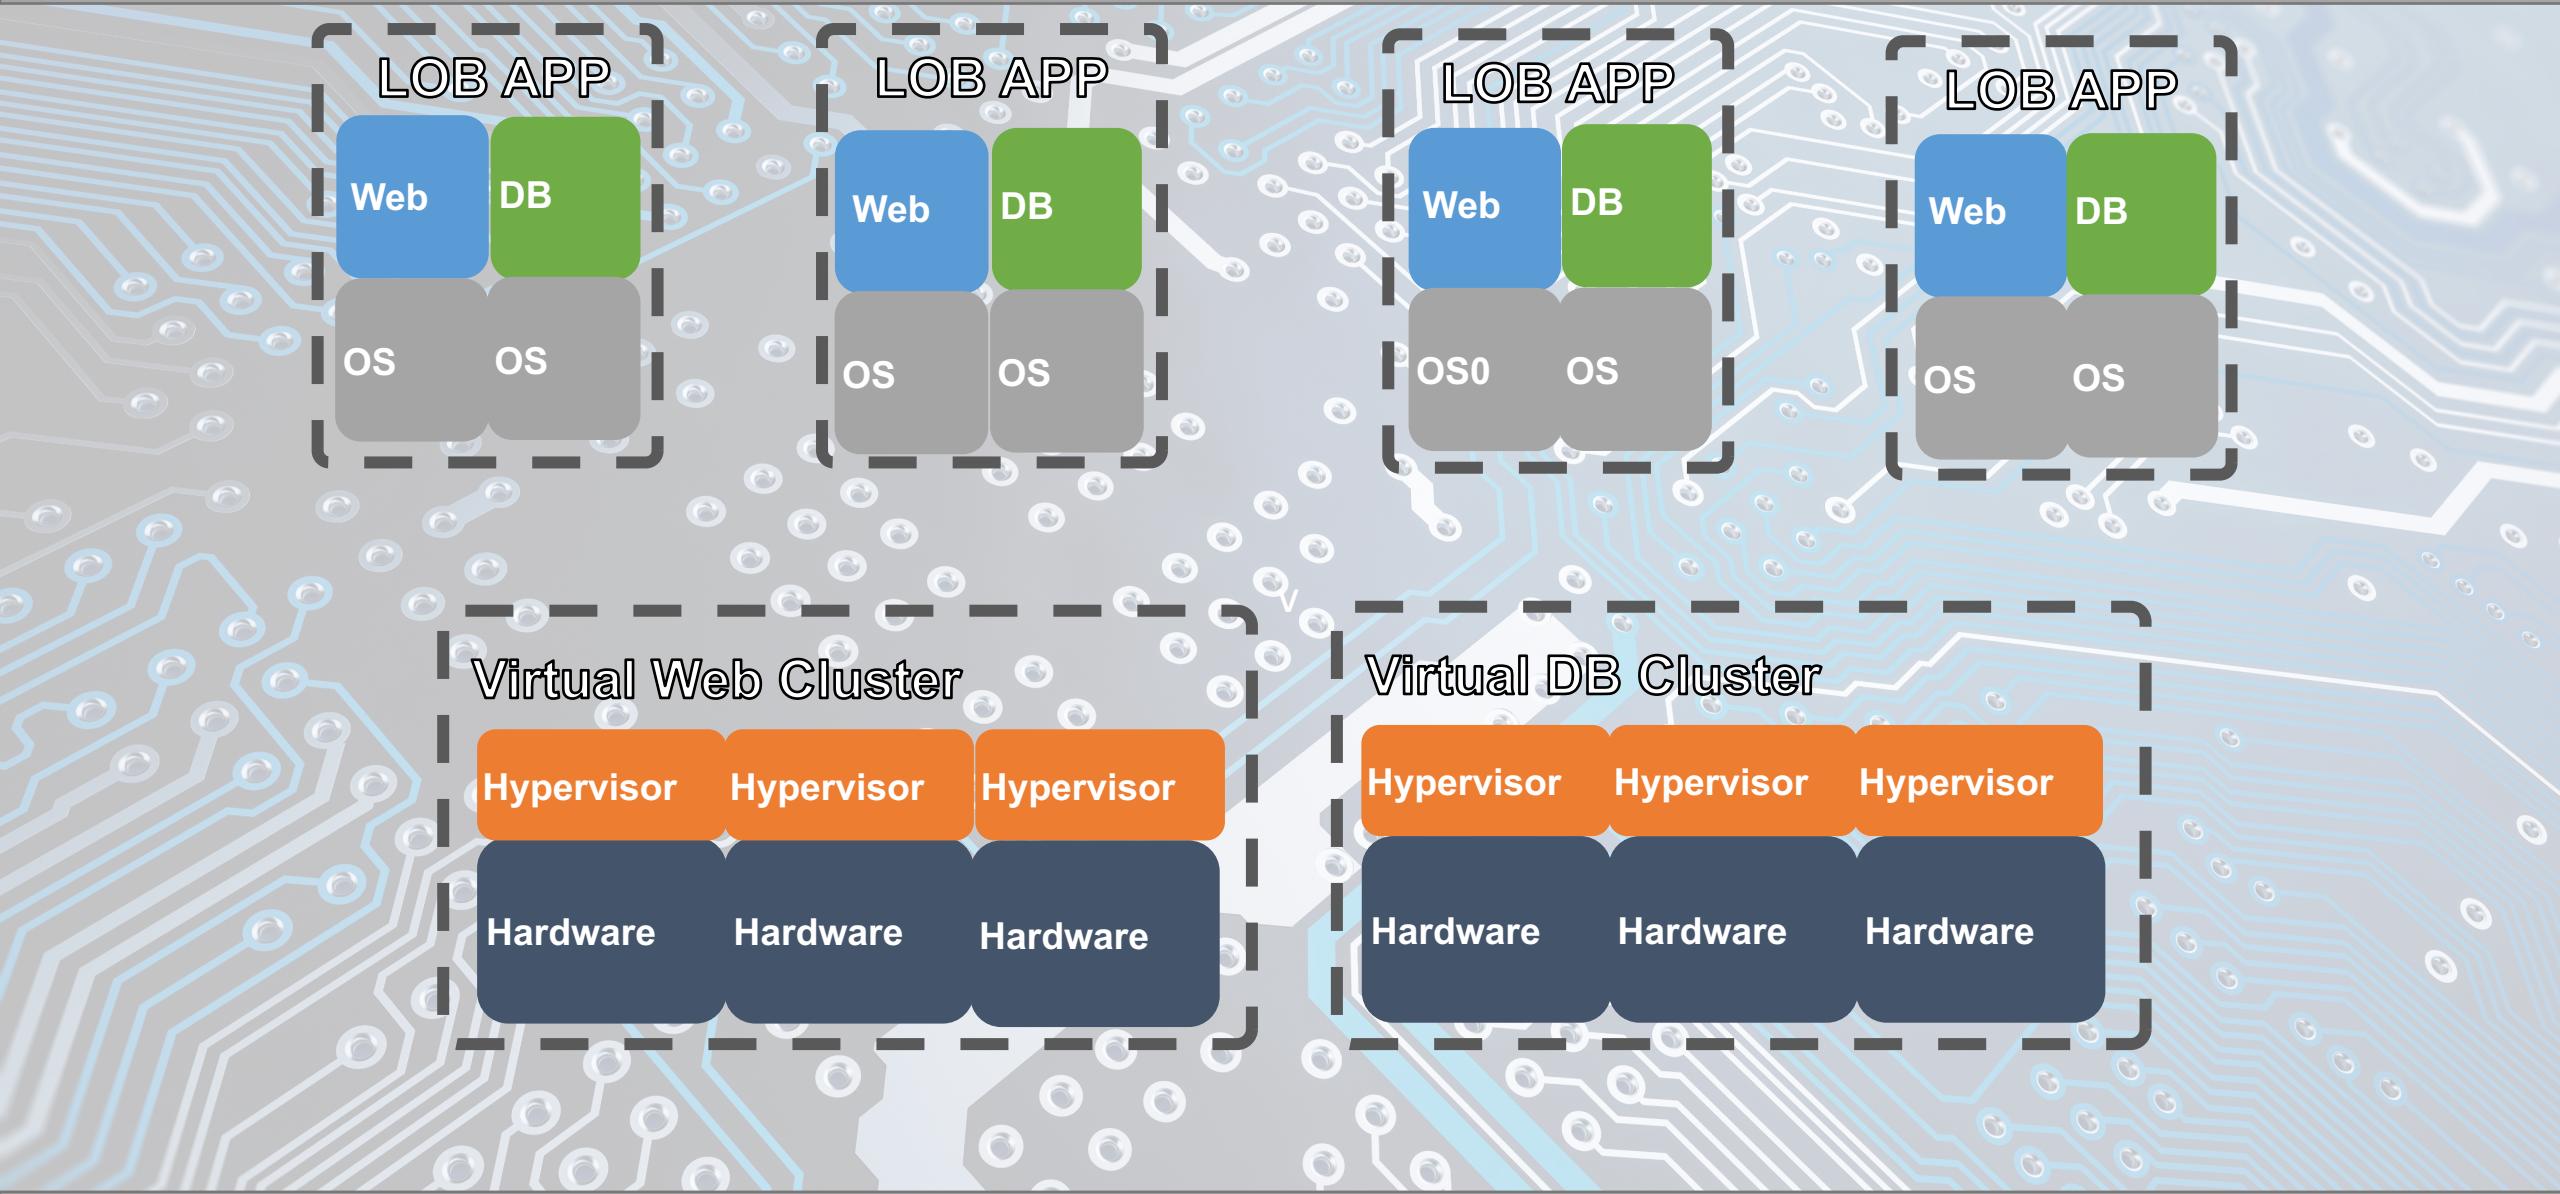
Hardware

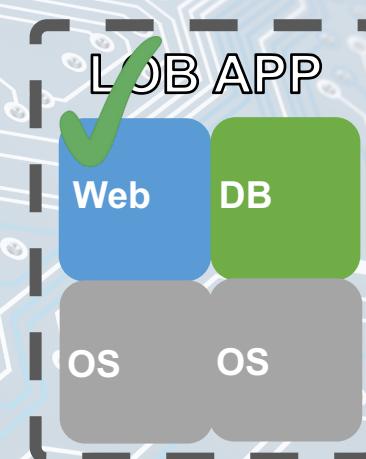
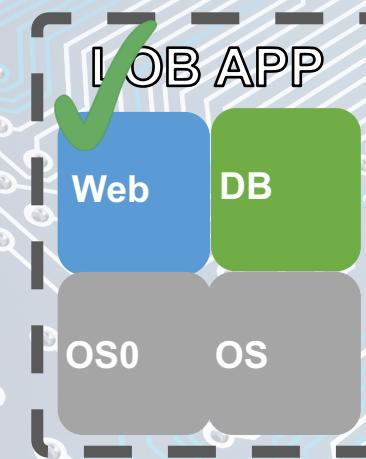
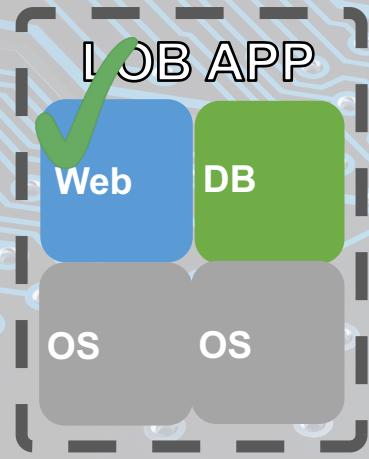
## Infrastructure Operations











Virtual Web Cluster



Virtual DB Cluster



## Platforms

Operating System

APIs

Libraries

File Systems

Hypervisor

Driver

Driver

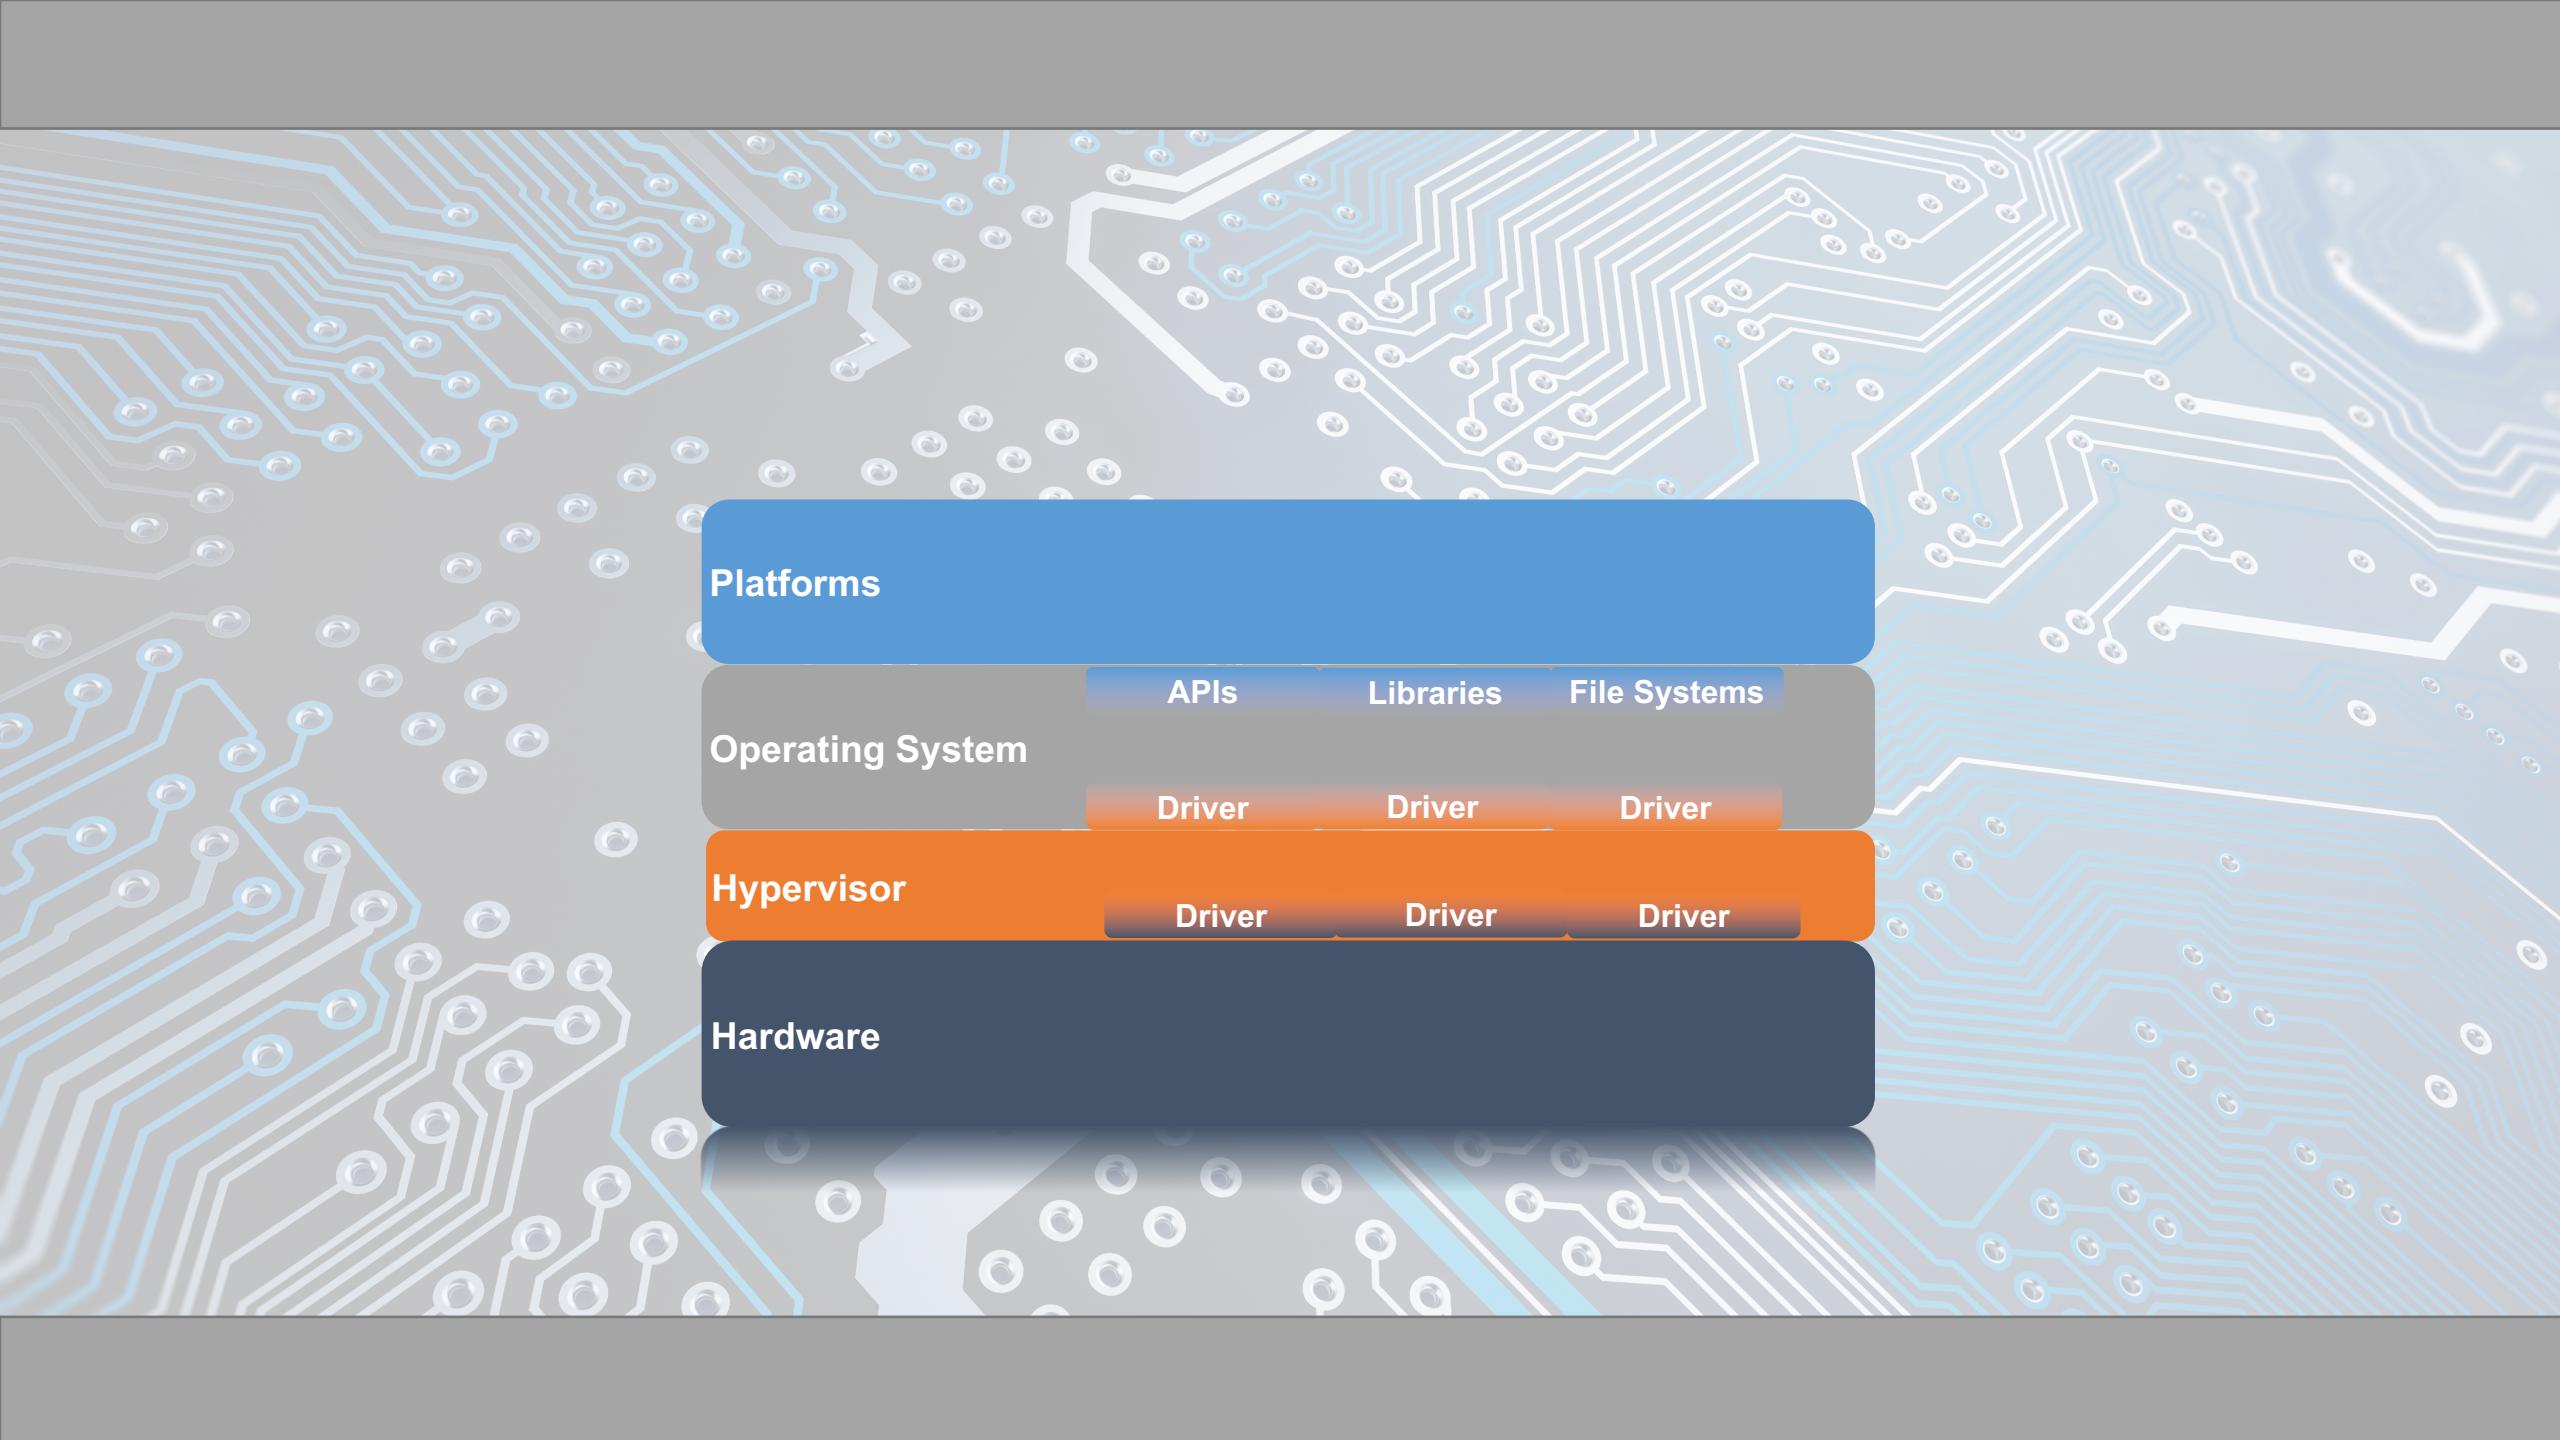
Driver

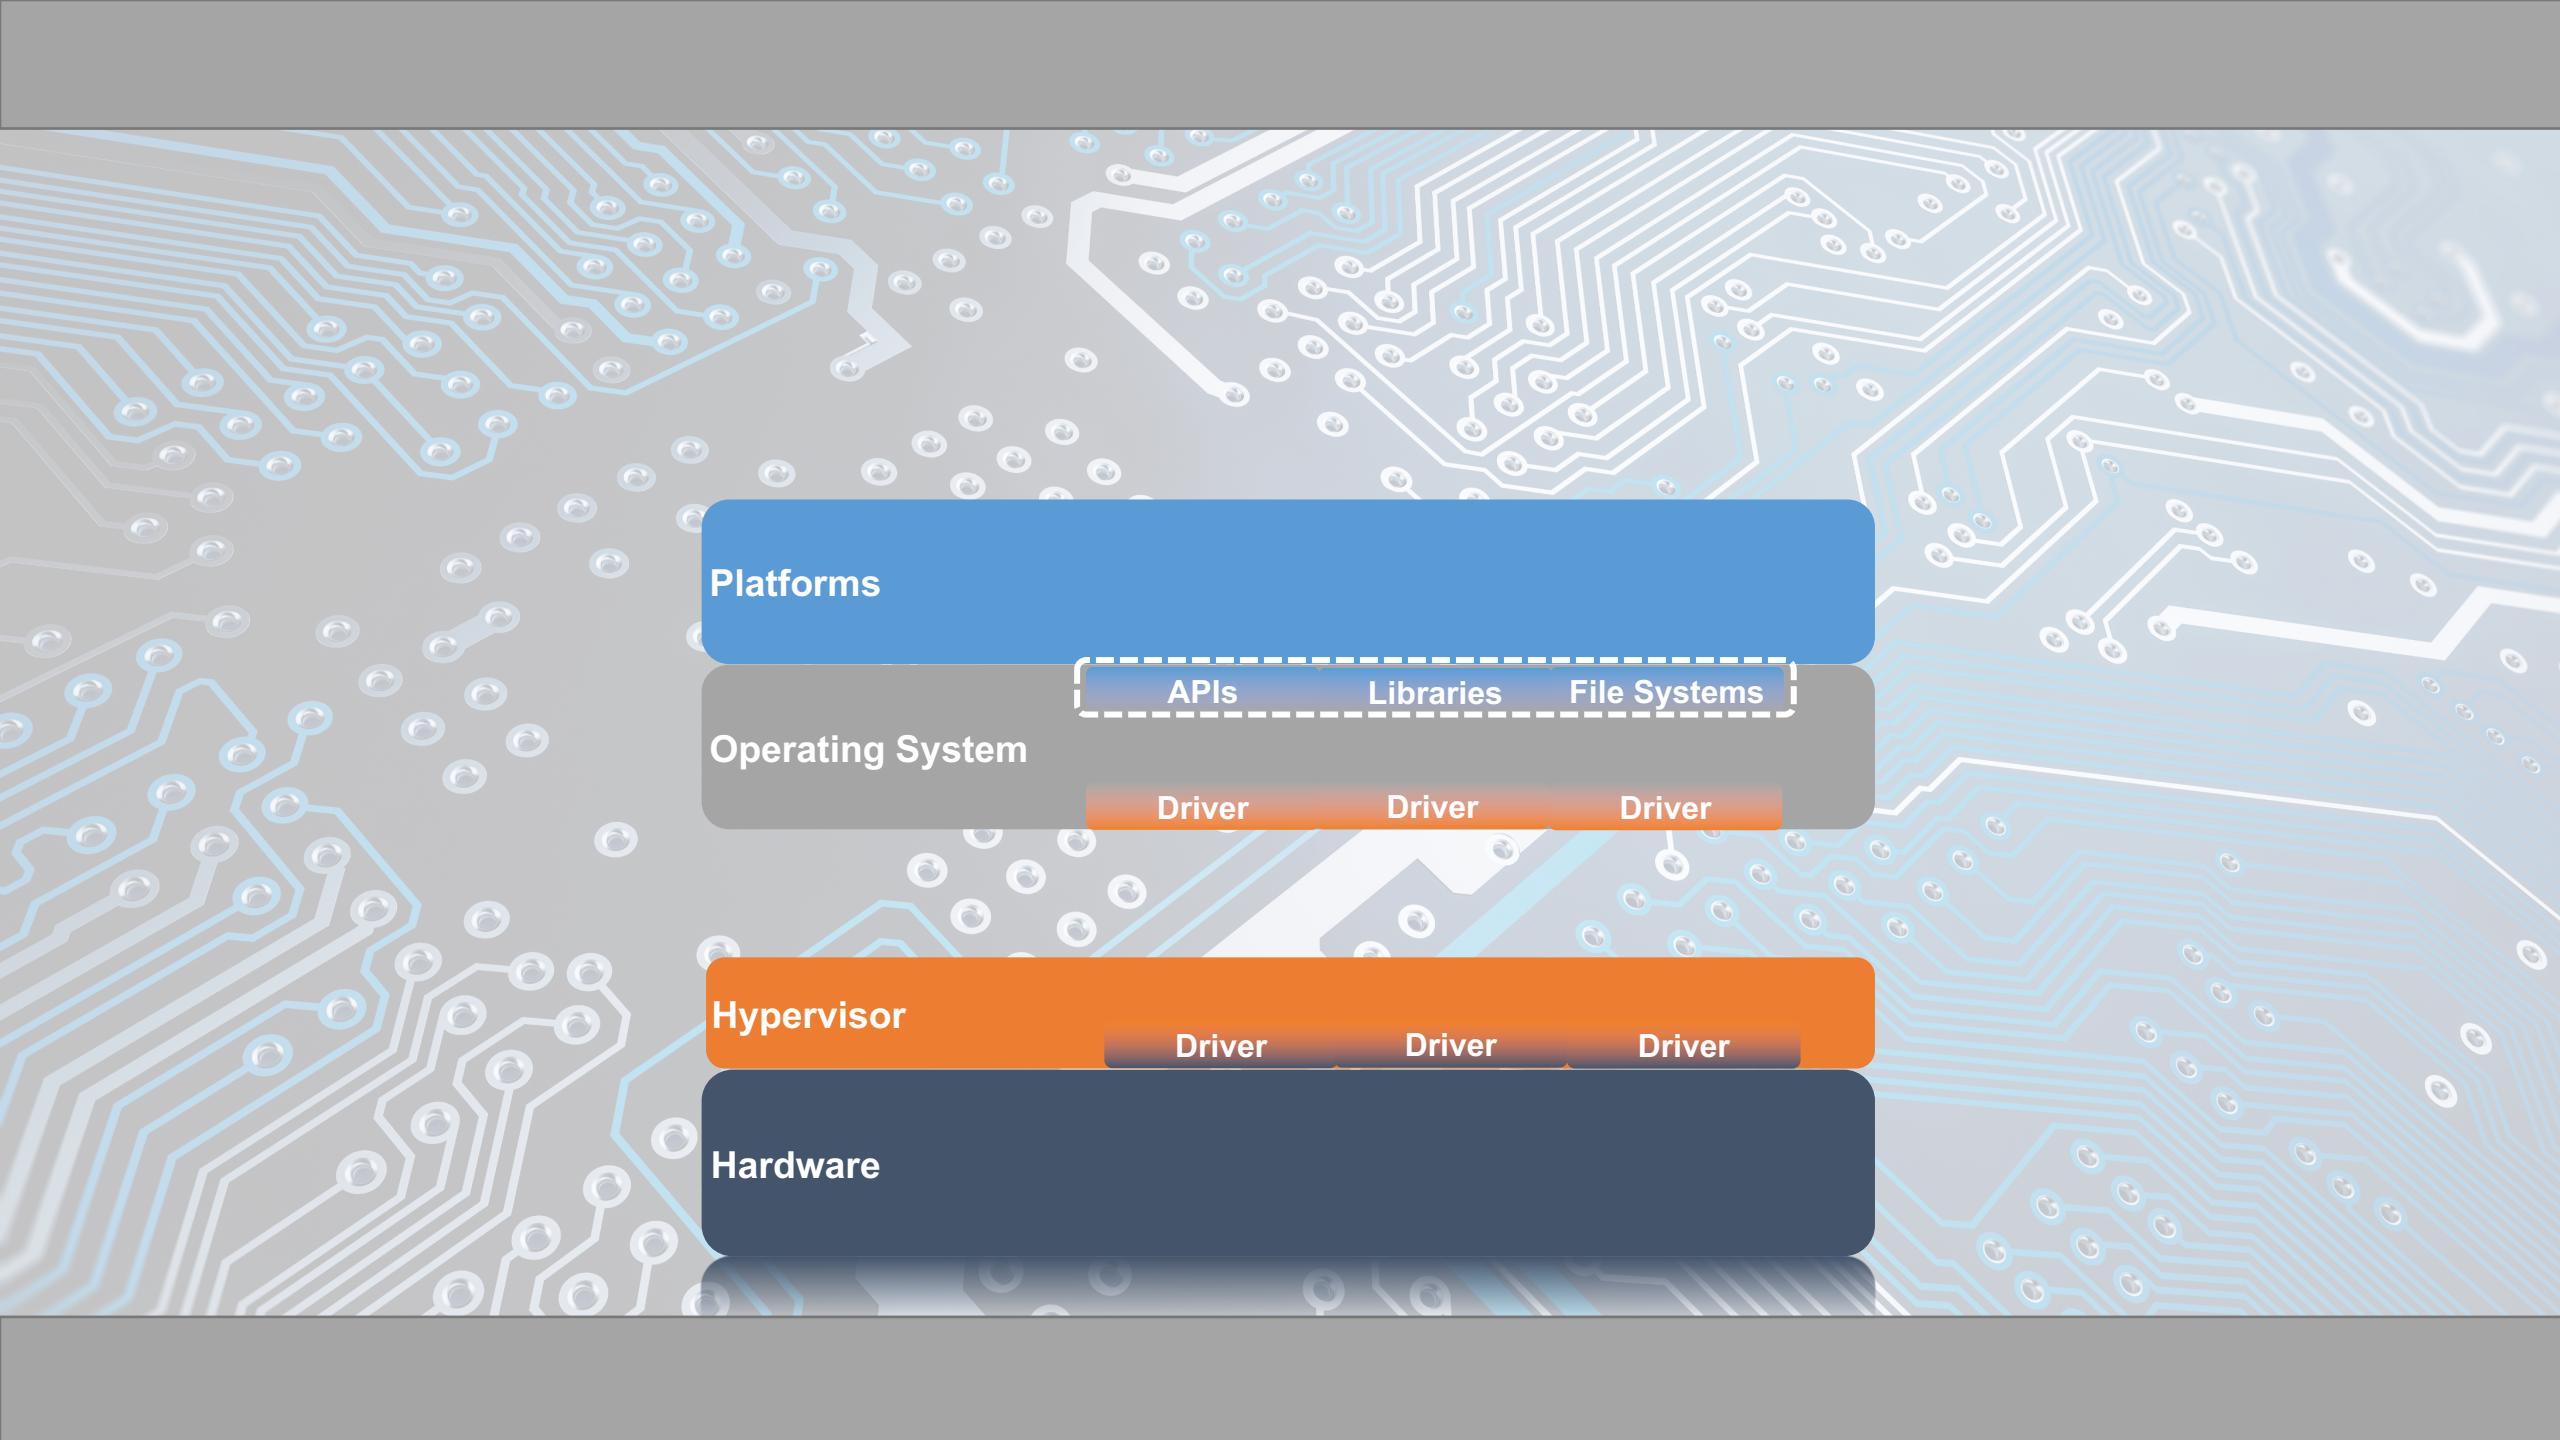
Hardware

Driver

Driver

Driver





A diagram illustrating the layers of a computing system architecture, set against a background of a printed circuit board (PCB) with blue traces and grey components. The layers are represented by rounded rectangles of increasing size from bottom to top.

- Hardware**: The base layer, colored dark grey.
- Hypervisor**: A layer above Hardware, colored orange. It contains three "Driver" labels, each aligned under a corresponding "APIs", "Libraries", and "File Systems" label from the Operating System layer.
- Operating System**: A layer above Hypervisor, colored grey. It contains three "Driver" labels, each aligned under a corresponding "APIs", "Libraries", and "File Systems" label from the Platforms layer.
- Platforms**: The top layer, colored blue. It contains three "Driver" labels, each aligned under a corresponding "APIs", "Libraries", and "File Systems" label from the Operating System layer.

APIs

Libraries

File Systems

Operating System

Driver

Driver

Driver

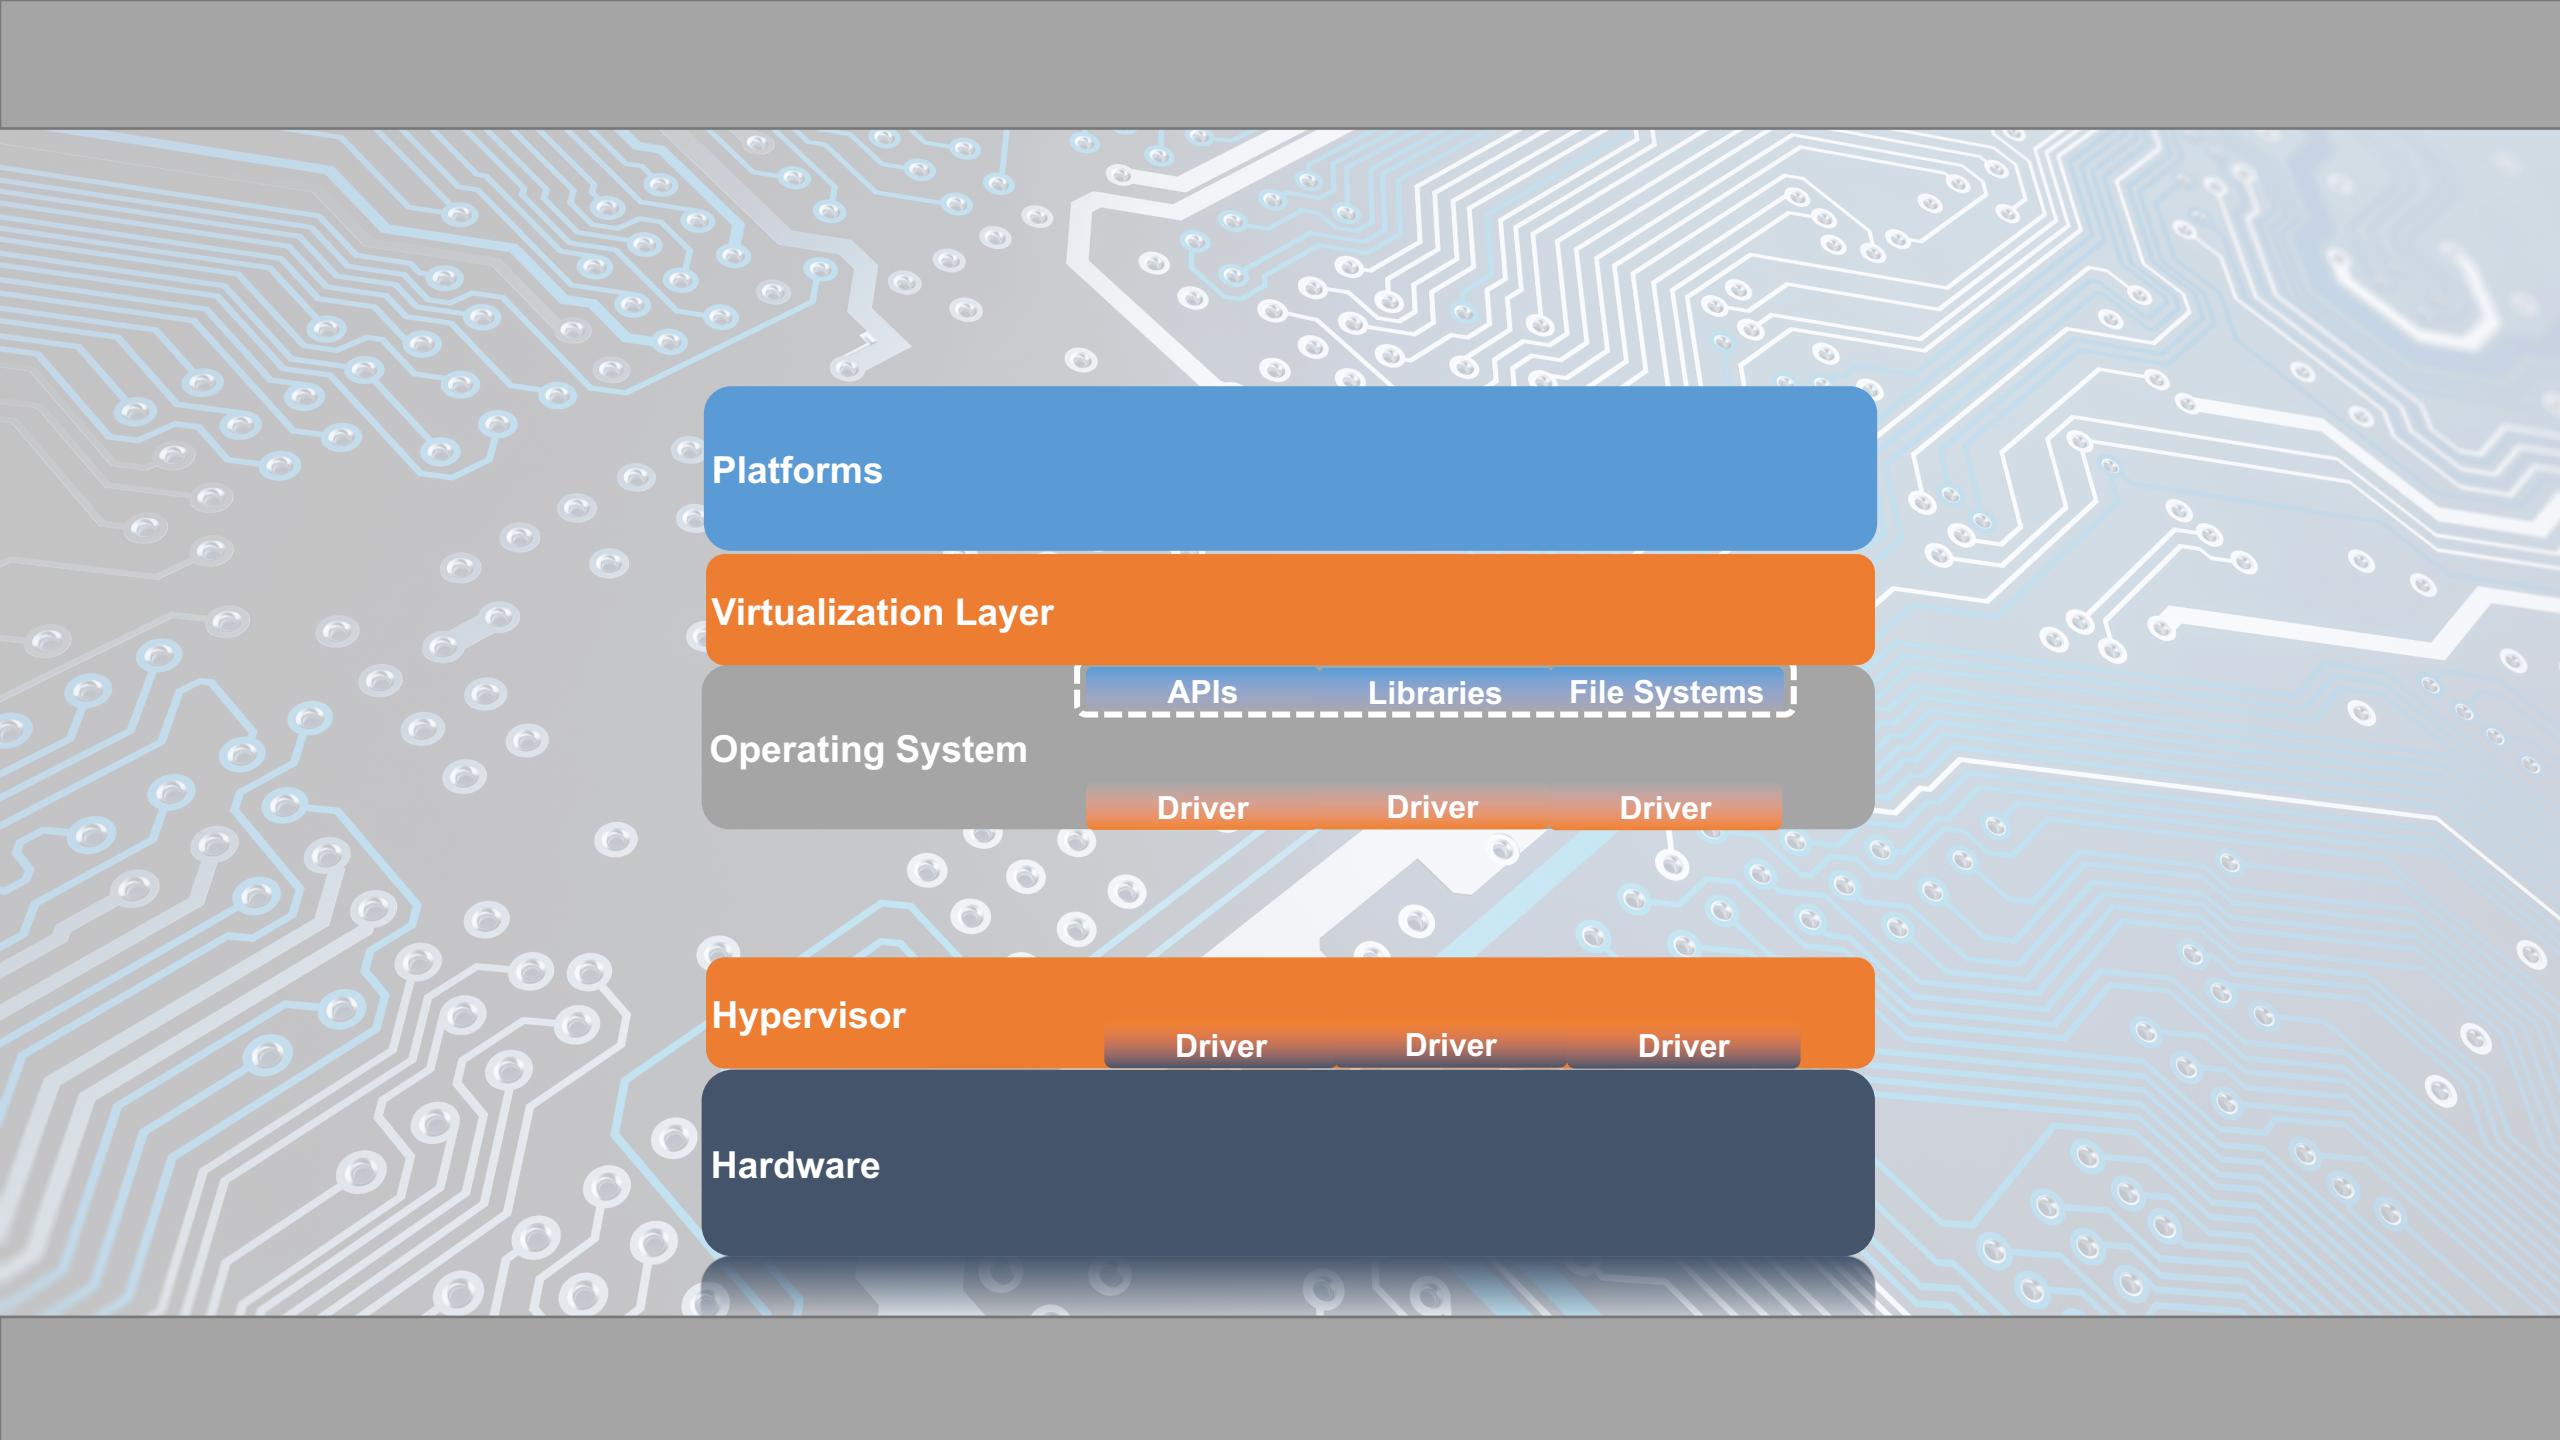
Hypervisor

Driver

Driver

Driver

Hardware



The diagram illustrates the architecture of nested virtualization layers, set against a background of a printed circuit board (PCB) with blue traces and grey components. It consists of five horizontal layers:

- Platforms**: The topmost layer, colored blue.
- Virtualization Layer**: The second layer from the top, colored orange.
- Operating System**: The third layer, colored grey. It contains three sub-components: **APIs**, **Libraries**, and **File Systems**, which are separated by dashed lines and colored blue, orange, and blue respectively.
- Hypervisor**: The fourth layer, colored orange. It contains three **Driver** components, which are colored orange.
- Hardware**: The bottommost layer, colored dark grey.

**Platforms**

**Container**

**Operating System**

APIs

Libraries

File Systems

Platforms

Container

Operating System

APIs

Libraries

File Systems

**Platforms**

APIs

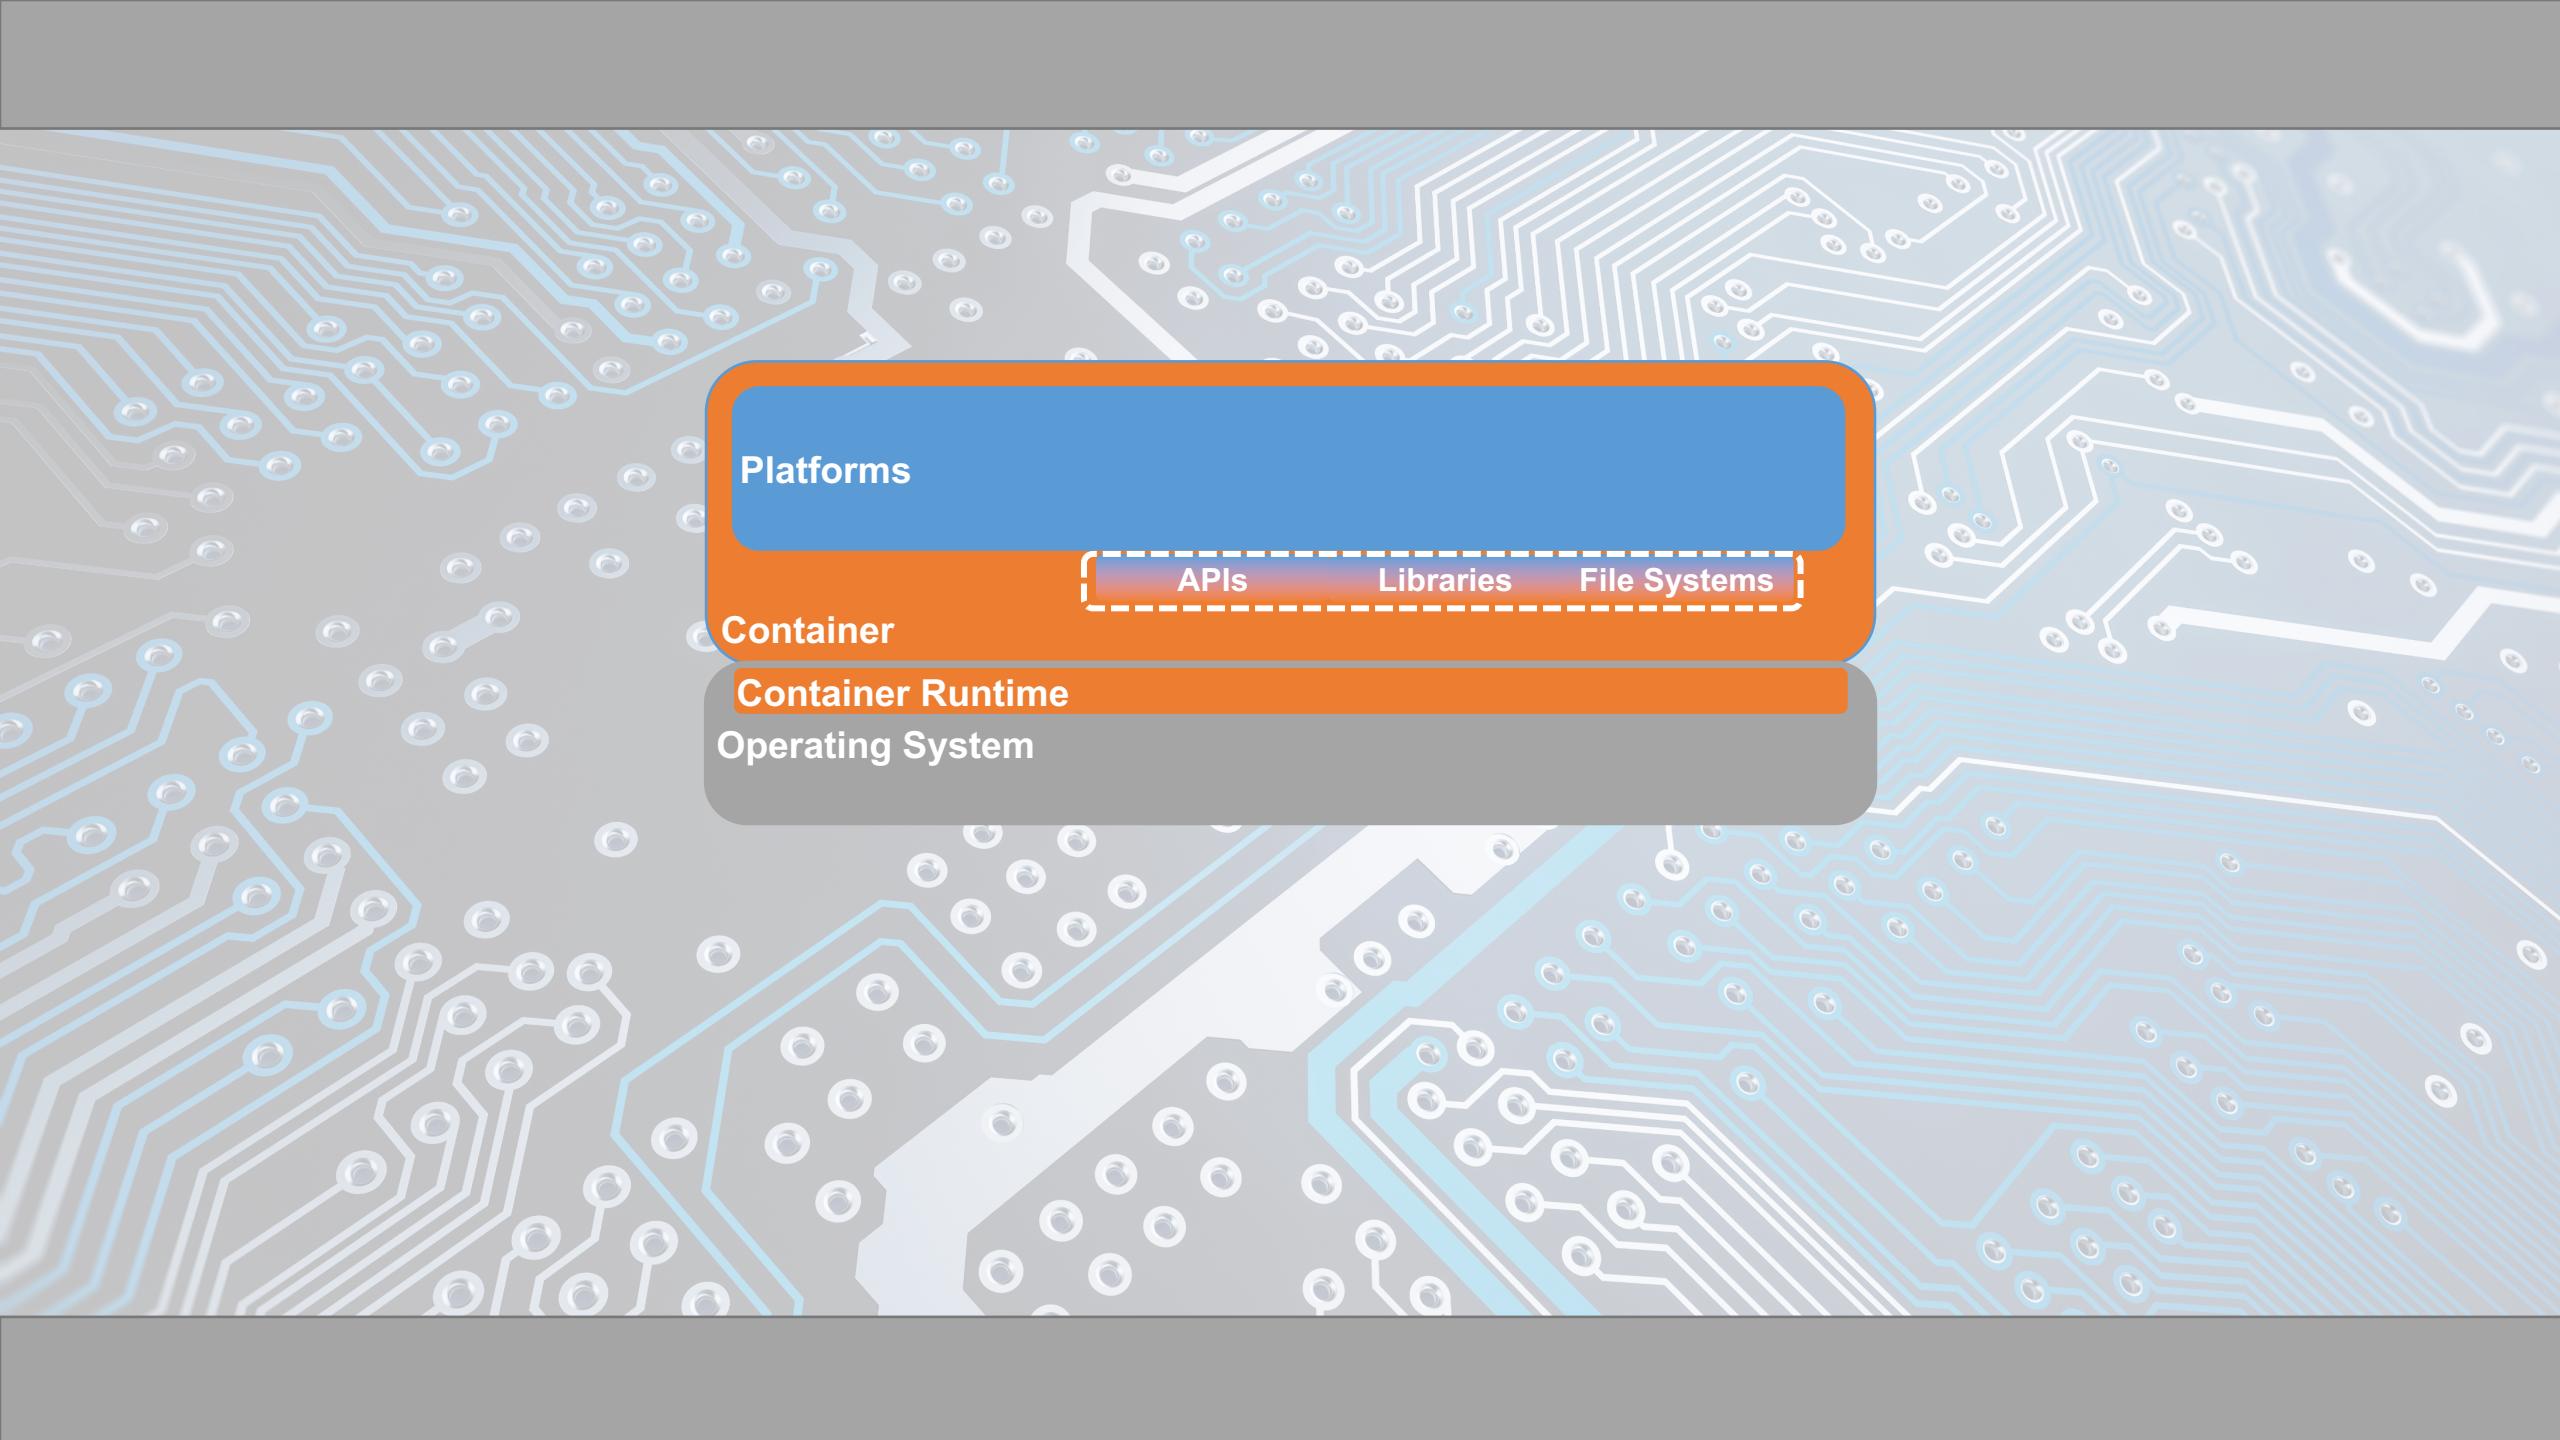
Libraries

File Systems

**Container**

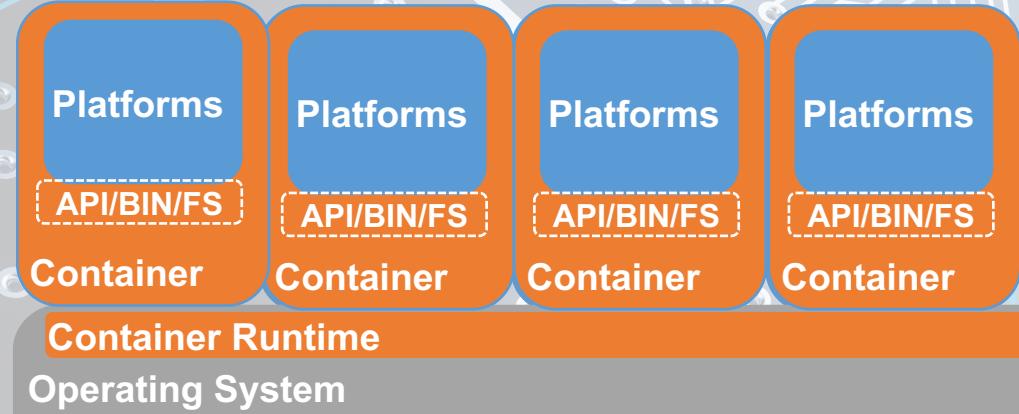
**Container Runtime**

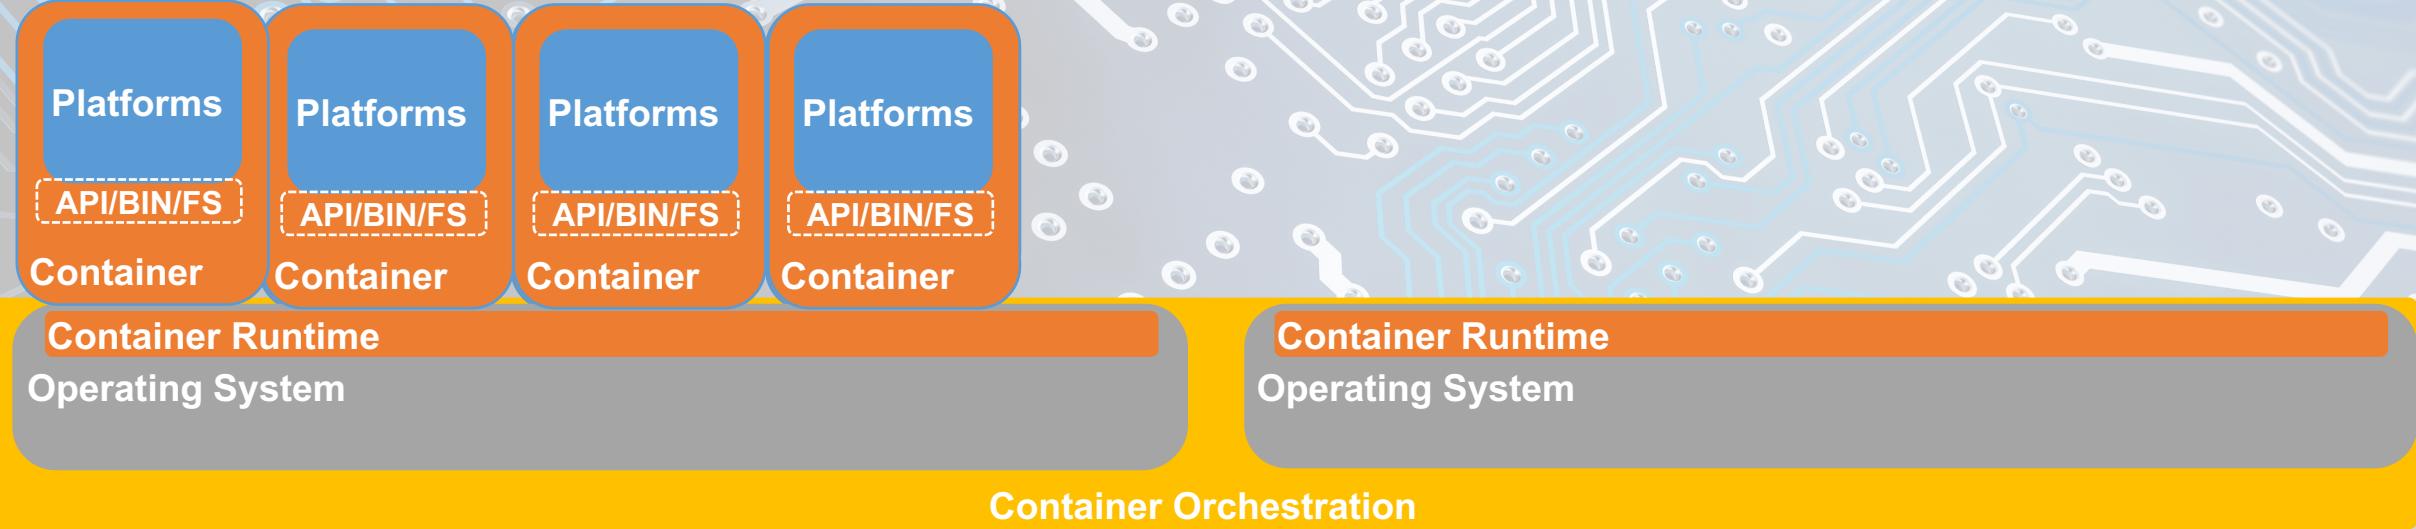
**Operating System**

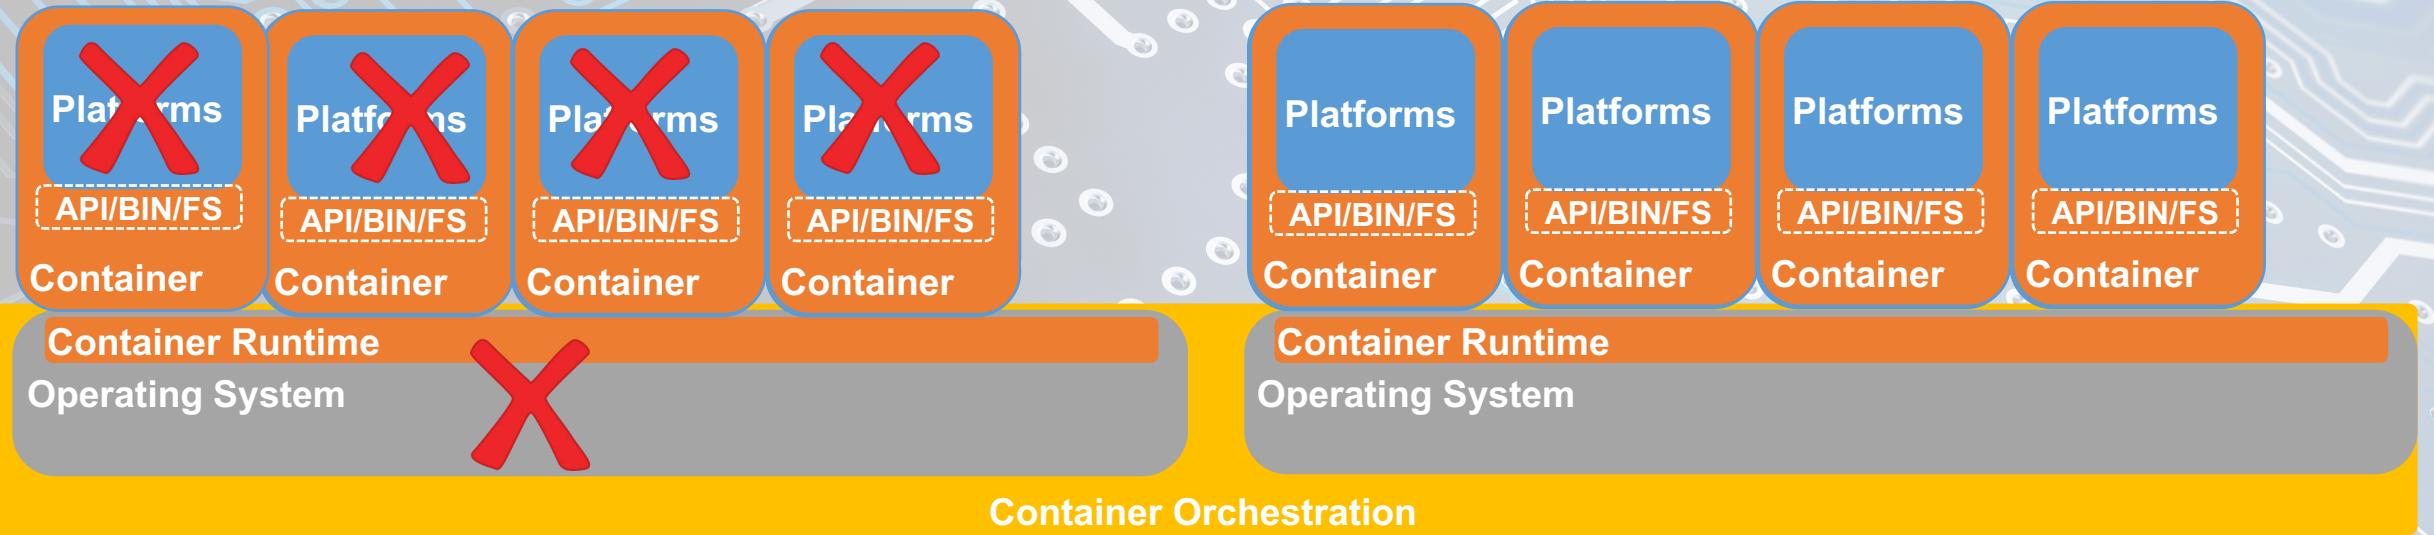


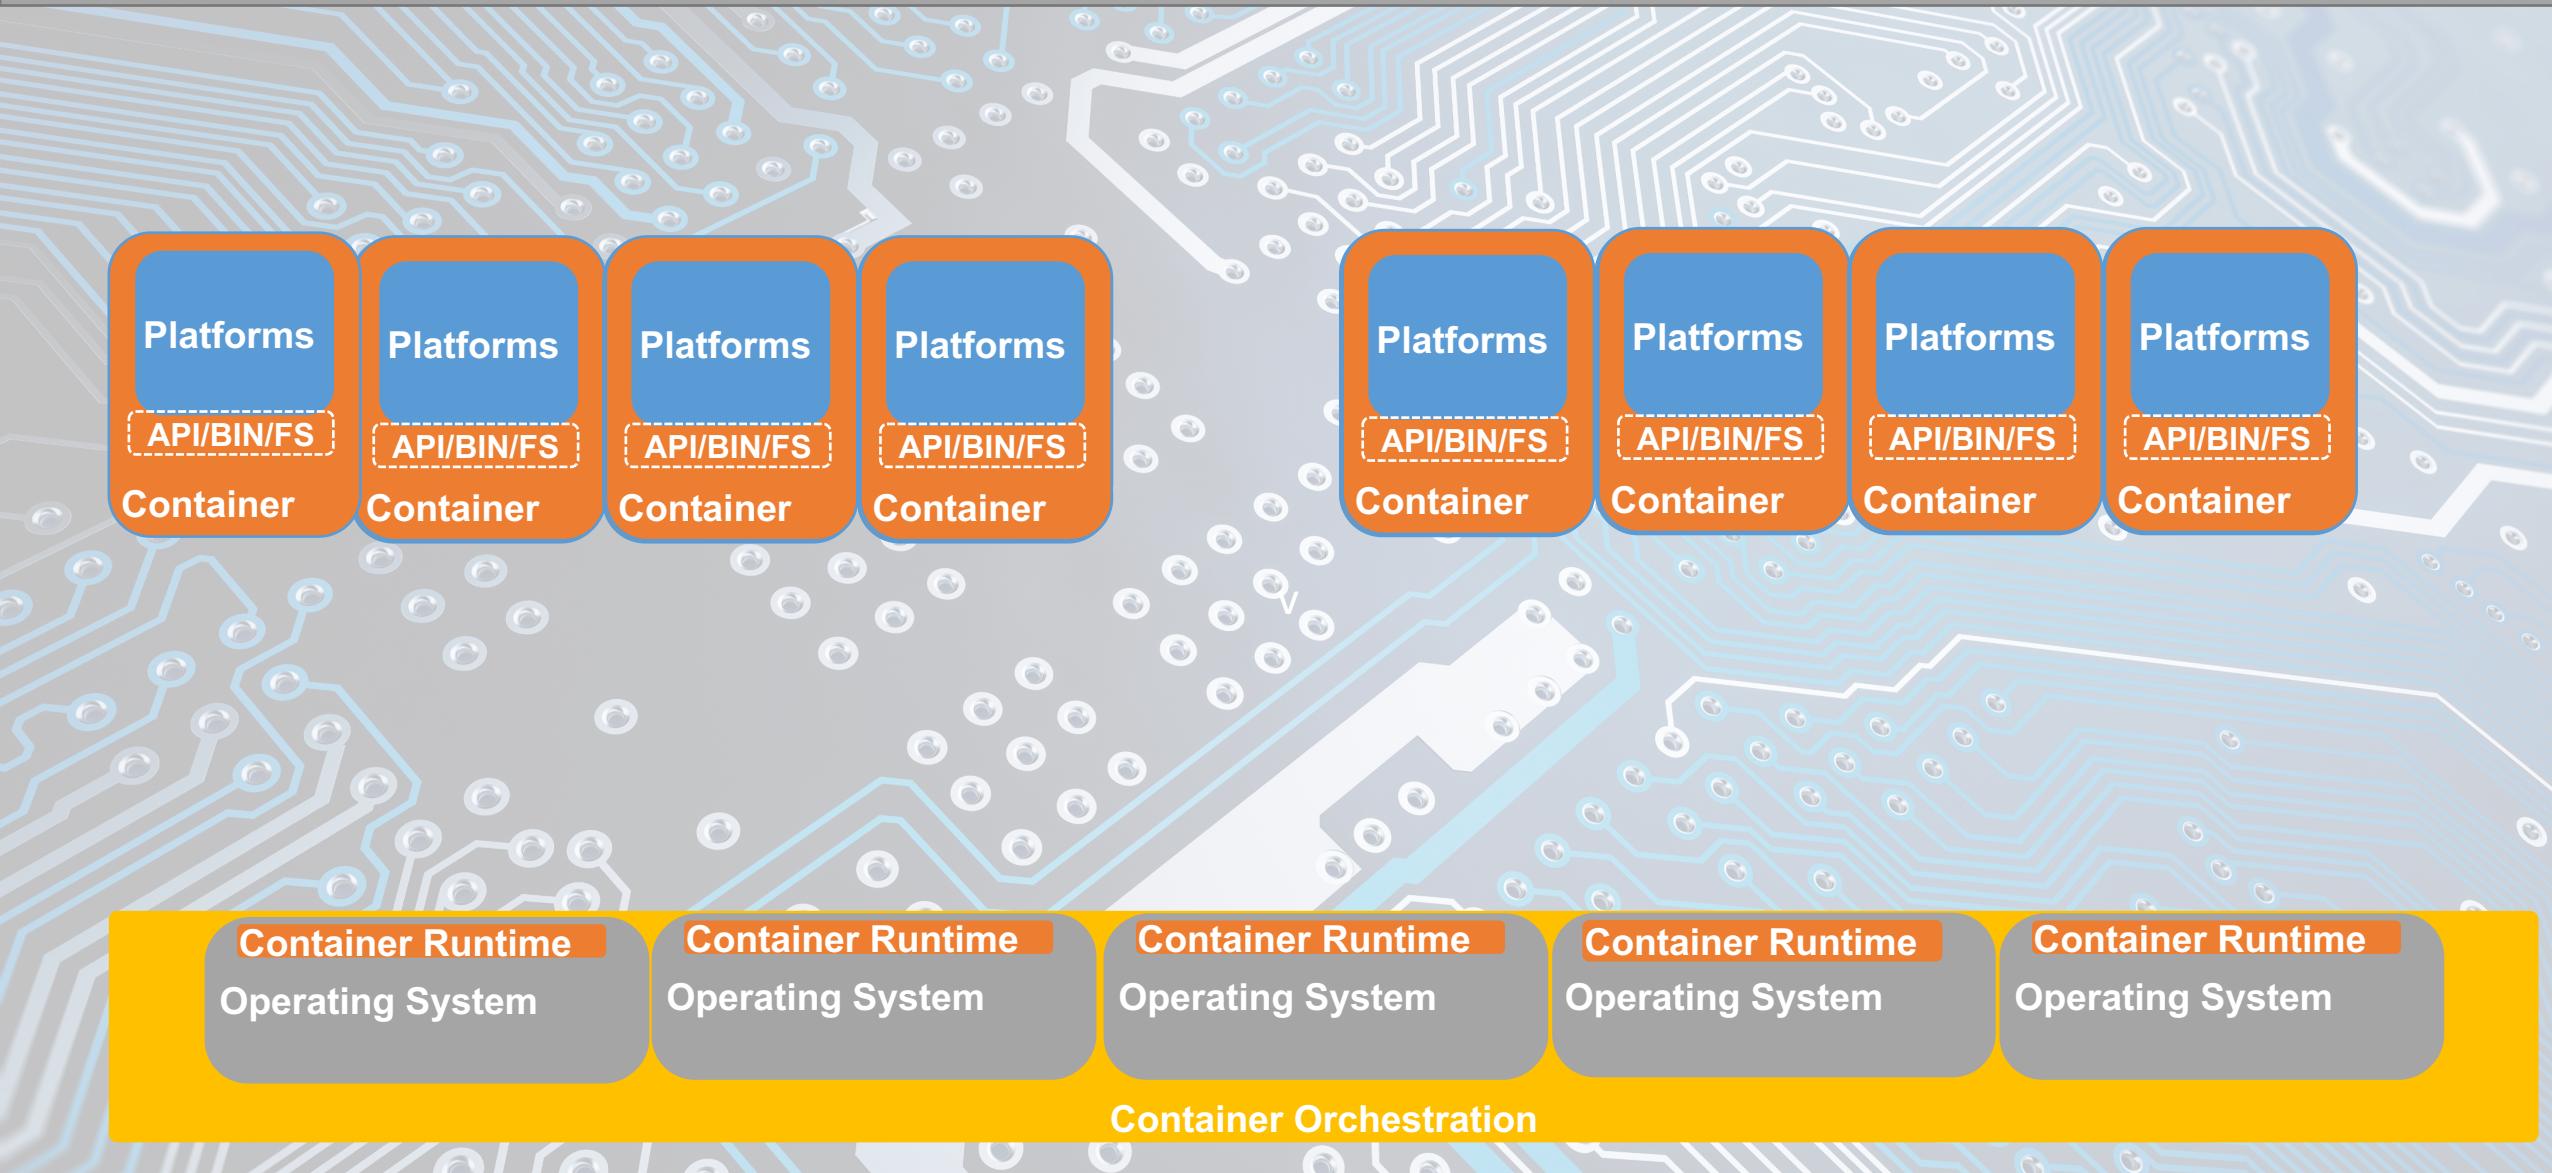


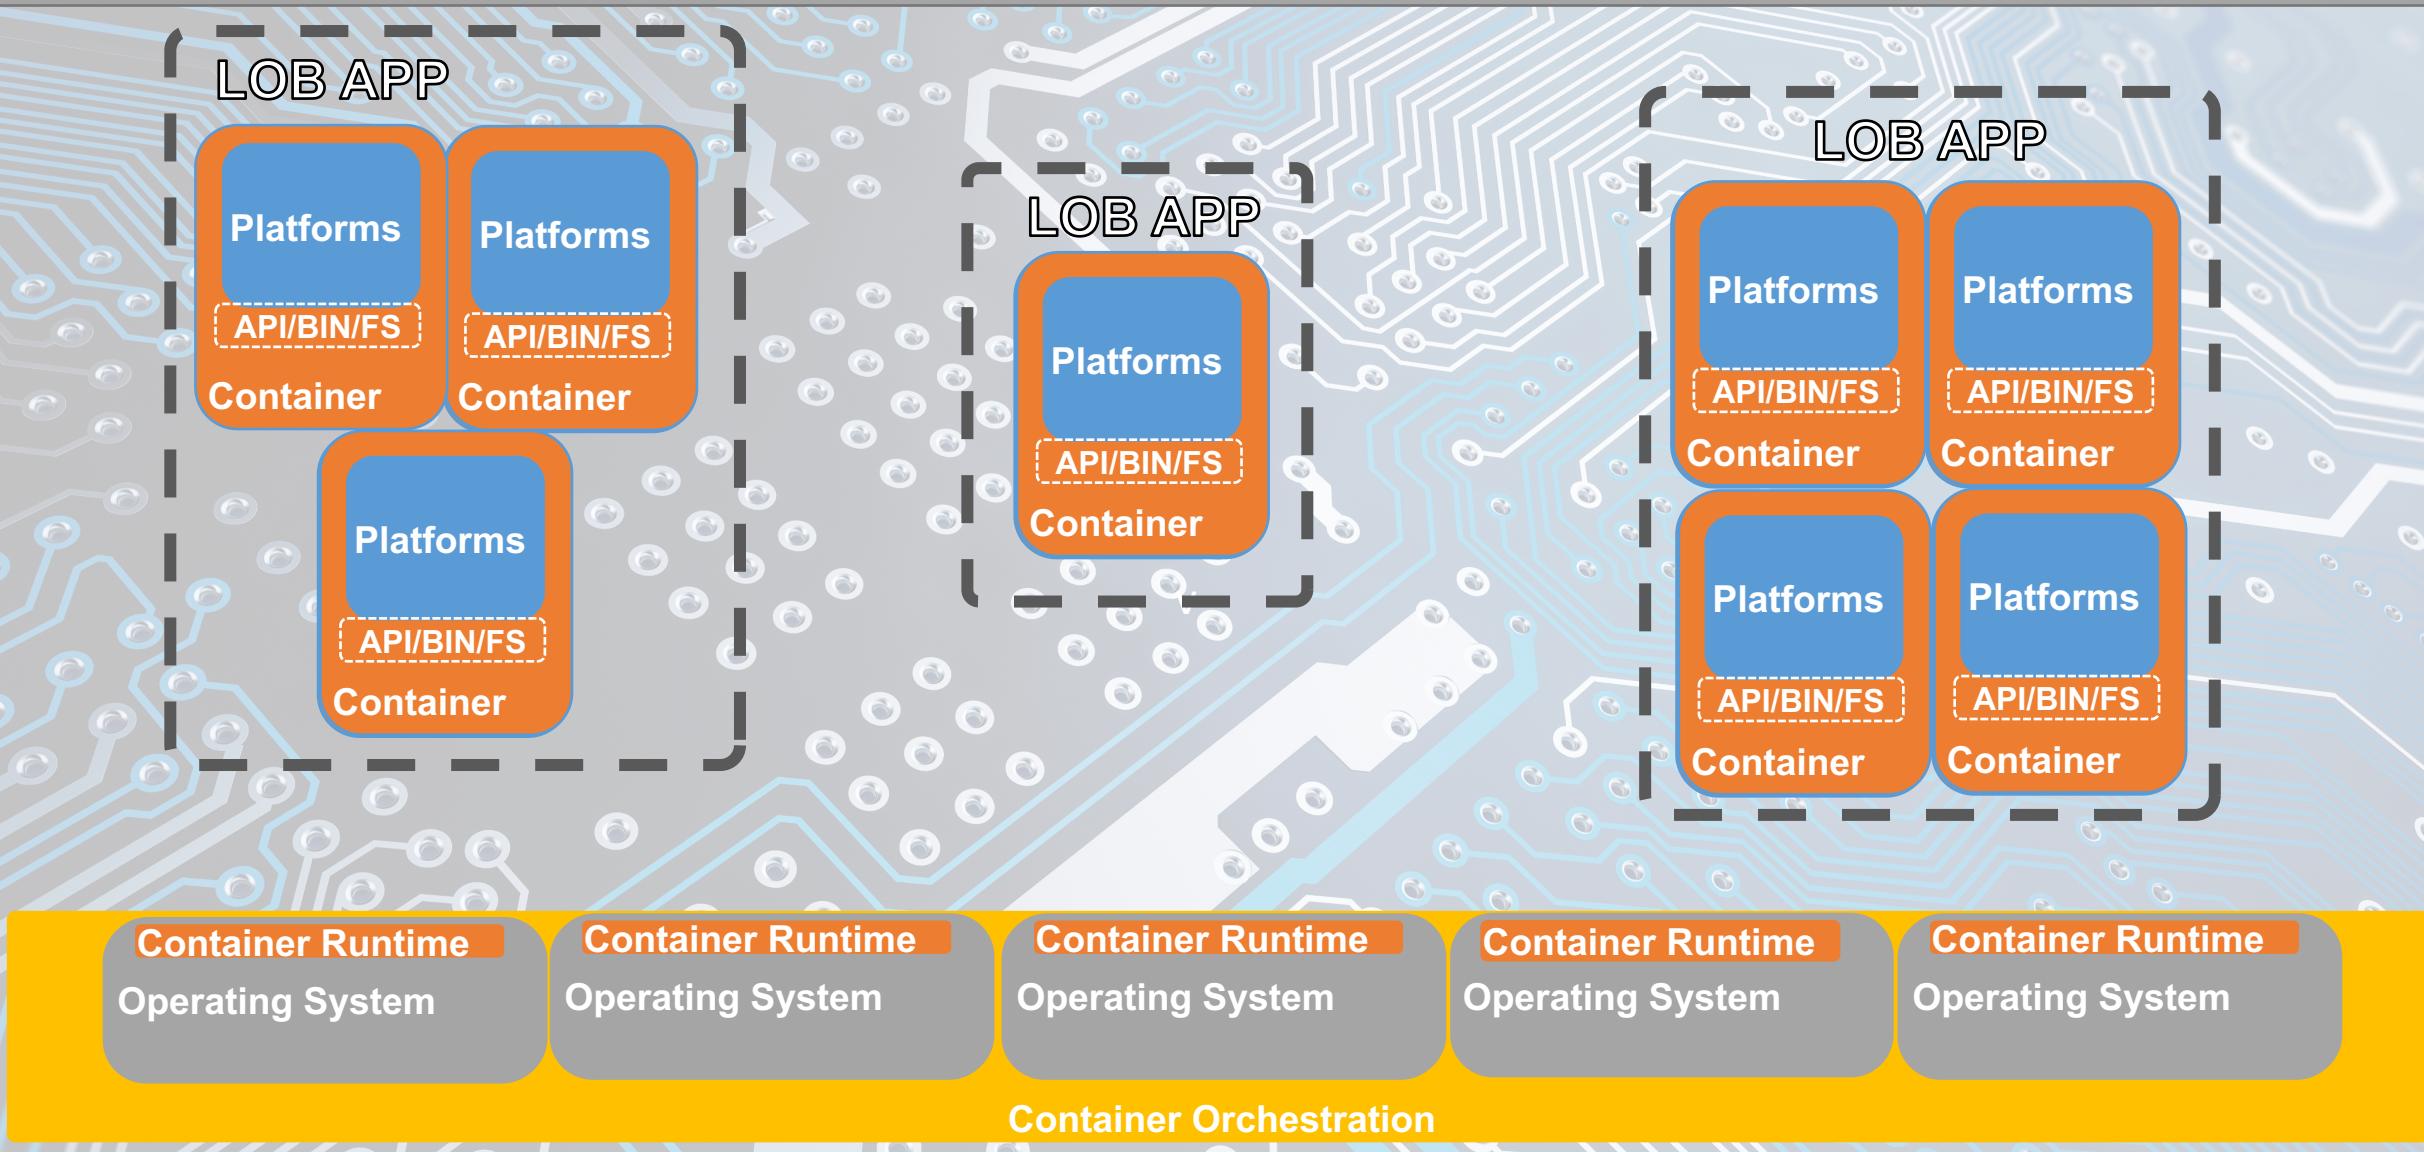
**Container Runtime**  
Operating System

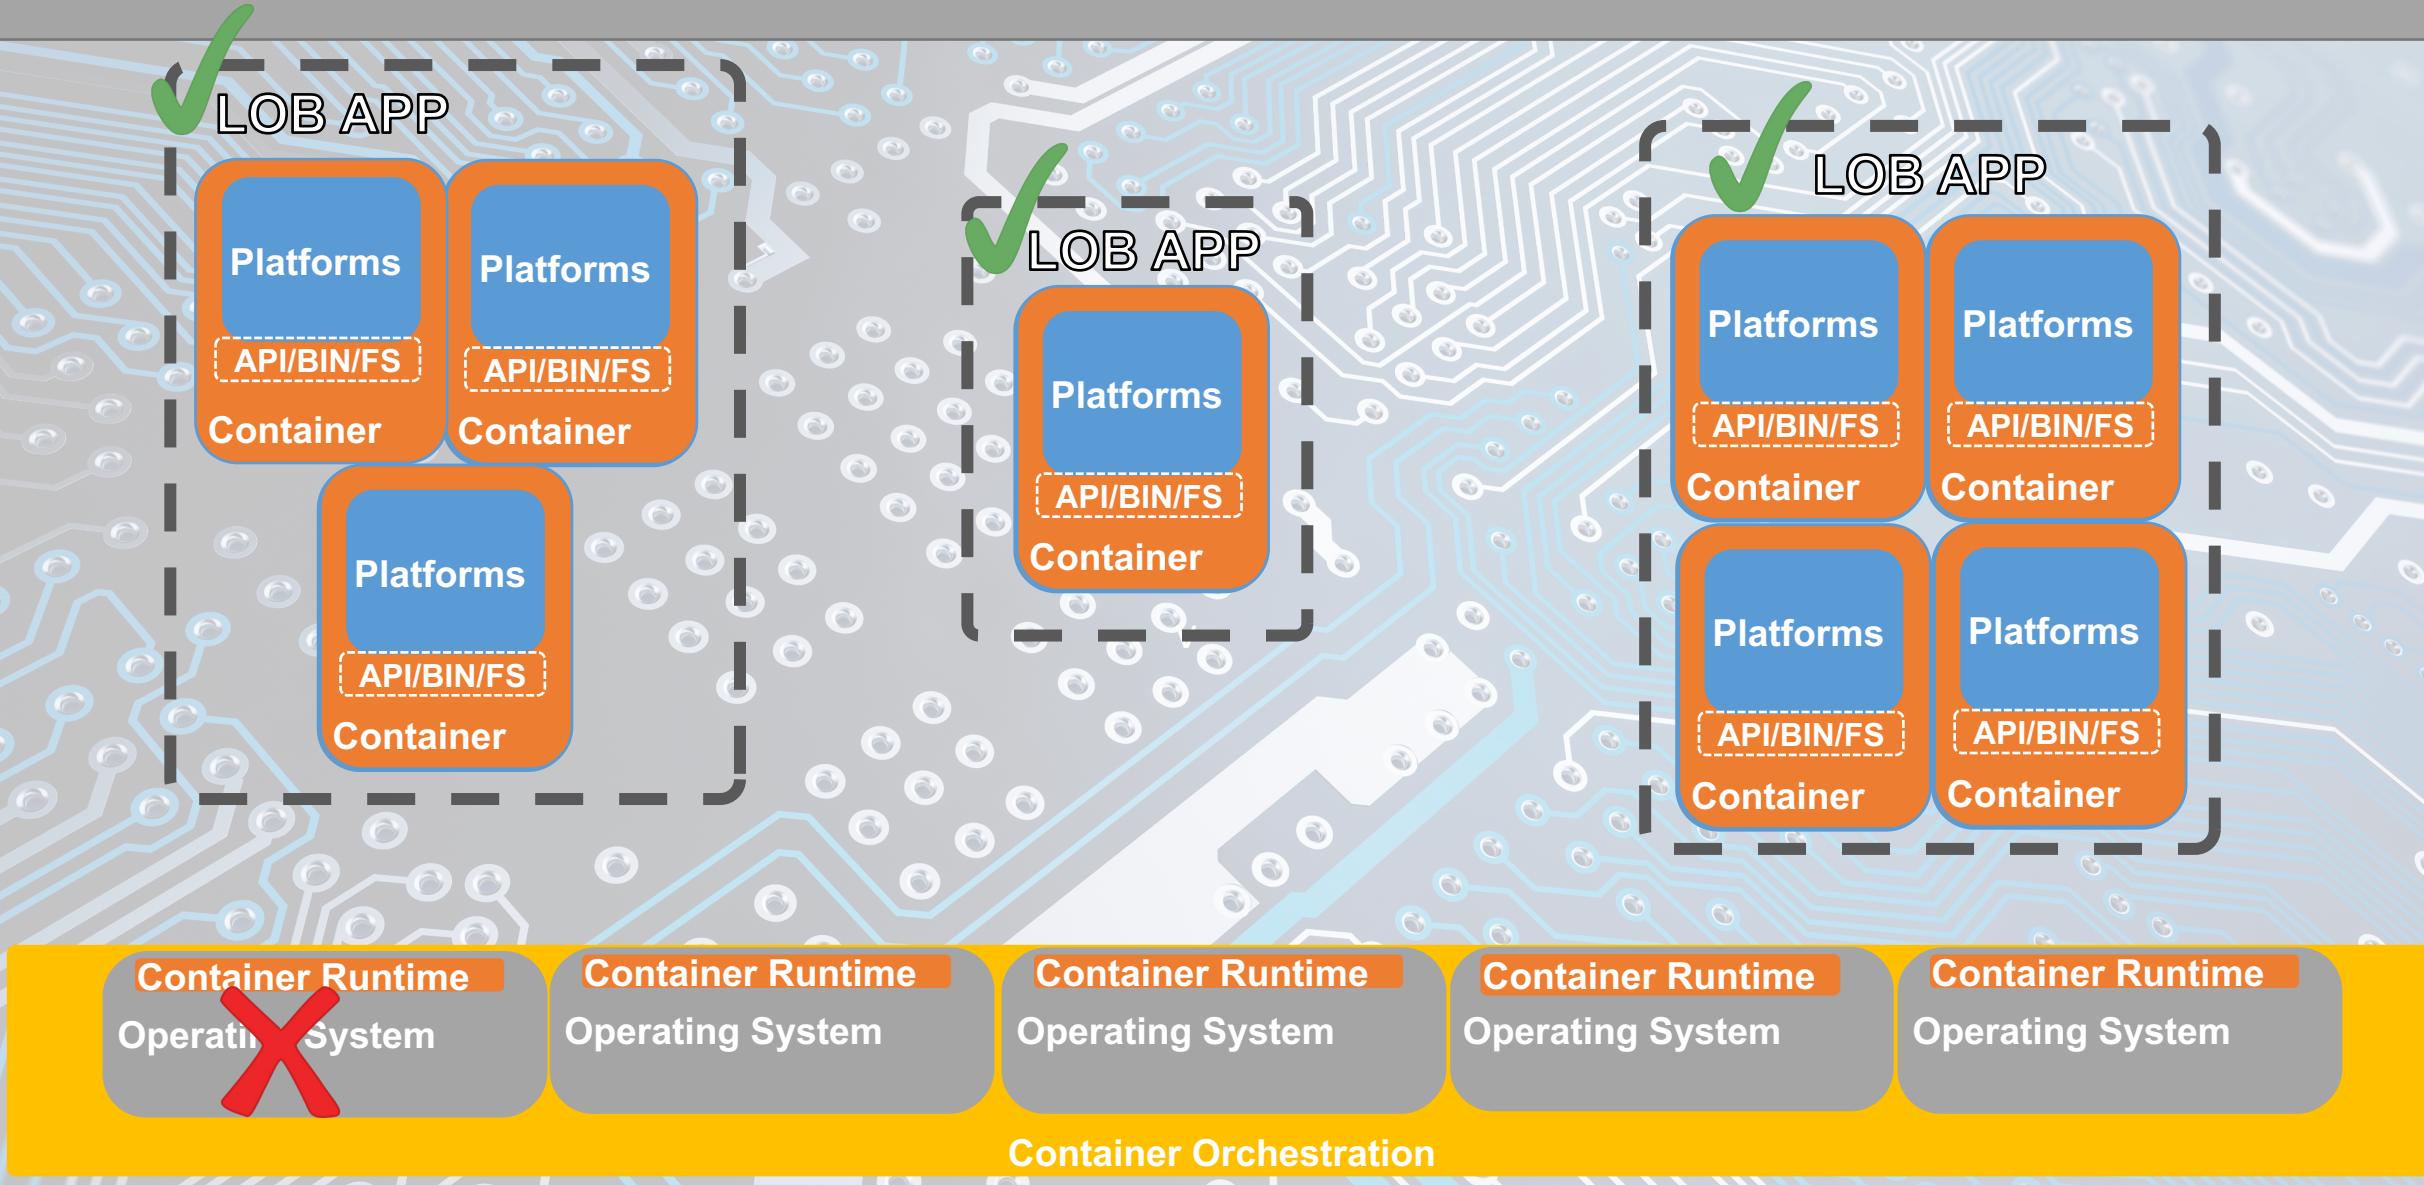


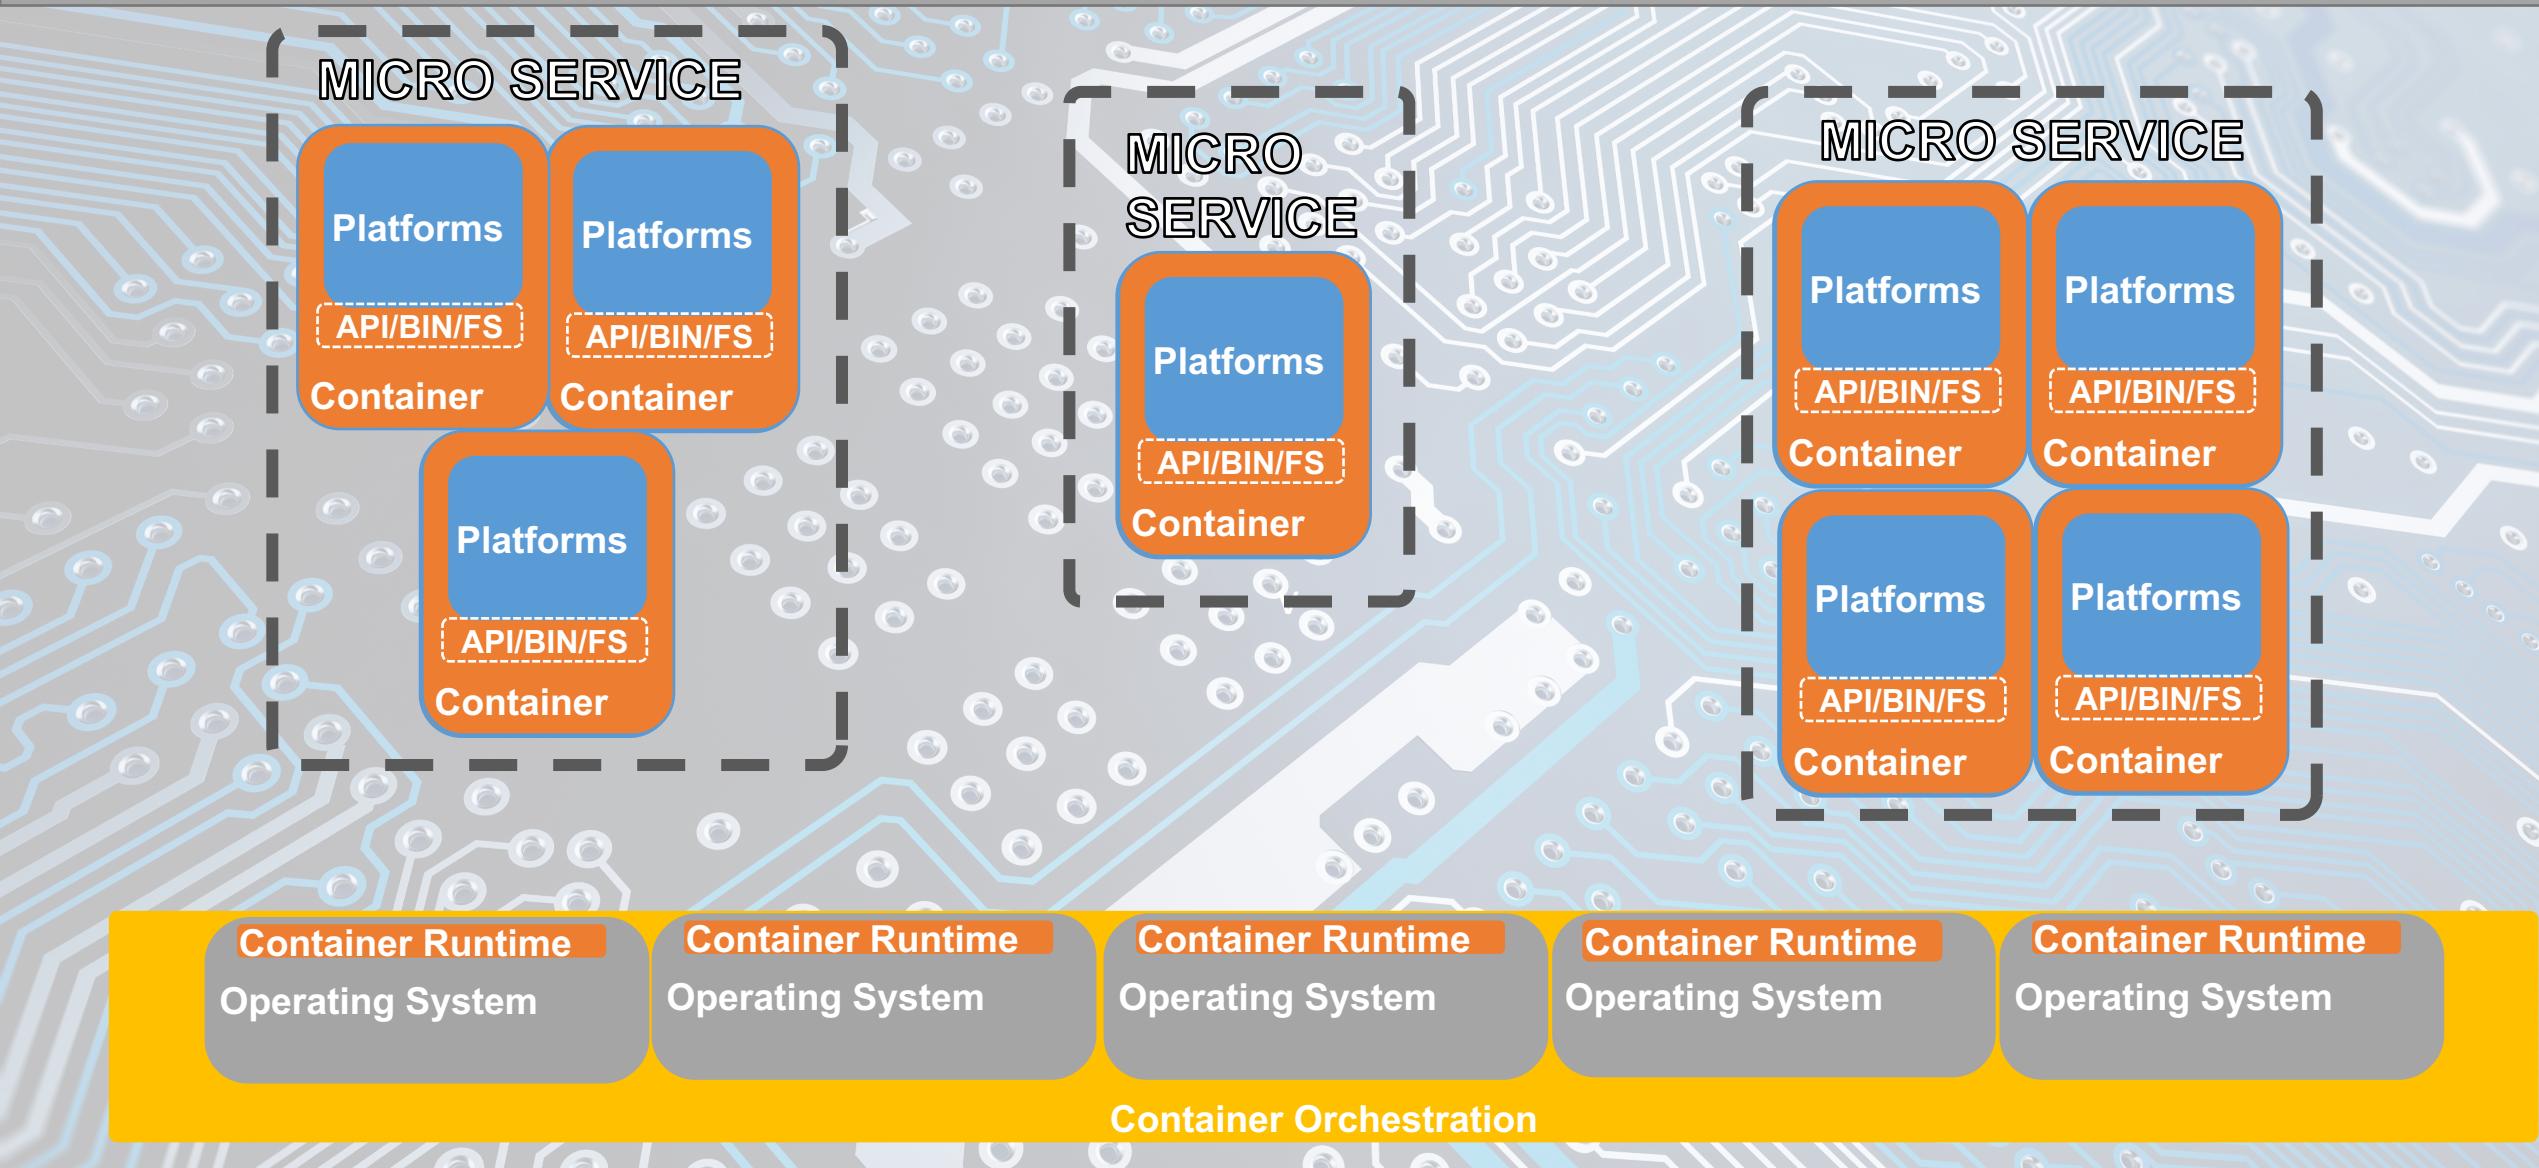


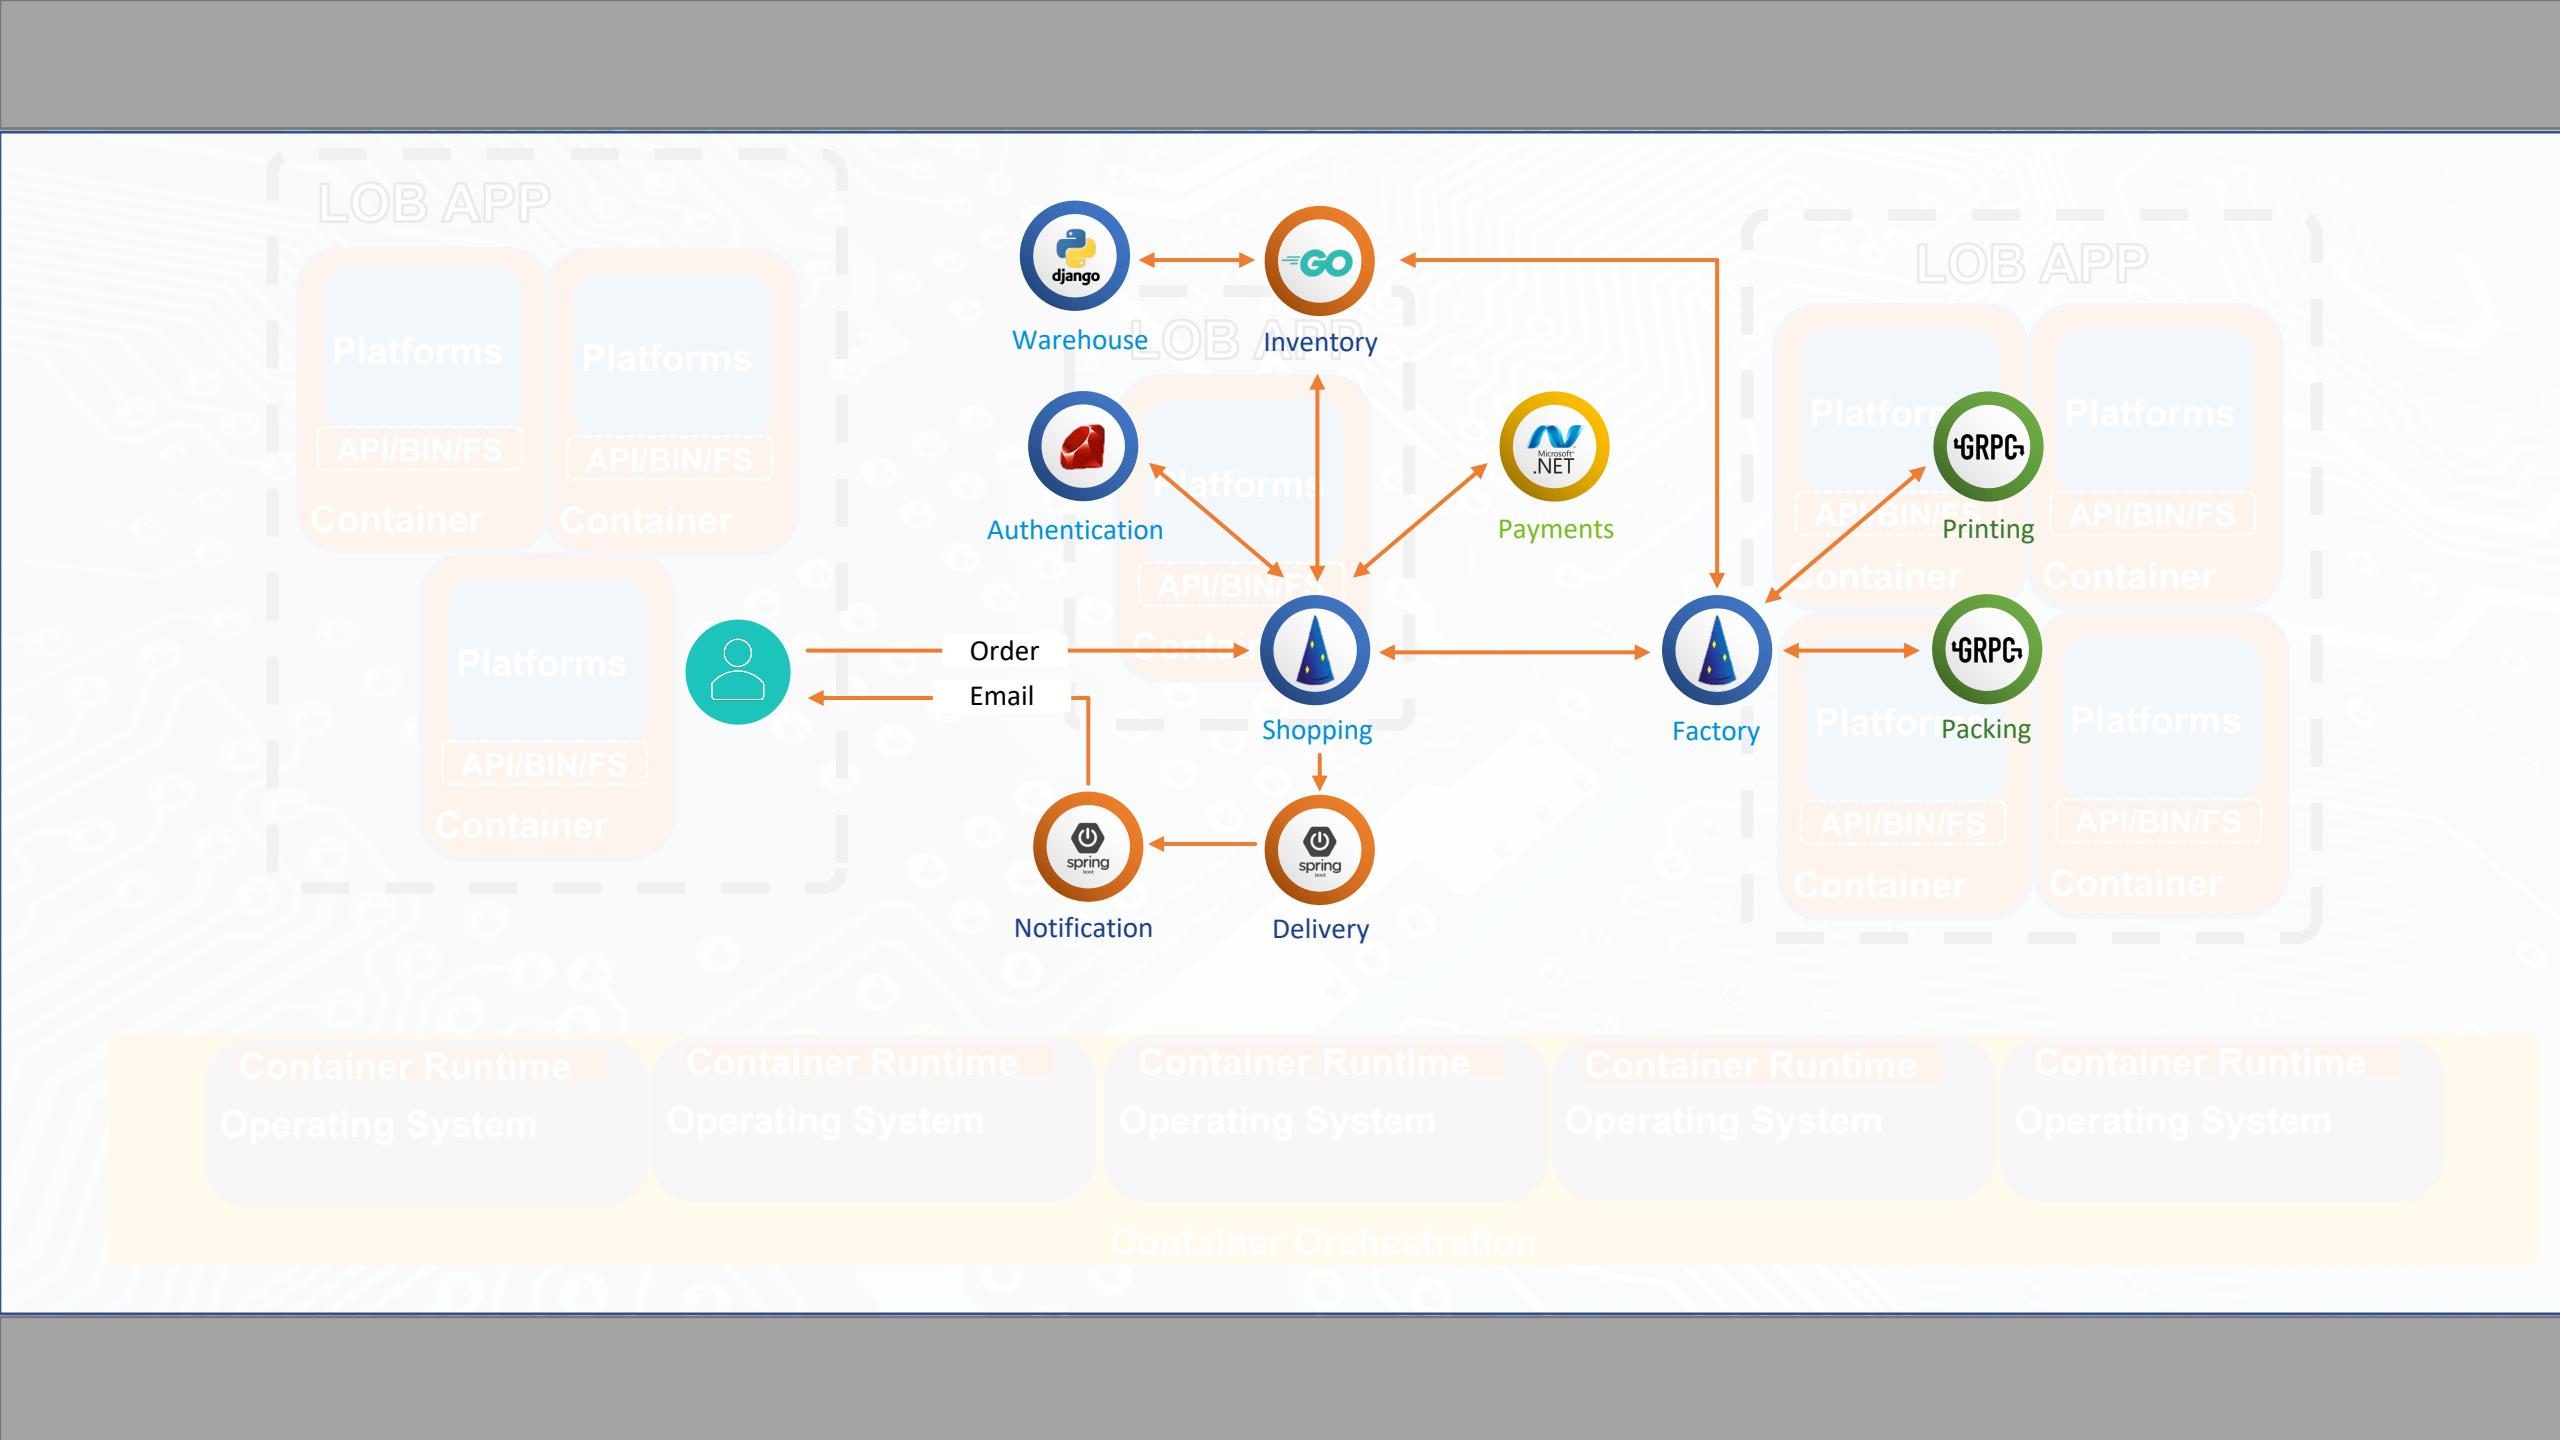


















1979

Development of “chroot”, allowed  
storage level process isolation

1979 Development of “chroot”, allowed storage level process isolation

2000

“jail” command release, extending  
isolation capability to users,  
network & more.

1979 Development of “chroot”, allowed storage level process isolation  
2000 “jail” command release, extending isolation capability to users, network & more.

2003

Google begins work on BORG

1979 Development of “chroot”, allowed storage level process isolation

2000 “jail” command release, extending isolation capability to users, network & more.

2003 Google begins work on BORG

2004

Solaris Zones released, providing user, process and filesystem isolation.

1979 Development of “chroot”, allowed storage level process isolation

2000 “jail” command release, extending isolation capability to users, network & more.

2003 Google begins work on BORG

2004 Solaris Zones released, providing user, process and filesystem isolation.

2007

AIX workload partitions release.  
Providing containers for Power.

network & more.

2003 Google begins work on BORG

2004 Solaris Zones released, providing user, process and filesystem isolation.

2007 AIX workload partitions release. Providing containers for Power.

2013

Docker announced at PyCon

network & more.

2003 Google begins work on RQBC

2004 Solaris  
isolation

2007 AIX wo

IN TARGET  
0 HUG

PYCON US 2013

The future of Linux Containers

presented by

Solomon Hykes



DIAMOND SPONSOR



heroku

Google



filesystem

s for Power.

on

network & more.

2003 Google begins work on BORG

2004 Solaris Zones released, providing user, process and filesystem isolation.

2007 AIX workload partitions release. Providing containers for Power.

2013

Docker announced at PyCon

network & more.

2003 Google begins work on BORG

2004 Solaris Zones released, providing user, process and filesystem isolation.

2007 AIX workload partitions release. Providing containers for Power.

2013 Docker announced at PyCon

2015

Google replaces “LMCTFY” with  
Docker engine (libcontainer)

2007 AIX workload partitions release. Providing containers for Power.  
2013 Docker announced at PyCon

2015

Google replaces “LMCTFY” with  
Docker engine (libcontainer)

Kubernetes 1.0 Released

2007 AIX workload partitions release. Providing containers for Power.

2013 Docker announced at PyCon

2015 Google replaces “LMCTFY” with Docker engine (libcontainer)

Kubernetes 1.0 Released

2017

Kubernetes becomes ‘defacto’  
orchestration engine & widespread  
adoption in enterprises

2007 AIX workload partitions release. Providing containers for Power.

2013 Docker announced at PyCon

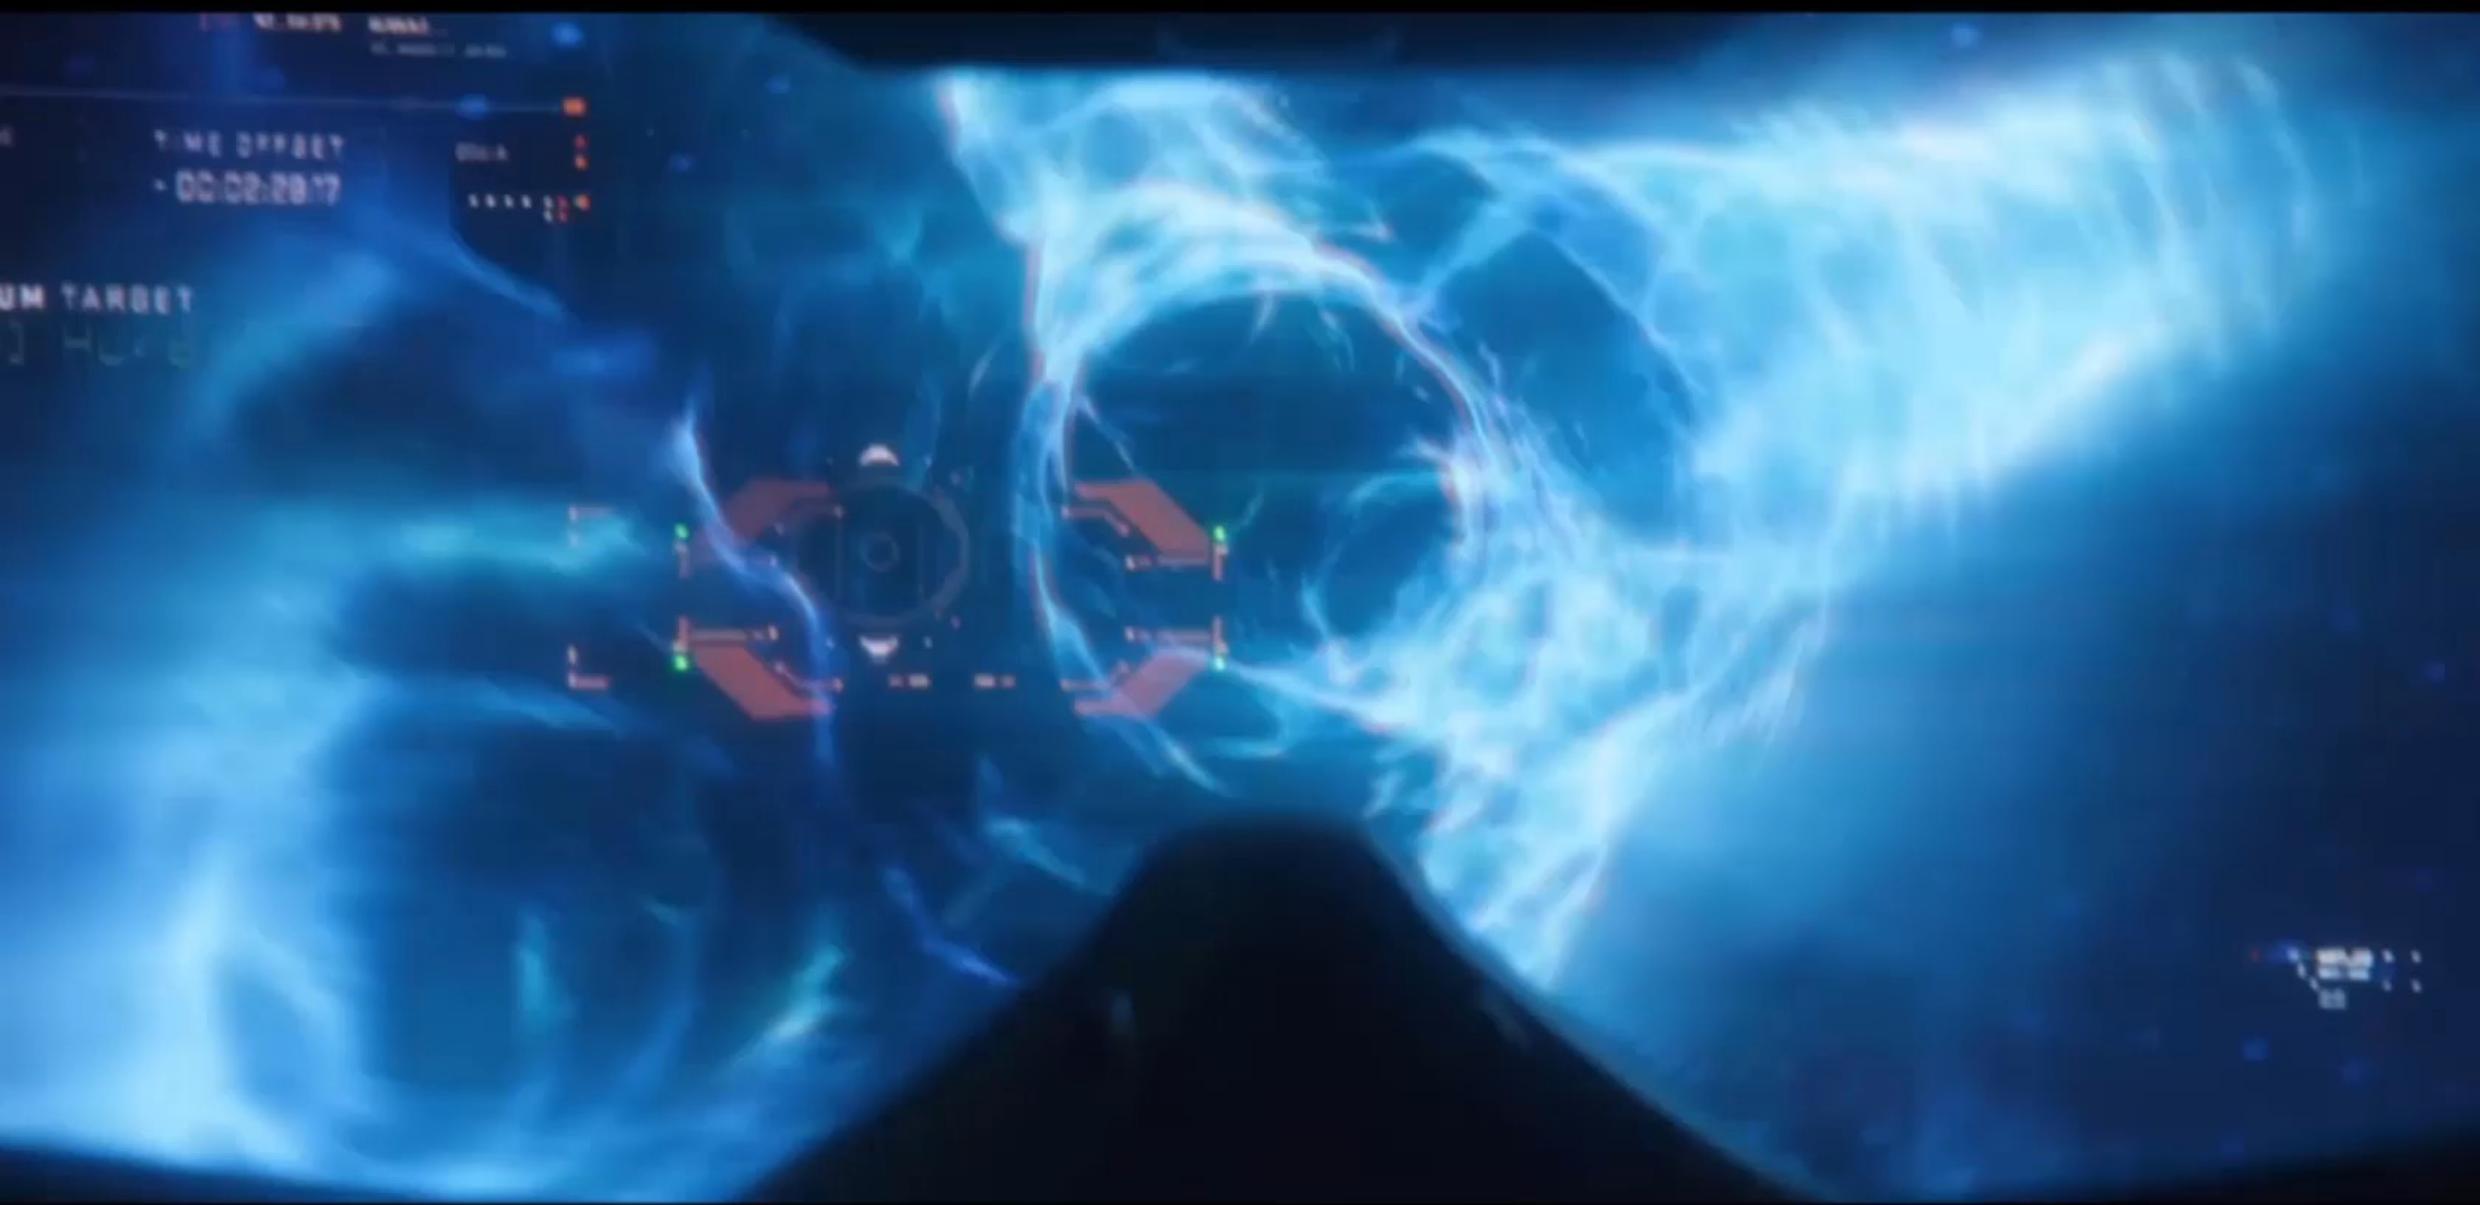
2015 Google replaces “LMCTFY” with Docker engine (libcontainer)

Kubernetes 1.0 Released

2017 Kubernetes becomes ‘defacto’ orchestration engine & widespread adoption in enterprises

# 2019

Vmware announced  
Tanzu & Project Pacific





# Containers vs VMs

Containers and virtual machines have similar resource isolation and allocation benefits, but function differently because containers virtualize the operating system instead of hardware. Containers are more portable and efficient.

## CONTAINERS

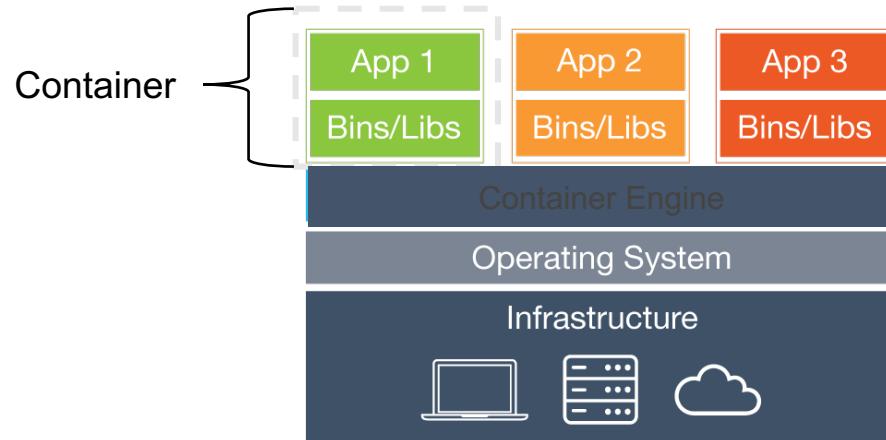
Containers are an abstraction at the app layer that packages code and dependencies together. Multiple containers can run on the same machine and share the OS kernel with other containers, each running as isolated processes in user space. Containers take up less space than VMs (container images are typically tens of MBs in size), can handle more applications and require fewer VMs and Operating systems.

## VIRTUAL MACHINES

Virtual machines (VMs) are an abstraction of physical hardware turning one server into many servers. The hypervisor allows multiple VMs to run on a single machine. Each VM includes a full copy of an operating system, the application, necessary binaries and libraries - taking up tens of GBs. VMs can also be slow to boot.

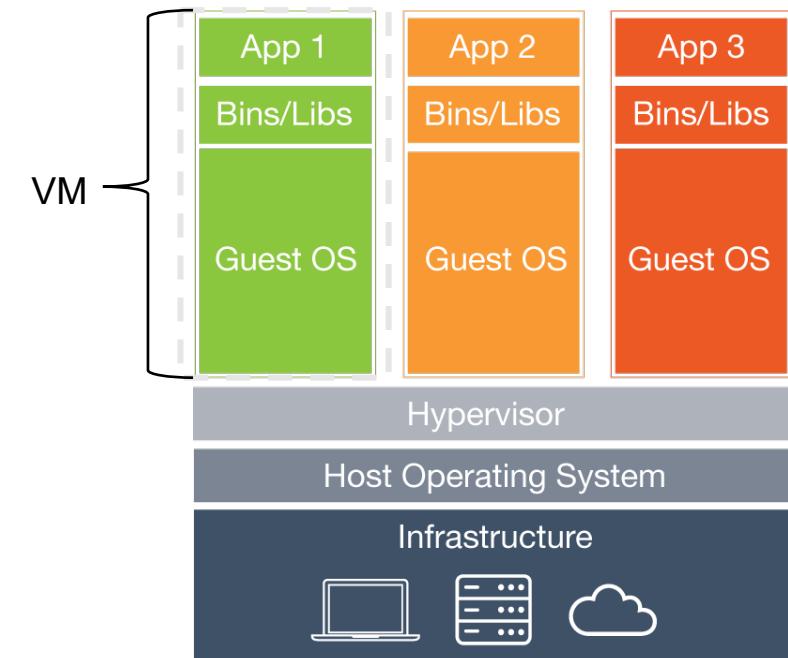
## CONTAINERS

- Shared Kernel
- Nascent technology created complexity
- Super fast to create a new container
- Containers are very portable



## VIRTUAL MACHINES

- Requires a full operating system
- Mature technology and management eco-system
- Slow to create a new VM
- VMs are not very portable



## DECLARATIVE

“I would like some tea”

...Define “tea”

Tea is an infusion of tea<sup>1</sup> leaves in a cup

<sup>1</sup>An infusion is obtained by letting the object steep a few minutes in hot water<sup>2</sup>

<sup>2</sup>Hot water is obtained by pouring it in a container and setting it on stove

## IMPERATIVE

Boil some water.

Pour it in a teapot.

Add tea leaves.

Steep for a while.

Serve in a cup

## DECLARATIVE

“I would like some tea”

- More complex
- Requires defining before hand
- Run automatically
- Knows current state of all components
- “Orchestration”
- YAML file

## IMPERATIVE

Boil some water.  
Pour it in a teapot.  
Add tea leaves.  
Steep for a while.  
Serve in a cup

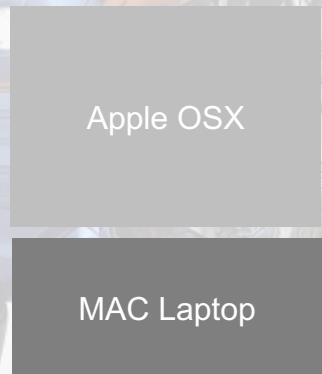
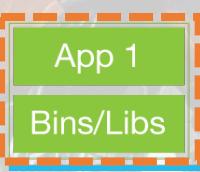
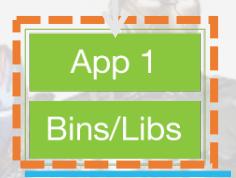
- Much simpler
- Requires scheduled action or human
- Run ad-hoc
- Lacks understanding of current state
- “Automation”
- Script

# Consistency and Decoupling Throughout the Process

Define the App Environment



Developers



Development

Unit Test

Integration

Prod

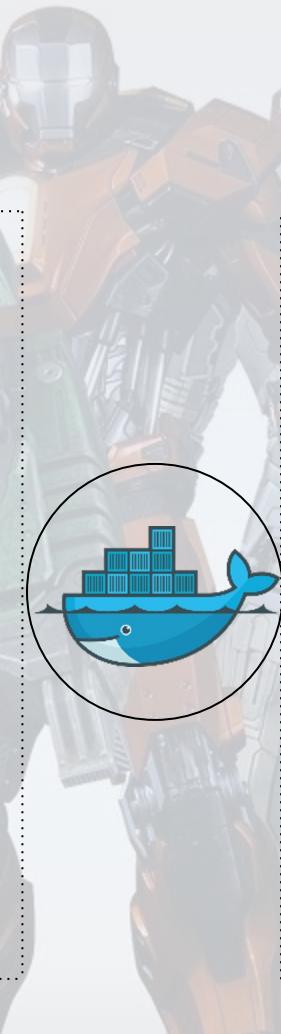
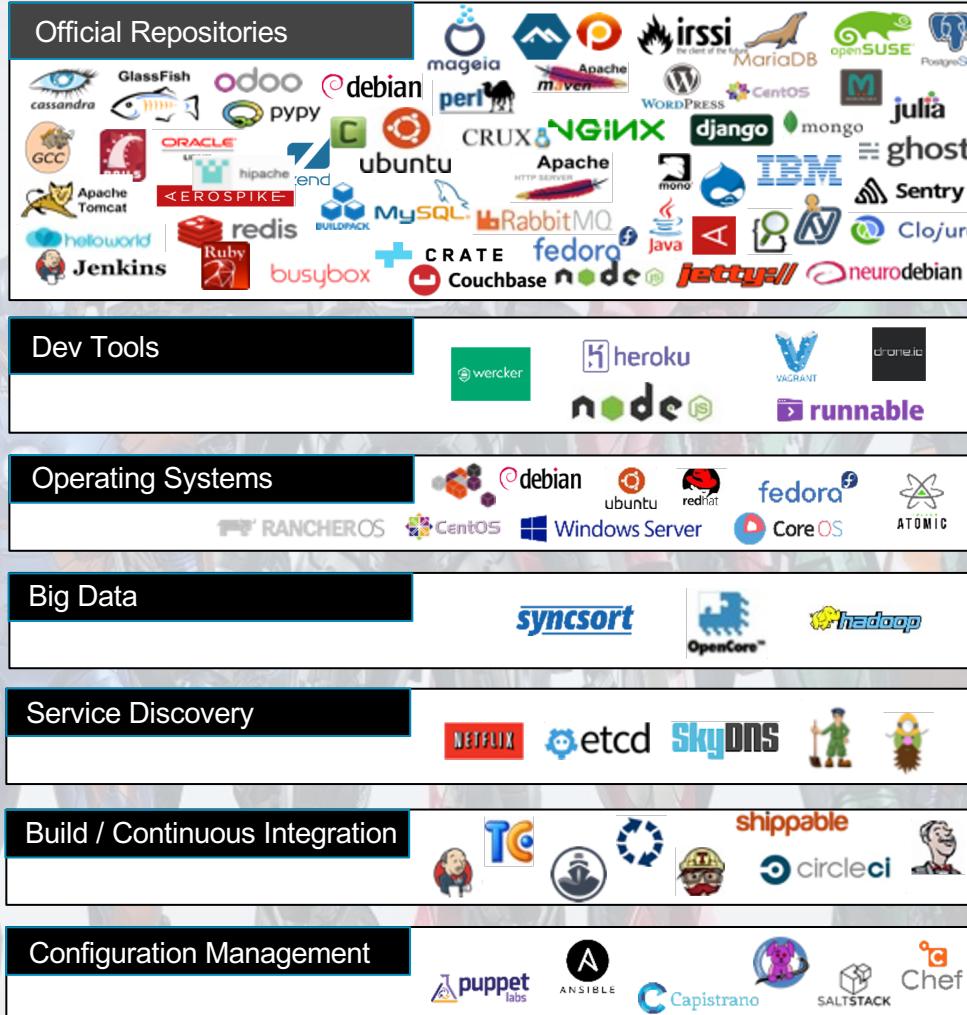
Operations



Container maintains consistent environment for application/code throughout the development and deployment process,

Container can be deployed and run virtually anywhere

# THE DOCKER ECOSYSTEM





Introducing  
**VMware Tanzu**

**Build** Modern Apps

**Run** Enterprise Kubernetes

**Manage** Kubernetes for Developers



# Introducing VMware Tanzu

Build Modern Apps  
Run Enterprise Kubernetes  
Manage Kubernetes for Developers



## Project Pacific: Rearchitecting vSphere with Native Kubernetes

Transforming vSphere into the App Platform of the Future

